

Gradient Descent Method

Data Intelligence and Learning ([DIAL](#)) Lab

Prof. Jongwuk Lee



Gradient Descent Method

Solving the Error Function

- How to **minimize the error** between $f(\mathbf{x})$ and y ?

$$E(\mathbf{w}) = \frac{1}{n} \sum_{(\mathbf{x}, y) \in \mathcal{D}} (y - f(\mathbf{x}; \mathbf{w}))^2$$

- $E(\mathbf{w})$ is **minimum** where **the derivative of $E(\mathbf{w})$ is zero**.

$$\frac{\partial E}{\partial w_0} = 0$$

$$\frac{\partial E}{\partial w_1} = 0$$

...

$$\frac{\partial E}{\partial w_d} = 0$$



Finding \mathbf{w} satisfying
these equations



Solving the Error Function

- Consider a simple linear regression with w_0 and w_1 .

$$\operatorname{argmin}_{w_0, w_1 \in [-\infty, \infty]} E(w_0, w_1) = \frac{1}{n} \sum_{i=1}^n \left(y^{(i)} - (w_1 x^{(i)} + w_0) \right)^2$$



Solving the Error Function

- Does the solution always exist?
- Yes, because the error function is **quadratic**.

$$E(\mathbf{w}) = \frac{1}{n} \sum_{(\mathbf{x}, y) \in \mathcal{D}} (y - f(\mathbf{x}; \mathbf{w}))^2$$

- Can we solve it if we choose **more complex models**?
- No, we may not.
- What should we do if we cannot solve it?



Solving the Optimization Problem



- The ML models can be trained **analytically** (e.g., **normal equation**) or are solved **numerically** (e.g., **gradient descent**).

- **Analytical solution**
 - ◆ It involves framing the problem in a **well-understood form** and calculating the **exact solution**.
 - ◆ In general, it is preferred because it is **faster**, and the solution is **exact**.

- **Numerical solution**
 - ◆ Make guesses for the solution.
 - ◆ It is necessary to validate whether it is solved well or not.

Challenge of Optimization Problems




➤ Finding the minimum value by the closed-form solution is **often difficult or impossible**.

- ◆ Quadratic functions in many variables

- System of equations for partial derivatives may be **ill-conditioned**.

- ◆ Other convex functions

- Global minimum exists, but there is **no closed-form solution**.
- E.g., logistic regression



- ◆ Nonlinear functions


- Partial derivatives are **not linear**.
- E.g., $f(x_1, x_2) = x_1(\sin(x_1 x_2)) + x_2^2$
- E.g., sum of transfer functions in neural networks

Solutions for Optimization Problems




➤ Several approximation methods

- ◆ **Gradient descent method**
- ◆ Newton method
- ◆ Gauss-Newton
- ◆ Levenberg-Marquardt
- ◆ BFGS
- ◆ Conjugate gradient



Convex



Non-convex


Intuition of Gradient Descent Method



What is the Gradient?


- The gradient of a function f is the collection of **all its partial derivatives into a vector**.
 - ◆ It is denoted as $\nabla f(\mathbf{w})$.
 - ◆ Each element in the gradient is the **slope of the function** along the direction of one of the variables.

$$\nabla f(\mathbf{w}) = \nabla f(w_1, w_2, \dots, w_d) = \begin{bmatrix} \frac{\partial f(\mathbf{w})}{\partial w_1} \\ \frac{\partial f(\mathbf{w})}{\partial w_2} \\ \vdots \\ \frac{\partial f(\mathbf{w})}{\partial w_d} \end{bmatrix}$$




How to Update w ?

- Randomly choose **an initial point w^0** .



How to Update w ?

- We want to update w satisfying $E(\mathbf{w}_{\text{new}}) < E(\mathbf{w})$.
 - ◆ When $\mathbf{w}_{\text{new}} = \mathbf{w} + \epsilon$, $E(\mathbf{w}_{\text{new}}) < E(\mathbf{w})$
- The slope at a position == differential value at a position
 - ◆ If $E(\mathbf{w})$ is differentiable, it is easy to find out the slope.




How to Update w ?

➤ Move the **reverse direction of the gradient**.


- ◆ If the slope is **positive**, move w to the **negative** direction.
- ◆ If the slope is **negative**, move w to the **positive** direction.

➤ When $w_{\text{new}} = w - \frac{\partial E}{\partial w}$, $E(w_{\text{new}}) < E(w)$




How to Update w ?

- Suppose that w is a vector and so is ϵ .
- When **the sum of two vectors can be quite large**, we need to control the size of ϵ .
 - ◆ This is important because if we make a large update ϵ to w , we **might miss out the global minimum** of error function $E(w)$.



Without the learning rate




With the learning rate

How to Update w ?


- Using a **learning rate** that limits the size of update for w

**learning rate:
controlling the step size**





How to Update w ?

- Move until **the gradient of $E(w)$ is zero.**



What is the Learning Rate?

- η is used to control the **step size** or **step length**.
- Too small η can incur **slow convergence**.
- Too large η can incur **overshoot** the minima and **diverge**.



Example: Choosing the Learning Rate



➤ Demo: <https://developers.google.com/machine-learning/crash-course/fitter/graph>

Experiment with different learning rates and see how they affect the number of steps required to reach the minimum of the loss curve. Try the exercises below the graph.

Set learning rate:



0.01

Execute single step:

Step


0

Let's try $\eta = 0.1, 1.0, 4.0$.

Reset the graph:

Reset

Loss vs. Weight



Gradient Descent (GD)

learning rate:
controlling the step size

Randomly choose an initial solution \mathbf{w}^0 ,

Repeat


$$\mathbf{w}^{t+1} = \mathbf{w}^t - \eta \left. \frac{dE}{d\mathbf{w}} \right|_{\mathbf{w}=\mathbf{w}^t}$$

Until the **stopping condition** is satisfied

- Fixed number of iterations
- $|E(\mathbf{w}^{t+1}) - E(\mathbf{w}^t)|$ is very small.

Finding Some Minimal Positions

- How to **minimize the error** between $f(x)$ and y ?
- Depending on w^0 , it can find different minimums.



Find out any of those minimums




Example: Gradient Descent Method

Simple Linear Model

- Finding a **linear model** that fits a given data

$$f(x; w_0, w_1) = w_1x + w_0$$



$$Data = \{(0.0, 0.0), (1.0, 1.0), (1.0, 2.0)\}$$

Solving the Optimization Problem



➤ **Finding $w = (w_0, w_1)$ that minimize $E(w_0, w_1)$**

◆ For simplicity, we use sum instead of mean.

$$E(w_0, w_1) = \sum_{(x,y) \in \mathcal{D}} (y - f(x; w_0, w_1))^2$$

$$f(x; w_0, w_1) = w_1 x + w_0$$

$$Data = \{(0.0, 0.0), (1.0, 1.0), (1.0, 2.0)\}$$

$$E(w_0, w_1) = (0.0 - f(0.0; w_0, w_1))^2 + (1.0 - f(1.0; w_0, w_1))^2 + (2.0 - f(1.0; w_0, w_1))^2$$

$$E(w_0, w_1) = (0.0 - w_0)^2 + (1.0 - (w_0 + w_1))^2 + (2.0 - (w_0 + w_1))^2$$

$$E(w_0, w_1) = 2w_1^2 + 3w_0^2 - 6w_1 - 6w_0 + 4w_1w_0 + 5$$

Calculating the Gradient

➤ Calculate the gradient for a given error function.

$$E(w_0, w_1) = 2w_1^2 + 3w_0^2 - 6w_1 - 6w_0 + 4w_1w_0 + 5$$

$$\frac{\partial E}{\partial w_1} = 4w_1 + 4w_0 - 6$$

$$\frac{\partial E}{\partial w_0} = 4w_1 + 6w_0 - 6$$

Randomly choose an initial solution, w_0^0, w_1^0 .

Repeat

$$w_0^{t+1} = w_0^t - \eta \left. \frac{\partial E}{\partial w_0} \right|_{w_0=w_0^t, w_1=w_1^t}$$

$$w_1^{t+1} = w_1^t - \eta \left. \frac{\partial E}{\partial w_1} \right|_{w_0=w_0^t, w_1=w_1^t}$$

Until the stopping condition is satisfied

Calculating the Gradient

- Find w_0, w_1 that minimizes the error function.
 - ◆ Choose a learning rate η , e.g., $\eta = 0.1$.
 - ◆ Initialize w_0^0, w_1^0 as random values, e.g., $w_0^0 = 1, w_1^0 = 1$

Randomly choose an initial solution, w_0^0, w_1^0 .

Repeat

$$w_0^{t+1} = w_0^t - \eta(4w_1^t + 6w_0^t - 6)$$

$$w_1^{t+1} = w_1^t - \eta(4w_1^t + 4w_0^t - 6)$$

Until the stopping condition is satisfied

How does the GD Work?

➤ Update w_0^t, w_1^t iteratively.

Randomly choose an initial solution, w_0^0, w_1^0 .

Repeat

$$w_0^{t+1} = w_0^t - \eta(4w_1^t + 6w_0^t - 6)$$

$$w_1^{t+1} = w_1^t - \eta(4w_1^t + 4w_0^t - 6)$$

Until the stopping condition is satisfied

$$w_0^0 = 1$$

$$w_1^0 = 1$$

$$w_0^1 = 1 - 0.1(4 \times 1 + 6 \times 1 - 6) = 0.6$$

$$w_1^1 = 1 - 0.1(4 \times 1 + 4 \times 1 - 6) = 0.8$$

$$w_0^2 = 0.6 - 0.1(4 \times 0.8 + 6 \times 0.6 - 6) = 0.54$$

$$w_1^2 = 0.8 - 0.1(4 \times 0.8 + 4 \times 0.6 - 6) = 0.84$$

$$w_0^3 = 0.54 - 0.1(4 \times 0.84 + 6 \times 0.54 - 6) = 0.480$$

$$w_1^3 = 0.84 - 0.1(4 \times 0.84 + 4 \times 0.54 - 6) = 0.888$$

How does the GD Work?

➤ Update w_0^t, w_1^t iteratively.

$$w_0^3 = 0.480$$
$$w_1^3 = 0.888$$

$$w_0^4 = 0.480 - 0.1(4 \times 0.888 + 6 \times 0.480 - 6) = 0.4368$$
$$w_1^4 = 0.888 - 0.1(4 \times 0.888 + 4 \times 0.480 - 6) = 0.9408$$

...

$$w_0^{100} = 0.00007713$$
$$w_1^{100} = 1.49989171$$

Randomly choose an initial solution, w_0^0, w_1^0 .

Repeat

$$w_0^{t+1} = w_0^t - \eta(4w_1^t + 6w_0^t - 6)$$


$$w_1^{t+1} = w_1^t - \eta(4w_1^t + 4w_0^t - 6)$$

Until the stopping condition is satisfied

Solution of the Linear Model

- Finding a **linear model** that fits a given data

$$f(x; w_0, w_1) = 1.49989171x + 0.00007713$$



$$Data = \{(0.0, 0.0), (1.0, 1.0), (1.0, 2.0)\}$$



Details: Gradient Descent Method

Taylor Series Expansion

➤ **Taylor series function** is an infinite sum of terms that are expressed by the function's derivatives at a single point.

◆ $f^{(k)}(a)$: k -th derivative of the function

$$f(x) = f(a) + \frac{f^{(1)}(a)}{1!}(x - a) + \frac{f^{(2)}(a)}{2!}(x - a)^2 + \dots + \frac{f^{(k)}(a)}{k!}(x - a)^k + \dots$$



$$f(x) = \sum_{i=0}^{\infty} \frac{f^{(i)}(a)}{i!}(x - a)^i$$


Example: Taylor Series Expansion

➤ $f(x) = e^x$ and $a = 0$

$$e^x = f(0) + \frac{f^{(1)}(0)}{1!}(x - 0) + \frac{f^{(2)}(0)}{2!}(x - 0)^2 + \dots + \frac{f^{(k)}(0)}{k!}(x - 0)^k + \dots$$



$$e^x = 1 + x + \frac{x^2}{2!} + \dots + \frac{x^k}{k!} + \dots$$



Taylor Series Expansion for $f(x + \epsilon)$



➤ Suppose that $x = x_1 + \epsilon$ and $a = x_1$

$$f(x) = f(a) + \frac{f^{(1)}(a)}{1!}(x - a) + \frac{f^{(2)}(a)}{2!}(x - a)^2 + \dots + \frac{f^{(k)}(a)}{k!}(x - a)^k + \dots$$

$$(x - a) = (x_1 + \epsilon - x_1) = \epsilon \quad \downarrow$$

$$f(x_1 + \epsilon) = f(x_1) + \frac{f^{(1)}(x_1)}{1!}\epsilon + \frac{f^{(2)}(x_1)}{2!}\epsilon^2 + \dots + \frac{f^{(k)}(x_1)}{k!}\epsilon^k + \dots$$



$$f(x_1 + \epsilon) = f(x_1) + \epsilon f^{(1)}(x_1) + O(\epsilon^2)$$

The rest of the terms are restricted by $O(\epsilon^2)$.

Gradient Descent in One Dimension



- The first-order approximation $f(x + \epsilon)$ is given by $f(x)$ and $f^{(1)}(x)$ at x .

$$f(x + \epsilon) = f(x) + \epsilon f^{(1)}(x) + O(\epsilon^2)$$

- We want to choose ϵ to preserve $f(x + \epsilon) \lesssim f(x)$

$$f(x + \epsilon) - f(x) = \epsilon f^{(1)}(x) + O(\epsilon^2) \lesssim 0$$



Assuming ϵ is too small,
 $O(\epsilon^2)$ is negligible.

$$f(x + \epsilon) - f(x) \approx \epsilon f^{(1)}(x) \lesssim 0$$

Gradient Descent in One Dimension



➤ Pick a fixed step size $\eta > 0$ and choose $\epsilon = -\eta f^{(1)}(x)$.

- ◆ If $f^{(1)}(x) \neq 0$, then $\epsilon f^{(1)}(x) = -\eta \left(f^{(1)}(x)\right)^2 < 0$, where $\left(f^{(1)}(x)\right)^2$ is always positive.

➤ When we choose η small enough,

$$f(x + \epsilon) - f(x) \approx \epsilon f^{(1)}(x) \lesssim 0$$



$$f(x - \eta f^{(1)}(x)) - f(x) \approx -\eta \left(f^{(1)}(x)\right)^2 \lesssim 0$$

Note that η is a small positive number.

Gradient Descent in One Dimension




➤ Finally, we arrive at

$$f(\mathbf{x} - \eta f^{(1)}(\mathbf{x})) - f(\mathbf{x}) \approx -\eta \left(f^{(1)}(\mathbf{x}) \right)^2 \lesssim 0 \quad \Rightarrow \quad f(\mathbf{x} - \eta f^{(1)}(\mathbf{x})) \lesssim f(\mathbf{x})$$

➤ This means that, if we use

$$\mathbf{x} \leftarrow \mathbf{x} - \eta f^{(1)}(\mathbf{x})$$

➤ The value of $f(\mathbf{x})$ might decline.



Gradient Descent in Multi-Dimension



- Let us consider the situation $\mathbf{x} \in \mathbb{R}^d$.

$$f(\mathbf{x} + \boldsymbol{\epsilon}) = f(\mathbf{x}) + \boldsymbol{\epsilon}^T \nabla f(\mathbf{x}) + O(\boldsymbol{\epsilon}^2) \approx f(\mathbf{x}) + \boldsymbol{\epsilon}^T \nabla f(\mathbf{x})$$

- Up to second order terms in $\boldsymbol{\epsilon}$, the direction of steepest decent is given by **the negative gradient** $-\nabla f(\mathbf{x})$.
- Choosing a suitable learning rate $\eta > 0$ yields the gradient descent algorithm.

$$\mathbf{x} \leftarrow \mathbf{x} - \eta \nabla f(\mathbf{x})$$




Convex Function

What is the Convex Function?

- A function f is convex \Leftrightarrow the function f is **below** any line segment between two points on f .

$$f(tx + (1 - t)x') \leq tf(x) + (1 - t)f(x') \text{ for any } x, x' \in X \text{ and } t \in [0, 1]$$




What is the Convex Function?

➤ A function $f: X \rightarrow \mathbb{R}$ is defined as a convex function if


$$f(tx + (1 - t)x') \leq tf(x) + (1 - t)f(x') \text{ for any } x, x' \in X \text{ and } t \in [0, 1]$$

$$f\left(\frac{x+x'}{2}\right) \leq \frac{f(x)+f(x')}{2} \text{ for any } x, x' \in X \text{ and } t = 0.5$$

Midpoint convex





Convex



Non-convex

Examples of Convex Functions

- Quadratic function: $f(x) = x^2$
- Exponential functions: $f(x) = 2^x$
- Negative logarithm function: $f(x) = -\ln x$



- Convex optimization overview

- ◆ <http://web.stanford.edu/class/cs224n/readings/cs229-cvxopt.pdf>


Convex vs. Non-convex Function

➤ Convex function


- ◆ $\nabla f(\mathbf{w}^*) = 0 \Leftrightarrow \mathbf{w}^*$ is a global minimum.
- ◆ Example: linear regression, logistic regression

➤ Non-convex function

- ◆ $\nabla f(\mathbf{w}^*) = 0 \Leftrightarrow \mathbf{w}^*$ is a global min, local min or saddle point (also called stationary points.)
- ◆ Most algorithms only converge to stationary points.
- ◆ Example: neural networks




Convex




Non-convex

Convex vs. Non-convex Function


- Given $y = x_1^4 - 2x_1^2 + x_2^2$, it has **two symmetric local min $(-1,0)$ and $(1,0)$** , and a **saddle point $(0,0)$** between them.



Saddle point



Convex



Non-convex

Q&A

