

# Neural Networks Basics

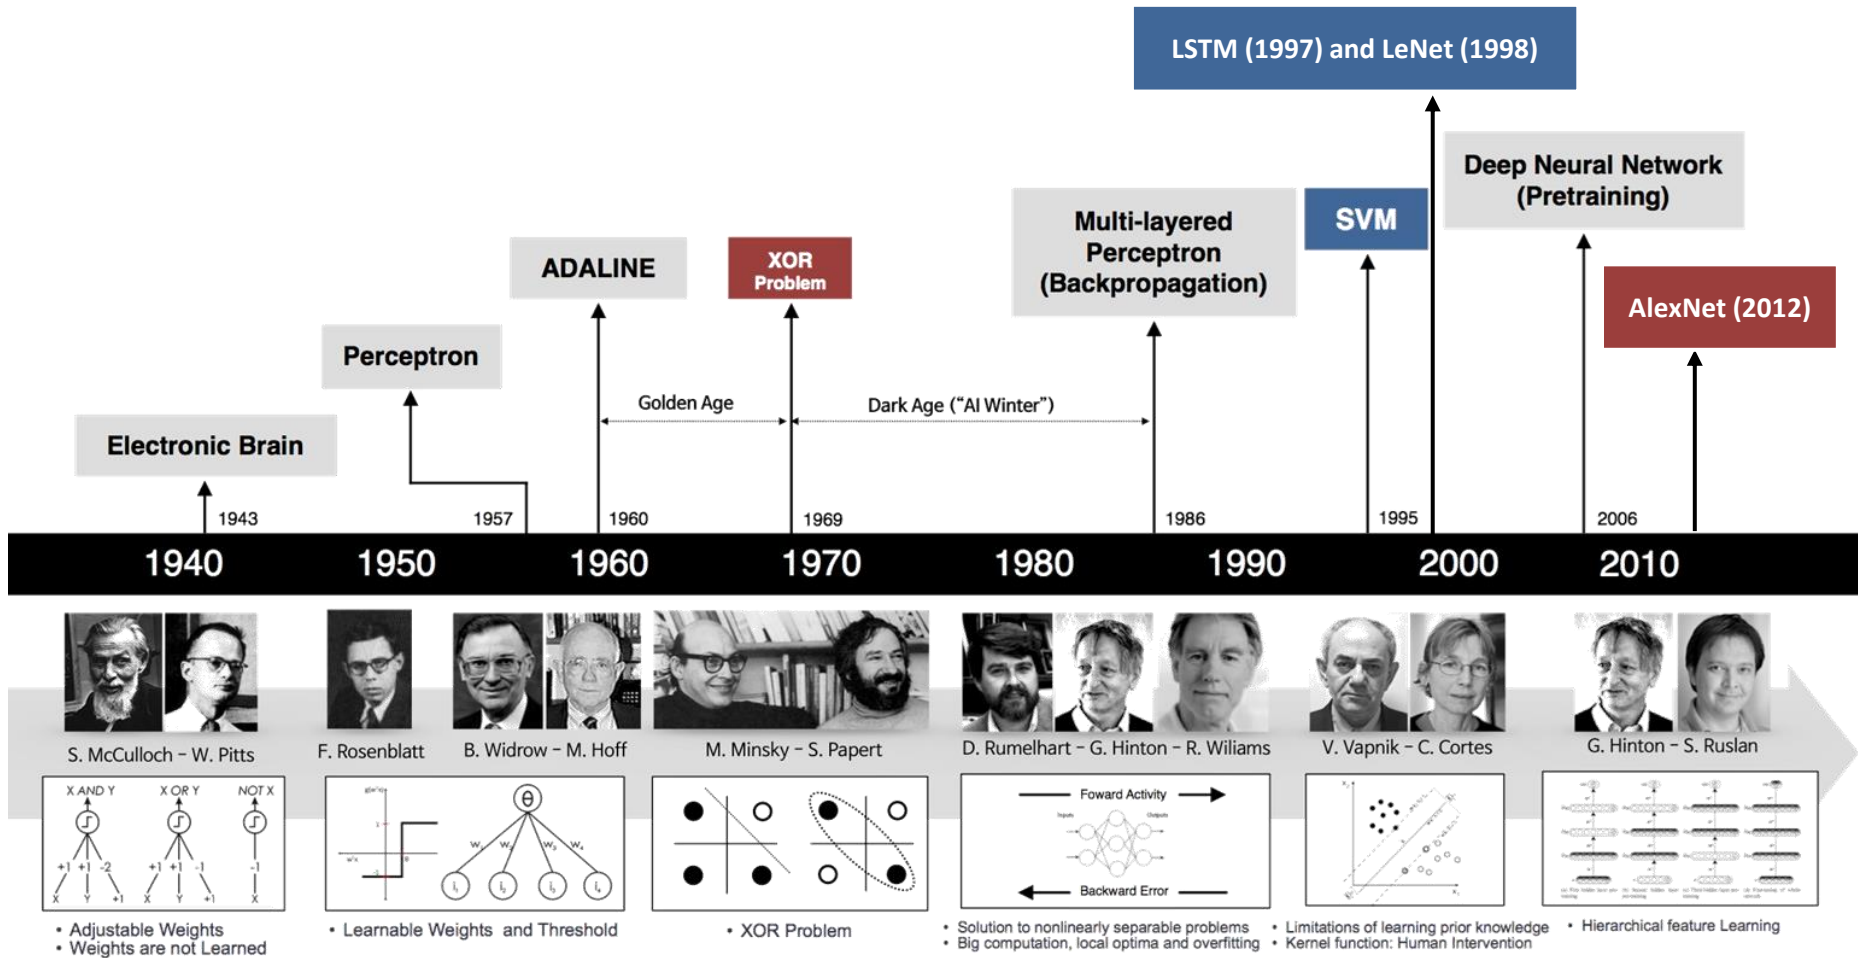
Data Intelligence and Learning ([DIAL](#)) Lab

Prof. Jongwuk Lee




# Perceptron

# Brief History of Neural Networks



# Concept of Neurons

- Receive inputs from other neurons (via synapses).
- When input exceeds a threshold, “fires.”
  - ◆ Sends output along axon to other neurons.
- Human brain:  $10^{11}$  neurons





# Perceptron: Artificial Neuron (1957)



- A **neuron** is activated when the correlation between the input **x** and a pattern **w** exceeds a **threshold**.
- ◆ It is called an **artificial neuron**.
  - ◆ It mimics the function of human **neurons**.



Frank Rosenblatt



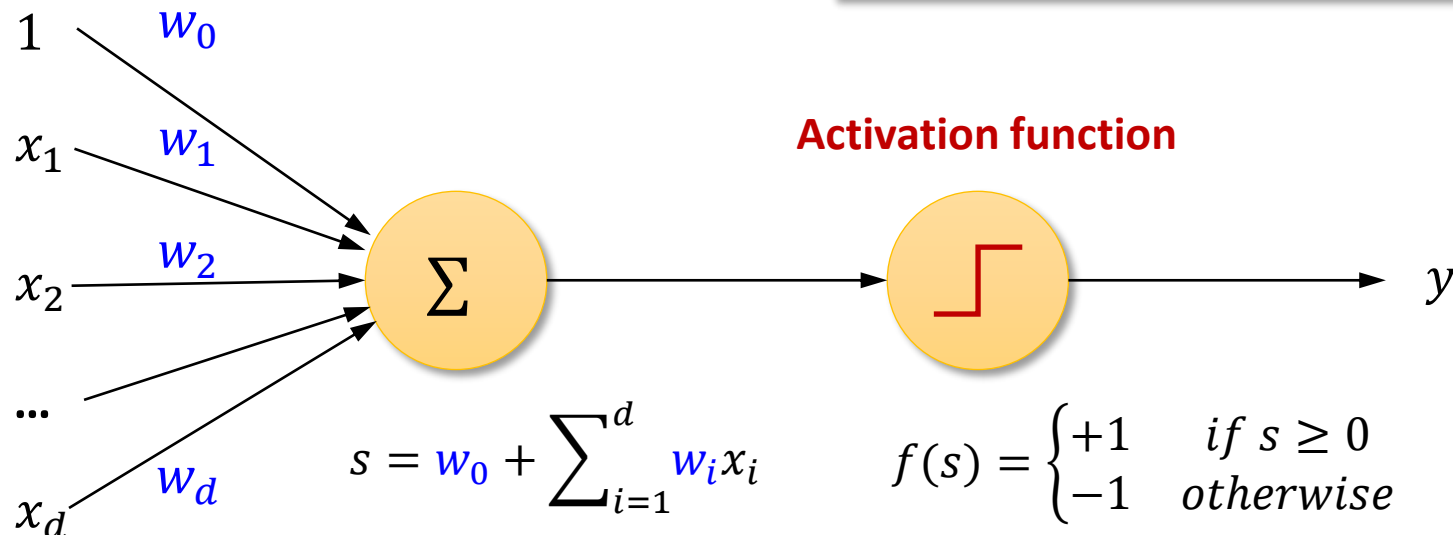
# What is the Perceptron?

- Linear combination of input  $x$ :

$$s = w_0 + \sum_{i=1}^d w_i x_i$$

- Nonlinear transformation of  $s$ :

$$y = f(s), \text{ where } \begin{cases} +1 & \text{if } s \geq 0 \\ -1 & \text{otherwise} \end{cases}$$





# What is the Perceptron?



➤ Linear combination of input  $\mathbf{x}$ :

$$s = \mathbf{w}^T \mathbf{x} = \sum_{i=0}^d w_i x_i$$

➤ Nonlinear transformation of  $s$ :



$$f(s) = \begin{cases} +1 & \text{if } s \geq 0 \\ -1 & \text{otherwise} \end{cases}$$

# Formulating a Learning Classifier

- We regard the output  $y$  into  $\{+1, -1\}$ .
  
- Find  $f(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$  that satisfies the condition.
  - ◆  $y = +1$  if  $\mathbf{w}^T \mathbf{x} \geq 0$
  - ◆  $y = -1$  if  $\mathbf{w}^T \mathbf{x} < 0$
  
- Prediction:  $\hat{y} = \text{sign}(f(\mathbf{x})) = \text{sign}(\mathbf{w}^T \mathbf{x})$ 
  - ◆ Note:  $\text{sign}(\cdot)$  returns  $+1$  or  $-1$ .



# Formulating a Learning Classifier

- Given training data  $\{(\mathbf{x}^{(i)}, y^{(i)}): 1 \leq i \leq n\}$
- Hypothesis function:  $\hat{y} = \text{sign}(\mathbf{w}^T \mathbf{x})$
- **0-1 loss function:** # of inputs such that  $y^{(i)} \neq \text{sign}(\mathbf{w}^T \mathbf{x}^{(i)})$

$$E(\mathbf{w}) = - \sum_{(\mathbf{x}^{(i)}, y^{(i)}) \in \mathcal{D}} y^{(i)} \mathbf{w}^T \mathbf{x}^{(i)} \mathbb{I}[\text{mistake on } \mathbf{x}^{(i)}]$$

$$\mathbb{I}[\text{mistake on } \mathbf{x}^{(i)}] = \begin{cases} 1 & \text{if } y^{(i)} \neq \text{sign}(\mathbf{w}^T \mathbf{x}^{(i)}) \\ 0 & \text{otherwise} \end{cases}$$

# Formulating a Learning Classifier

## ➤ Representing a simplified a 0-1 loss function

$$E(\mathbf{w}) = - \sum_{(\mathbf{x}^{(i)}, y^{(i)}) \in \mathcal{D}} y^{(i)} \mathbf{w}^T \mathbf{x}^{(i)} \mathbb{I}[\text{mistake on } \mathbf{x}^{(i)}]$$



$$E(\mathbf{w}) = \sum_{(\mathbf{x}^{(k)}, y^{(k)}) \in \mathcal{S}} -y^{(k)} (\mathbf{w}^T \mathbf{x}^{(k)})$$

$\mathcal{S}$  is a set of **incorrect samples**.

# Property of the 0-1 Loss Function

➤ For any sample  $(\mathbf{x}^{(k)}, y^{(k)}) \in \mathcal{S}$ ,  $-y^{(k)}(\mathbf{w}^T \mathbf{x}^{(k)}) > 0$ .

➤  $E(\mathbf{w}) \geq 0$

$$E(\mathbf{w}) = \sum_{(\mathbf{x}^{(k)}, y^{(k)}) \in \mathcal{S}} -y^{(k)}(\mathbf{w}^T \mathbf{x}^{(k)})$$

➤ If  $\mathbf{w}$  is optimal, then  $E(\mathbf{w}) = 0$ .

◆  $E(\mathbf{w})$  tends to increase with the number of **incorrect samples**.

# Computing the Derivative of $w$

➤ How to compute the derivative of  $w_j$

$$\frac{\partial E(\mathbf{w})}{\partial w_j} = \sum_{(\mathbf{x}^{(k)}, y^{(k)}) \in \mathcal{S}} \frac{\partial}{\partial w_j} \left( -y^{(k)} \left( w_0 x_0^{(k)} + \dots + \textcolor{red}{w_j x_j^{(k)}} + \dots + w_d x_d^{(k)} \right) \right)$$



Other values are regarded as constants.

$$\frac{\partial E(\mathbf{w})}{\partial w_j} = \sum_{(\mathbf{x}^{(k)}, y^{(k)}) \in \mathcal{S}} -y^{(k)} x_j^{(k)}, \quad \forall j = 0, 1, \dots, d$$

# Perceptron Learning Algorithm (PLA)



## ➤ Batch version

Initialize a random weight  $\mathbf{w}$ .

Repeat

$\mathcal{S} \leftarrow \emptyset$

**Step 1: finding a set of incorrect samples**

for  $i = 1$  to  $n$

$\hat{y} \leftarrow \text{sign}(\mathbf{w}^T \mathbf{x}^{(i)})$

if  $(\hat{y} \neq y^{(i)})$   $\mathcal{S} \leftarrow \mathcal{S} \cup (\mathbf{x}^{(i)}, y^{(i)})$

**Step 2: updating the weight  $\mathbf{w}$**

if  $(\mathcal{S} \neq \emptyset)$

$\Delta \mathbf{w} = - \sum_{(\mathbf{x}^{(k)}, y^{(k)}) \in \mathcal{S}} y^{(k)} \mathbf{x}^{(k)}$

$\mathbf{w} \leftarrow \mathbf{w} - \eta \Delta \mathbf{w}$

Until  $(\mathcal{S} = \emptyset)$

# Perceptron Learning Algorithm (PLA)



## ➤ Stochastic version

Initialize a random weight  $\mathbf{w}$ .

Repeat

$\text{quit} \leftarrow \text{true}$

    for  $i \leftarrow 1$  to  $n$

**Step 1: finding an incorrect sample**

$$\hat{y} \leftarrow \text{sign}(\mathbf{w}^T \mathbf{x}^{(i)})$$

**if**  $(\hat{y} \neq y^{(i)})$

$\text{quit} \leftarrow \text{false}$

$$\Delta \mathbf{w} = -y^{(i)} \mathbf{x}^{(i)}$$

$$\mathbf{w} \leftarrow \mathbf{w} - \eta \Delta \mathbf{w}$$

**Step 2: updating the weight  $\mathbf{w}$**

Until ( $\text{quit} = \text{true}$ )

# Perceptron Learning Algorithm (PLA)



- Start with a random weight vector.
- Given an input, predict its class.
  - ◆ For the mistake on a **positive sample**,

$$\mathbf{w}_{new} = \mathbf{w} - \eta \nabla E(\mathbf{w}; (\mathbf{x}^{(i)}, y^{(i)})) = \mathbf{w} + \eta \mathbf{x}^{(i)}$$

- ◆ For the mistake on a **negative sample**,

$$\mathbf{w}_{new} = \mathbf{w} - \eta \nabla E(\mathbf{w}; (\mathbf{x}^{(i)}, y^{(i)})) = \mathbf{w} - \eta \mathbf{x}^{(i)}$$

- Repeat it until there is no mistake.

# Perceptron Learning Algorithm (PLA)



- Start with a random weight vector.
- Given an input, predict its class.
  - ◆ For the mistake on a **positive sample**,

$$\mathbf{w}_{new} = \mathbf{w} - \eta \nabla E(\mathbf{w}; (\mathbf{x}^{(i)}, y^{(i)})) = \mathbf{w} + \mathbf{1}\mathbf{x}^{(i)}$$

- ◆ For the mistake on a **negative sample**,

$$\mathbf{w}_{new} = \mathbf{w} - \eta \nabla E(\mathbf{w}; (\mathbf{x}^{(i)}, y^{(i)})) = \mathbf{w} - \mathbf{1}\mathbf{x}^{(i)}$$

When  $\eta = 1$

- Repeat it until there is no mistake.



# Recap: Learning a Linear Classifier



➤ Execute the algorithm until no mistakes are encountered.

Randomly choose an initial solution  $\mathbf{w}^0$ .

For  $t = 0, 1, \dots$

Find a **mistake sample**  $(\mathbf{x}^{(i)}, y^{(i)})$  of  $\mathbf{w}^t$   
 $\text{sign}(\mathbf{w}^T \mathbf{x}^{(i)}) \neq y^{(i)}$


Correct the mistake by

$$\mathbf{w}^{t+1} = \mathbf{w}^t + y^{(i)} \mathbf{x}^{(i)}$$


Until no more mistake is found

Return last  $\mathbf{w}^t$  as the learned model.

If  $y = +1$




If  $y = -1$



# AND Operation

## ➤ How to represent the AND operation



$$w_1 = 1.0, w_2 = 1.0, w_0 = -1.5$$


| $x_1$ | $x_2$ | $\Sigma$ | $y$ |
|-------|-------|----------|-----|
| 0     | 0     | -1.5     | 0   |
| 0     | 1     | -0.5     | 0   |
| 1     | 0     | -0.5     | 0   |
| 1     | 1     | +0.5     | 1   |

$$1.0x_1 + 1.0x_2 - 1.5 = 0$$

# OR Operation

## ➤ How to represent the OR operation

$$w_1 = 1.0, w_2 = 1.0, w_0 = -0.5$$




| $x_1$ | $x_2$ | $\Sigma$ | $y$ |
|-------|-------|----------|-----|
| 0     | 0     | -0.5     | 0   |
| 0     | 1     | 0.5      | 1   |
| 1     | 0     | 0.5      | 1   |
| 1     | 1     | 1.5      | 1   |

$$1.0x_1 + 1.0x_2 - 0.5 = 0$$


# Linear Separability

- If PLA halts (i.e., no more mistakes),
  - ◆ **(necessary condition)**  $D$  allows some  $w$  to make no mistake.
- Call such  $D$  **linearly separable**.




Linear separable

Good!



Linear non-separable

Need a linear model that allows some errors.




Linear non-separable

Need a non-linear model.

# XOR Operation (1969)



➤ How to train a linear decision boundary?



**“No one on earth had found a viable way to train”** by Marvin Minsky

- Guaranteed to converge if **linearly separable**.
- Otherwise, many simple functions are NOT learnable.




# Multilayer Perceptron (MLP)

# Multilayer Perceptron (MLP)



- It is a neural network of **multiple** artificial neurons.



# Multilayer Perceptron (MLP)



## ➤ Structure

### ◆ Input layer


- Simply pass the input values to the next layer.
- # of nodes = # of inputs

### ◆ Hidden layer

- There can be several hidden layers.
- # of nodes should be given.

### ◆ Output layer


- # of nodes = # of outputs



➤ MLP is also called the **feed-forward neural network**.



# Multilayer Perceptron (MLP)



$$y = \frac{1}{1 + \exp\left(-\left(w_{31} \underbrace{\left(\frac{1}{1 + \exp(-(w_{11}x_1 + w_{12}x_2 + w_{10}))}\right)}_{h_1} + w_{32} \underbrace{\left(\frac{1}{1 + \exp(-(w_{21}x_1 + w_{22}x_2 + w_{20}))}\right)}_{h_2} + w_{30}\right)\right)}$$

# How to Solve XOR Operation?

$$w_{11} = 1.0, w_{12} = 1.0, w_{10} = -1.5 \quad w_{21} = 1.0, w_{22} = 1.0, w_{20} = -0.5 \quad w_{31} = -1.0, w_{32} = 1.0, w_{30} = -0.5$$

$$h_1 = \text{sigmoid}(x_1 + x_2 - 1.5)$$


$$h_2 = \text{sigmoid}(x_1 + x_2 - 0.5)$$

$$\hat{y} = \text{sigmoid}(-x_1 + x_2 - 0.5)$$

| $x_1$ | $x_2$ | $\Sigma$ | $h_1$ |
|-------|-------|----------|-------|
| 0     | 0     | -1.5     | 0     |
| 0     | 1     | -0.5     | 0     |
| 1     | 0     | -0.5     | 0     |
| 1     | 1     | +0.5     | 1     |



| $x_1$ | $x_2$ | $\Sigma$ | $h_2$ |
|-------|-------|----------|-------|
| 0     | 0     | -0.5     | 0     |
| 0     | 1     | 0.5      | 1     |
| 1     | 0     | 0.5      | 1     |
| 1     | 1     | 1.5      | 1     |

| $h_1$ | $h_2$ | $\Sigma$ | $\hat{y}$ |
|-------|-------|----------|-----------|
| 0     | 0     | -0.5     | 0         |
| 0     | 1     | 0.5      | 1         |
| 0     | 1     | 0.5      | 1         |
| 1     | 1     | -0.5     | 0         |



# How to Solve XOR Operation?


➤ A single neuron divides a region into **two subregions**.




| $x_1$ | $x_2$ | $\Sigma$ | $h_1$ |
|-------|-------|----------|-------|
| 0     | 0     | -1.5     | 0     |
| 0     | 1     | -0.5     | 0     |
| 1     | 0     | -0.5     | 0     |
| 1     | 1     | +0.5     | 1     |

# How to Solve XOR Operation?

➤ A single neuron divides a region into **two subregions**.




| $x_1$ | $x_2$ | $\Sigma$ | $h_2$ |
|-------|-------|----------|-------|
| 0     | 0     | -1.5     | 0     |
| 0     | 1     | 0.5      | 1     |
| 1     | 0     | 0.5      | 1     |
| 1     | 1     | 1.5      | 1     |




# How to Solve XOR Operation?

## ➤ Combining two neurons with AND operation




# How to Solve XOR Operation?



Mapping the **regions** in the original space  
to the **points** in another space


# Non-linear Classifier using MLP

➤ How to classify two classes?



# Non-linear Classifier using MLP


## ➤ Hyperspace partitioning



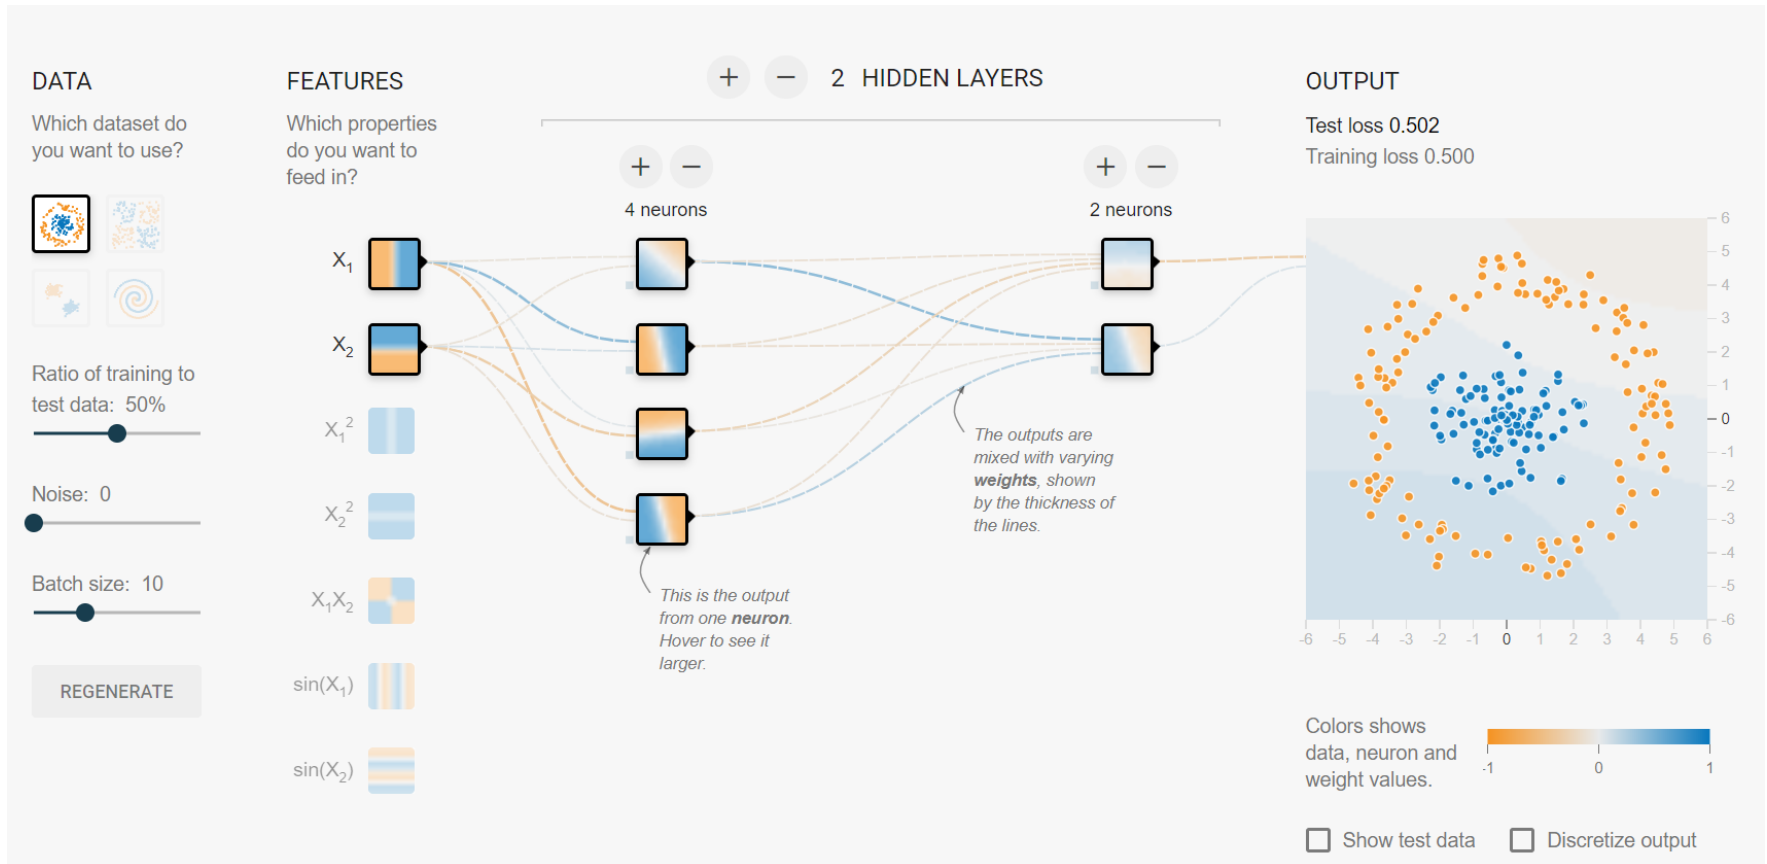


# Non-linear Classifier using MLP

➤ Use three nodes in the hidden layer.



# Example: Visualization of MLP





# Deep Neural Networks (DNNs)

# Recap: Machine Learning 1-2-3

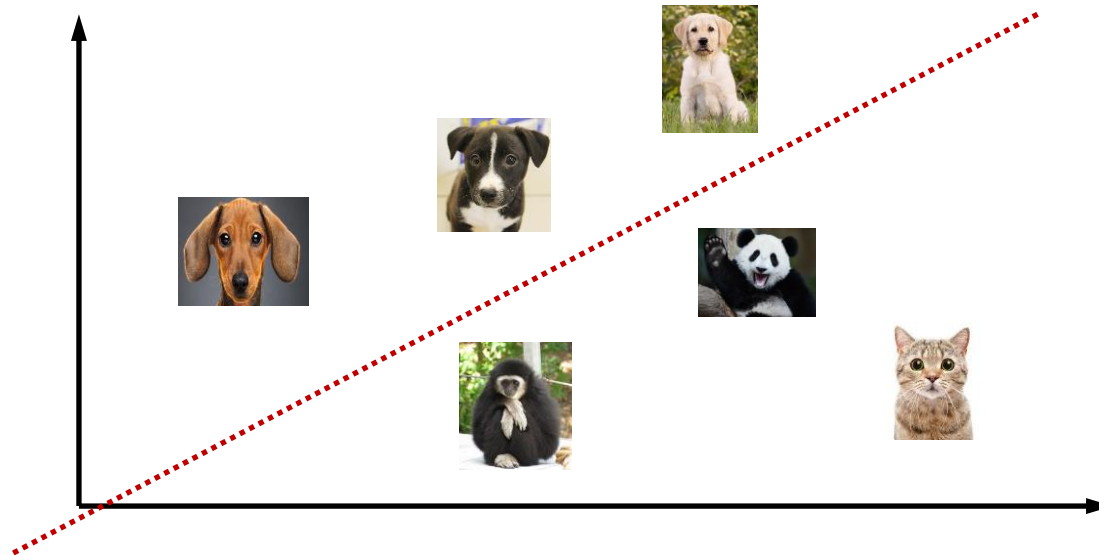
- Collect data and extract features.
  - Choose the **hypothesis function  $\mathcal{H}$**  and the **loss function  $\mathcal{L}$** .
  - Find an optimal parameter that minimizes the empirical loss.
- 
- Q: How to represent feature vectors for images?
  - A: It is difficult to design the vector.



# Simplest Case: Linear Classifiers



- Consider the original perceptron model on raw data.
- To classify the image of “**dog**” (+1) or “**not dog**” (-1), find the line that separates the +1’s from the -1’s.



# Non-Linear Features, Linear Classifiers





➤ Most problems need **non-linear features**.

- ◆ Image classification, machine translation, speech recognition, etc.

➤ How to represent non-linear features?

- ◆ Non-linear kernels in SVM
- ◆ Explicit design of features, e.g., SIFT, HOG




➤ Then, linear machines are good enough.

# How to Represent Good Features?




- Raw data live high dimensionalities. However, data lie in low-dimensional manifolds.
- Can we discover a **manifold** for our data?
  - ◆ Hypothesis: **Semantically similar things lie closer together than semantically dissimilar things.**




# Example: Manifolds of Digits

- There are good features (manifolds).
- Each image has 784 dimensions on the MNIST dataset.
  - ◆  $28 \text{ pixels} \times 28 \text{ pixels} = 784 \text{ dimensions}$



Entangled features  
(PCA)



Disentangled features  
(t-SNE)



# Why Learn Features?



## ➤ Manually designed features

- ◆ Often, take **a lot of time to** implement
- ◆ Often, take a lot of time to validate
- ◆ Often, they are **incomplete**, as one cannot know if they are optimal for the task



## ➤ Learned features

- ◆ If data is enough, easy to learn
- ◆ Compact and specific **to the task**



## ➤ Time spent for designing features vs. **Time spent on designing architectures**

# Shallow Learning vs. Deep Learning



**Artificial intelligence**

**Hard-coded knowledge & the rule-driven approach is used.**



**Machine learning**

**Extract the patterns & co-relation between features.**

It is difficult to know whether all features are effective and necessary.



**Representation learning**

**Learns the representation instead of the features.**

There are many factors of variations for a given representation.




**Deep learning**

**Complex representations are expressed in terms of simple representations.**

# Shallow Learning vs. Deep Learning



## ➤ Traditional machine learning




Hand-crafted feature extractor




Trainable classifier



## ➤ End-to-end learning → features are also **learned from data!**



Trainable feature extractor




Trainable classifier



# Deep Neural Networks (DNNs)



- It consists of a pipeline of successive, differentiable modules (transformations).
  - ◆ Each module's output is used as the input for the next module.
- Learns hierarchical representation from data.
- Each sequential module produces a higher abstraction feature.



# Deep Neural Networks (DNNs)

- A family of **parametric**, **non-linear**, and **hierarchical** representation learning functions
  - ◆ They are optimized with **stochastic gradient descent** to **encode domain knowledge**, i.e., domain invariances and stationarity.

$$a_L(x; \theta_{1,...,L}) = h_L(h_{L-1}(\dots h_1(x, \theta_1), \theta_{L-1}), \theta_L)$$


- Given training corpus  $\{X, Y\}$ , find optimal parameters.

$$\theta^* \leftarrow \operatorname{argmin}_{\theta} \sum_{(x,y) \in (X,Y)} \mathcal{L}(y, a_L(x; \theta_{1,...,L}))$$

# Deep Neural Networks (DNNs)

- A series of **hierarchically** connected functions
- This hierarchy can be very **complex**!

$$a_L(x; \theta_{1,...,L}) = h_L(h_{L-1}(\dots h_1(x, \theta_1), \theta_{L-1}), \theta_L)$$




*Feedforward architecture*

# Deep Neural Networks (DNNs)





➤ What if we go deeper?



# Input Layer of DNNs

- The input is represented as a vector.
  - ◆ Sometimes, it requires preprocessing, e.g., **normalization**




**X**

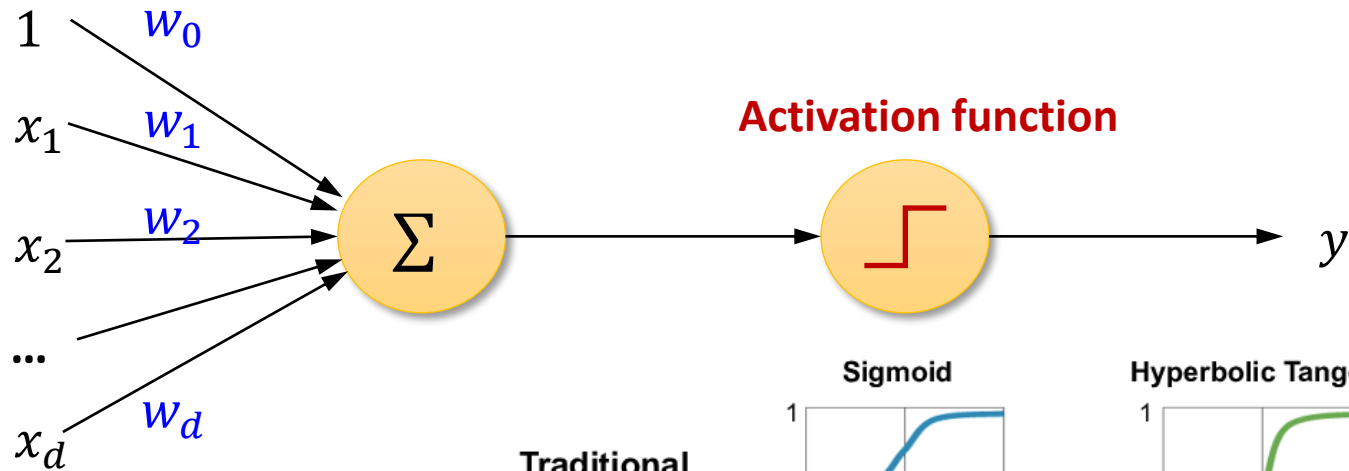


# Hidden Layers of DNNs

- Each neuron takes a weighted linear combination of the previous layer.
  - ◆ Provide a single **aggregated** value to the next layer.





# Nodes at the Hidden Layers






There are various activation functions.

**Traditional Non-Linear Activation Functions**



**Modern Non-Linear Activation Functions**





$\alpha = \text{small const. (e.g. 0.1)}$

# Output Layer of DNNs

➤ For output, there are three cases:

- ◆ Regression:  $y = w^T h$
- ◆ Binary classification:  $y = \sigma(w^T h)$
- ◆ Multi-class classification:  $y = \text{softmax}(w^T h)$



# Q&A

