

# Run TPC-C and Analyze the Results

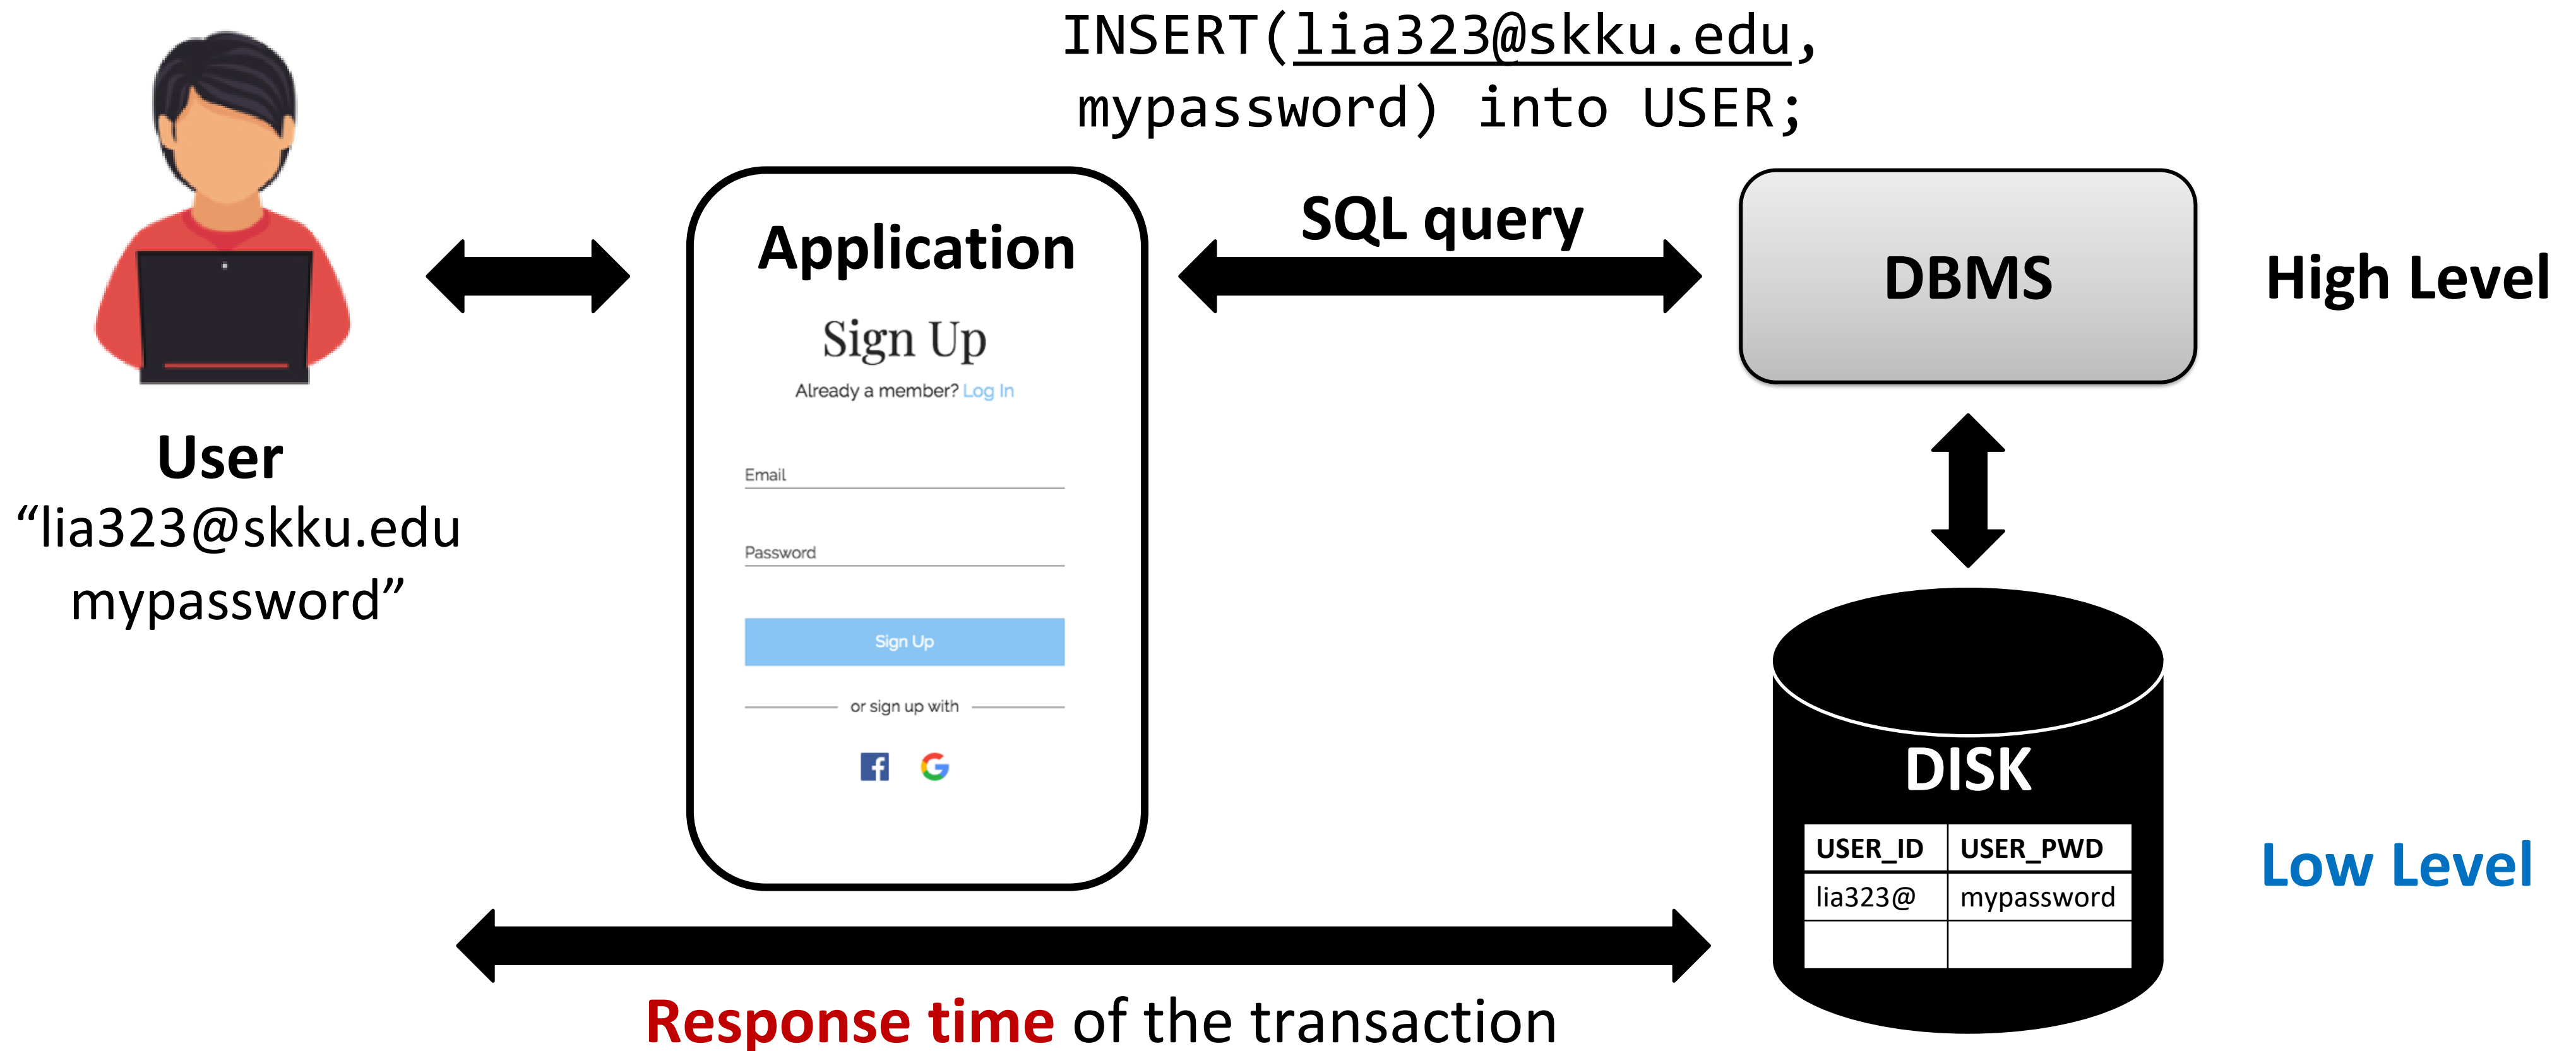
Bo-Hyun Lee  
[lia323@skku.edu](mailto:lia323@skku.edu)



**VLDB**  
Lab.



# DBMS Performance



# Benchmarking: What and Why

- Benchmarking is a process of **measuring** the performance of a given application and **comparing** it with other similar workloads
  - to discover if there is a performance gap
  - to improve its performance
- Benchmark is *domain-specific*:
  - The more general the benchmark, the less useful it is
  - A benchmark is a distillation of the essential attributes of a particular workload
- Desirable attributes:
  - Relevant → It should be meaningful within the target domain
  - Understandable/acceptable → Vendors and users embrace it
  - Scalable/coverable → It should not oversimplify the typical environment



# What is the TPC?

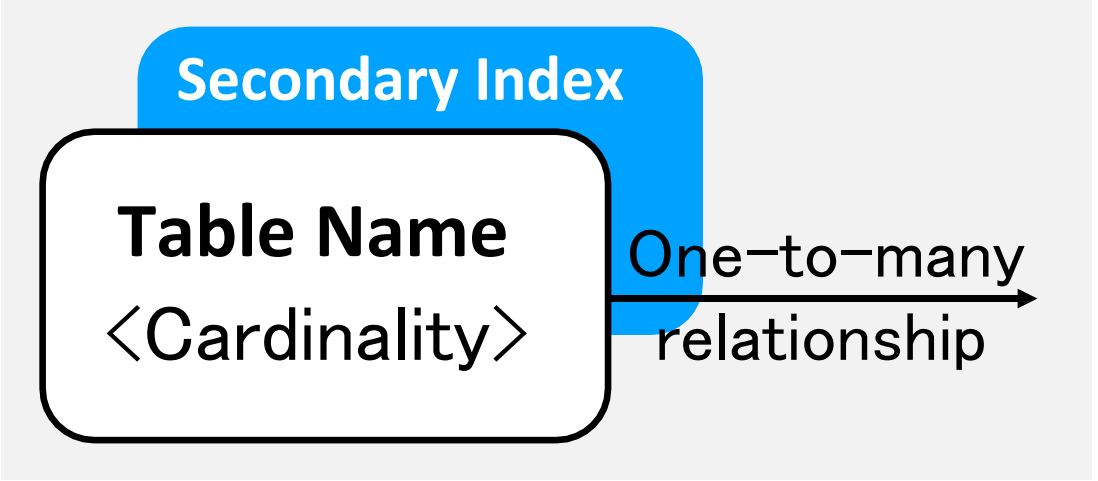
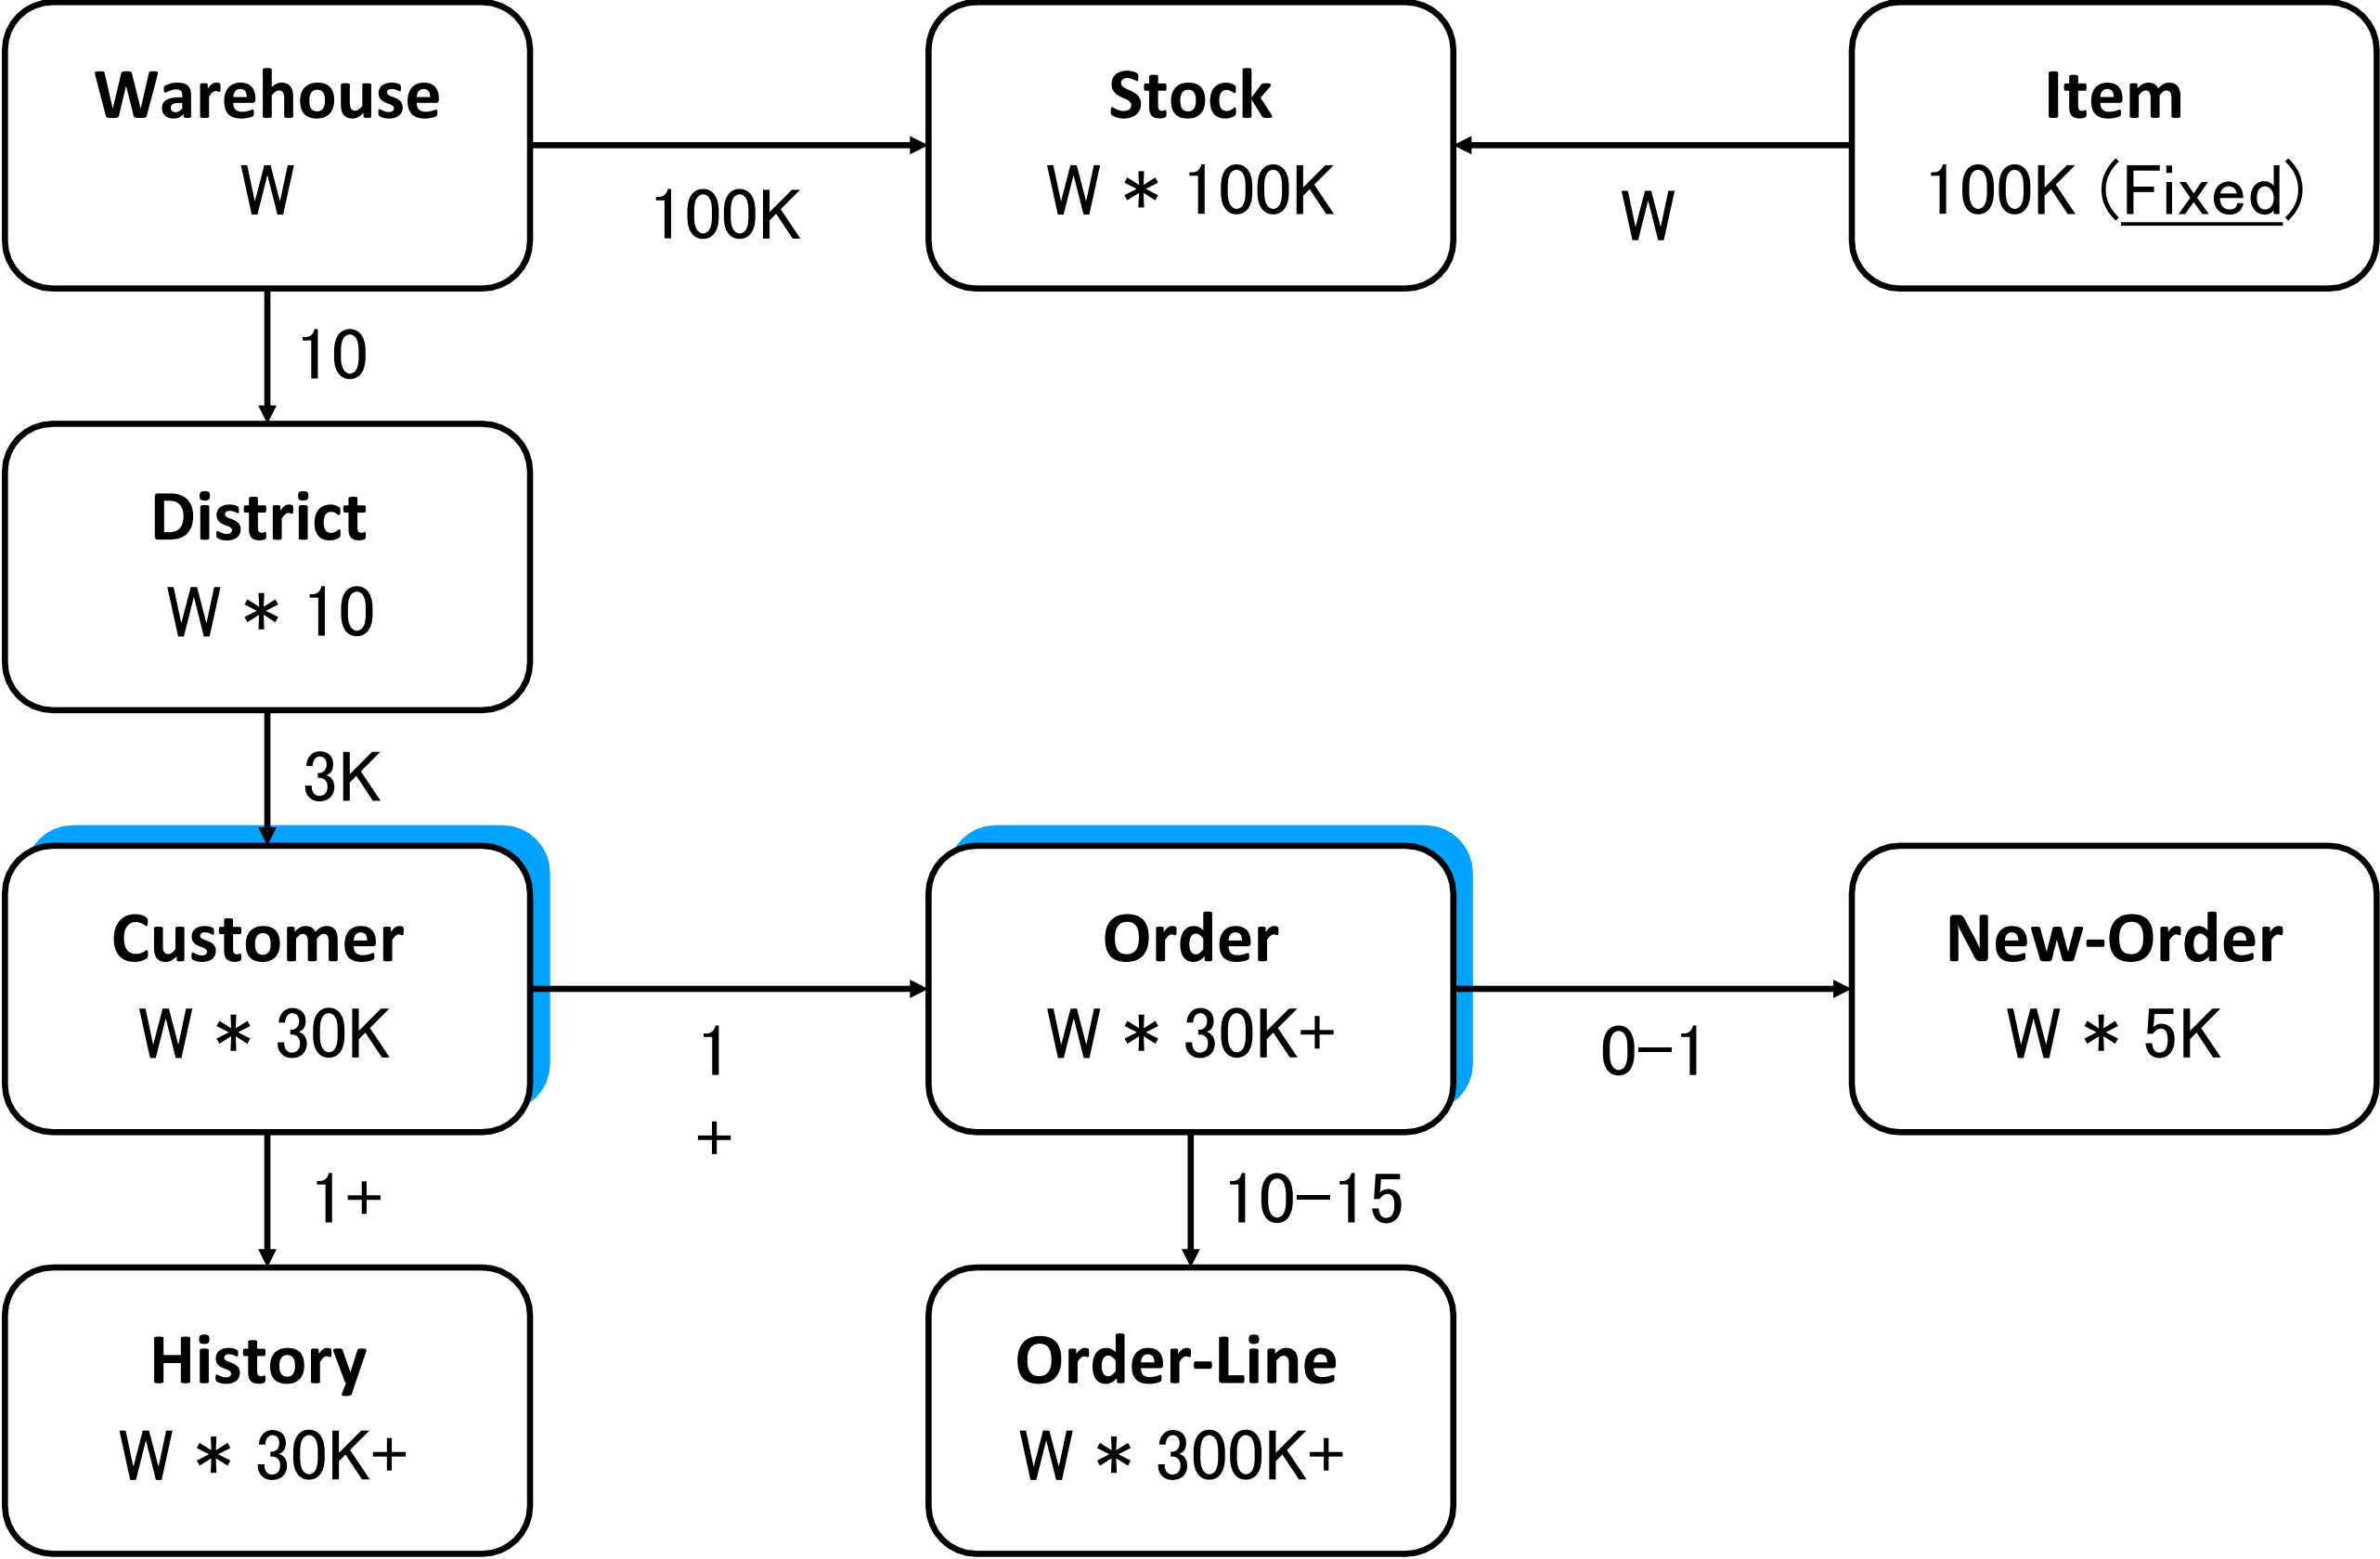
- TPC = Transaction Processing Performance Council
  - Founded in Aug 1988 by Omri Serlin and 8 vendors
- De facto industry-standards body for enterprise benchmark

[illegible]

# TPC-C

- A 29-year-old **industry-standard OLTP benchmark** used to measure the performance of databases
  - It simulates an *e-commerce* or *retail* company
- 5 types of well-defined transactions:
  - New-Order (Read/Write), Payment (Read/Write), Delivery (Read/Write), Order-Status (Read Only), Stock-Level (Read Only)
- **Throughput of TPC-C** = The number of **New-Order** transactions executed per minute
  - Transactions per minute Count (TpmC)
- **Random I/O intensive** workload
  - 65% Reads, 35% writes

# TPC-C Entity/Relationship Diagram





# New Order Transaction

1. Select(whouse-id) from Warehouse
  2. Select(dist-id, whouse-id) from District
  3. Update(dist-id, whouse-id) in District
  4. Select(customer-id, dist-id, whouse-id) from Customer
  5. Insert into Order
  6. Insert into New-Order
  7. For each item (10 items):
    - (a) Select(item-id) from Item
    - (b) Select(item-id, whouse-id) from Stock
    - (c) Update(item-id, whouse-id) in Stock
    - (d) Insert into Order-Line
  8. Commit
- Place an order for on average 10 items from a warehouse
- Insert the order
- Update the corresponding stock level for each item



# Payment Transaction

1. `Select(whouse-id) from Warehouse`
  2. `Select(dist-id,whouse-id) from District`
  - 3.(a) Case 1: `Select(customer-id,dist-id,whouse-id) from Customer`
  - (b) Case 2: `Non-Unique-Select(customer-name,dist-id,whouse-id) from Customer`
  4. `Update(whouse-id) in Warehouse`
  5. `Update(dist-id,whouse-id) in District`
  6. `Update(customer-id,dist-id,whouse-id) in Customer`
  7. `Insert into History`
  8. `Commit`
- Process a payment for a customer
- Update balances and other data





# Order Status Transaction

- 1.(a) Case 1: `Select(customer-id,dist-id,whouse-id) from Customer`  
(b) Case 2: `Non-Unique-Select(customer-name,dist-id,whouse-id) from Customer`
  2. `Select(Max(order-id),customer-id) from Order`
  3. for each item in the order:
    - (a) `Select(order-id) from Order-Line`
  4. Commit
- Return the status of a customer's last order



# Delivery Transaction

1. For each district within the warehouse (i.e. ten times):
  - (a) `Select(Min(order-id), whouse-id, dist-id) from New-Order`
  - (b) `Delete(order-id) from New-Order`
  - (c) `Select(order-id) from Order`
  - (d) `Update(order-id) Order`
  - (e) For each item in the order (i.e. ten times):
    - i. `Select(order-id) from Order-Line`
    - ii. `Update(order-id) Order-Line`
  - (f) `Select(customer-id) from Customer`
  - (g) `Update(customer-id) Customer`

Process orders corresponding to 10 pending orders,  
one for each district, with 10 items per order

2. Commit



# Stock Level Transaction

```
SELECT d_next_o_id INTO :o_id
FROM District
WHERE d_w_id = :w_id AND d_id = :d_id ;

SELECT COUNT(DISTINCT (s_i_id)) INTO :stock_count
FROM Order-Line, Stock
WHERE
  o_l_w_id = :w_id AND
  o_l_d_id = :d_id AND o_l_o_id < :o_id AND
  o_l_o_id ≥ (:o_id - 20) s_w_id = :w_id AND
  s_i_id = o_l_i_id AND s_quantity < :threshold ;
```

Examine the quantity of stock for the items  
ordered by each of the last 20 orders in a district



# Relation Access Pattern: S/I/U/D

| Relation Name | New Order                | Payment                | Order Status | Delivery                   | Stock Level | Comment        |
|---------------|--------------------------|------------------------|--------------|----------------------------|-------------|----------------|
| Warehouse     | Select(1)                | Select(1)<br>Update(1) |              |                            |             | Small Table    |
| District      | Select(1)<br>Update(1)   | Select(1)<br>Update(1) |              |                            | Select(1)   | Small Table    |
| Customer      | Select(1)                | Select(1)<br>Update(1) | Select(2,2)  | Select(10)<br>Update(10)   |             | Skewed Update  |
| Stock         | Select(10)<br>Update(10) |                        |              |                            | Select(200) | Skewed Update  |
| Item          | Select(10)               |                        |              |                            |             | Skewed RD-Only |
| Order         | Insert(1)                |                        | Select(1)    | Select(10)<br>Update(10)   |             | Growing        |
| New-order     | Insert(1)                |                        |              | Select(10)<br>Delete(10)   |             | Cyclic Reuse   |
| Order-line    | Insert(10)               |                        | Select(10)   | Select(100)<br>Update(100) | Select(200) | Growing        |
| History       |                          | Insert(1)              |              |                            |             | Growing        |



# OLTP vs. OLAP

- **OLTP (On-Line Transaction Processing):**
  - Handle a transactional system with operational data using a lot of short transactions (i.e., based on SELECT, INSERT, UPDATE, DELETE)
  - Mixed read/write workloads
  - Examples: ATM machines, online banking/booking/shopping, etc.
- **OLAP (On-Line Analytical Processing):**
  - Handle an analytical system with historical data using complex queries (i.e., based on SELECT)
  - Heavy read workloads (for large data)
  - Examples: Sales analysis, market research, forecasting, etc.



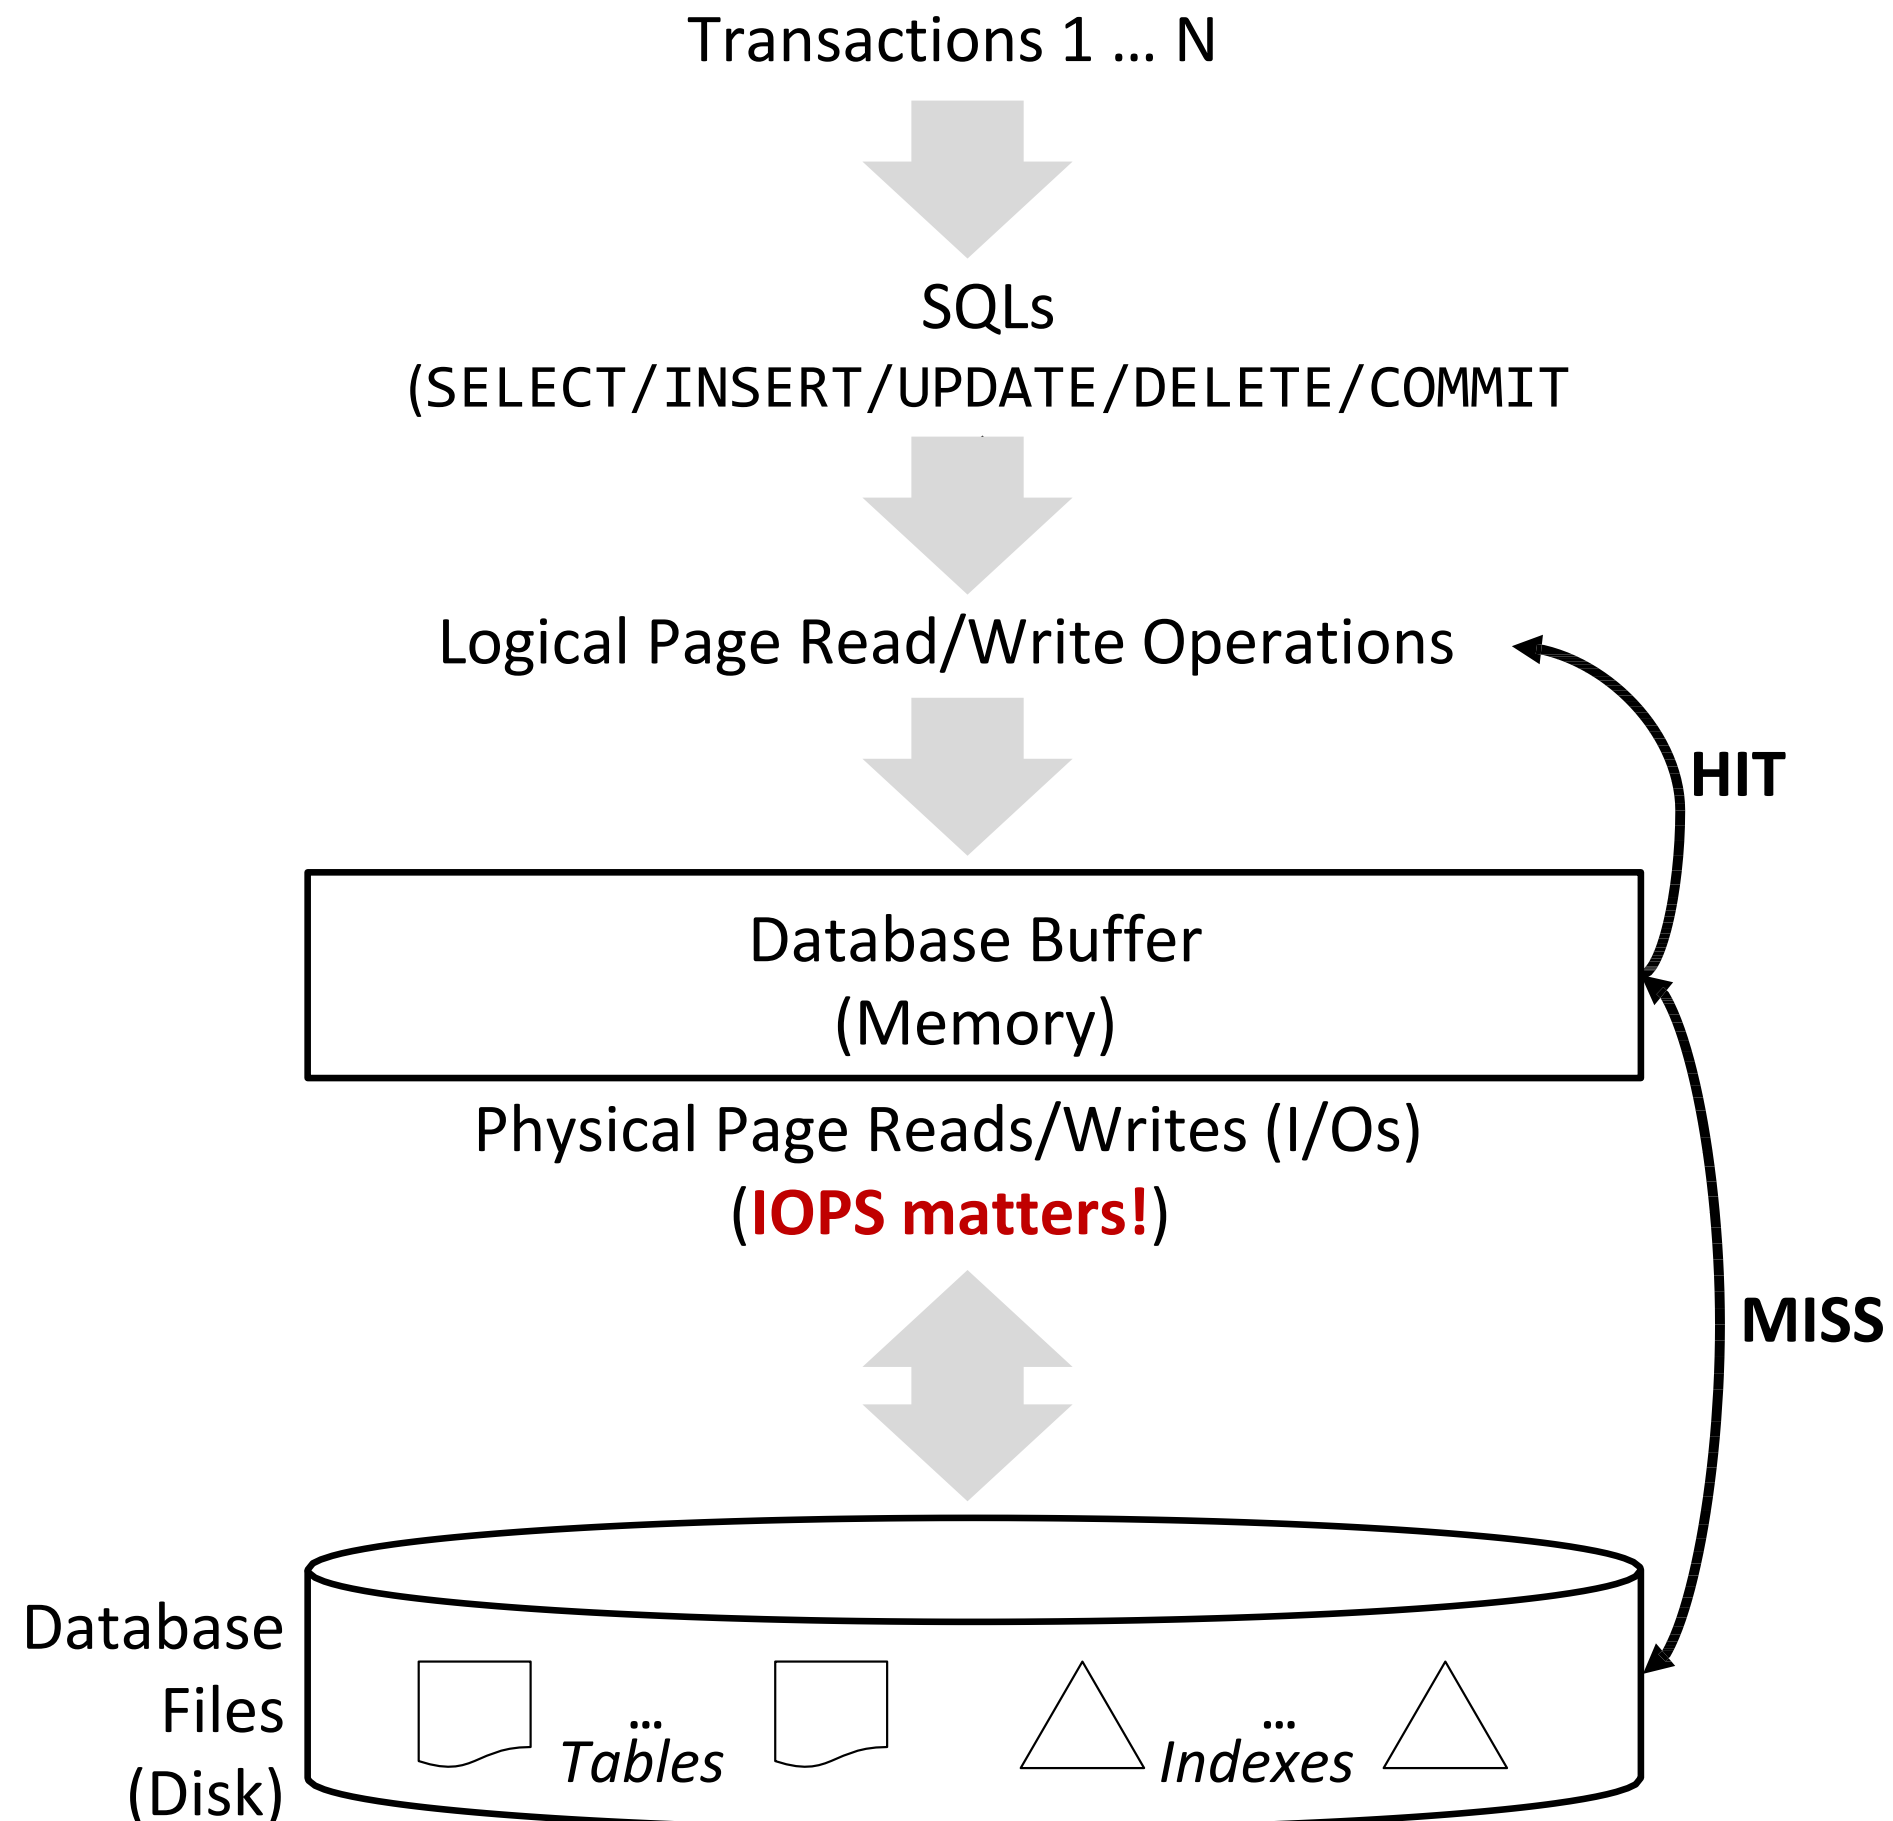


# OLTP vs. OLAP: Access Patterns

- Both OLTP and OLAP uses *indexes*, but they access and handle data differently
- OLTP is a **random read/write** workload
  - It consists of random accesses with typically 2KB~16KB request sizes
  - It uses *index scans*
- OLAP is a **heavy-read** workload
  - Typically, it consists of sequential accesses to a large amount of data
  - It mainly uses *full table scans*



# Database I/O Architecture



- A transaction
  - = A sequence of SQL statements
  - = A sequence of Reads and Writes
- SELECT reads tuples from page(s)
- INSERT/DELETE/UPDATE change records in page(s)
  - Thus, they access one or more pages
- When page(s) is in buffer (i.e., **HIT**): DRAM operation
- Otherwise (i.e., **MISS**): Disk I/Os
  - In case of dirty victim, write the page to storage
  - Read page(s) from storage



# Run the TPC-C Benchmark and Analyze the Results

- This week, you will learn to monitor the system performance while running the TPC-C benchmark on MySQL
- You will also learn what those performance metrics mean
- Refer to week 2 contents in <https://github.com/LeeBohyun/SWE3033-S20223>



# Reference

- 1 MySQL, “MySQL Community Downloads”, <https://dev.mysql.com/downloads/mysql/>
- 2 TPC, “TPC BENCHMARK C”, [http://www.tpc.org/tpc\\_documents\\_current\\_versions/pdf/tpc-c\\_v5.11.0.pdf](http://www.tpc.org/tpc_documents_current_versions/pdf/tpc-c_v5.11.0.pdf)
- 3 Veronica Lagrange, Changho Choi, Vijay Balakrishnan, “Accelerating OLTP performance with NVMe SSDs”, SDC 2016, [https://www.snia.org/sites/default/files/SDC/2016/presentations/solid\\_state\\_storage/VeronicaLaGrange\\_Accelerating\\_OLTP\\_Performance\\_V6.pdf](https://www.snia.org/sites/default/files/SDC/2016/presentations/solid_state_storage/VeronicaLaGrange_Accelerating_OLTP_Performance_V6.pdf)
- 4 Percona-Lab, “tpcc-mysql”, Github repository, <https://github.com/Percona-Lab/tpcc-mysql>
- 5 Scott T. Leutenegger and Daniel Dias, “A modeling study of the TPC-C benchmark”, SIGMOD Rec. 22, 2 (June 1, 1993), 22–31
- 6 Most of the slides are made by Mijin An([meeeejin@gmail.com](mailto:meeeejin@gmail.com))

