

Buffer Management of MySQL/InnoDB

(2) Write Requests Part 2

Bo-Hyun Lee
lia323@skku.edu

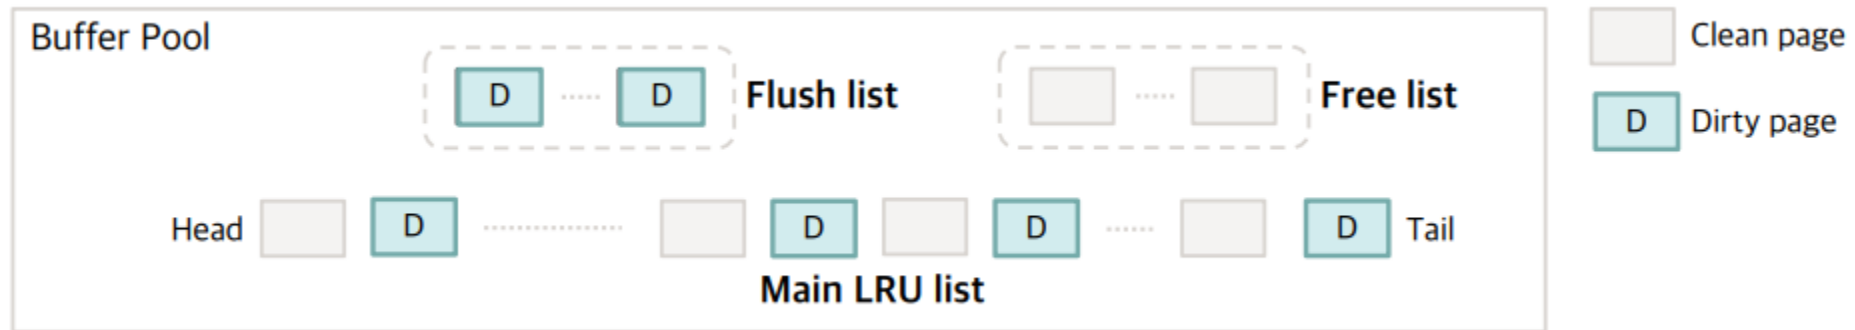


Index

- Three Types of Disk Write
- InnoDB Redo Log, LSN, and Checkpoint
- Flush List Flush
- Innodb parameters
- PA1



Lists of Buffer Blocks



- **Free list**
 - Contains free (empty) buffer frames
- **LRU list**
 - Contains all the blocks holding a file page
- **Flush list**
 - Contains the blocks holding file pages that have been **modified** in the memory but not written to disk yet (i.e., **dirty**)

Three Types of Disk Writes

- **Single Page Flush:**
 - A single write request issued by the foreground user process
 - Used as a victim for replacement
- **LRU Tail Flush:**
 - Asynchronous write requests issued by the background process
 - For cold page eviction
- **Flush List Flush (i.e., fuzzy checkpoint):**
 - Asynchronous write requests issued by the background process
 - For database recovery upon failure



Disk write type 3:

Flush List Flush




What is the Role of REDO Log?

- Each time data is changed, the page containing the data is modified in memory
- The page is denoted as *dirty*
- In case of an unexpected failure, we cannot lose all the changes
 - *But the data in memory is gone!*
- This is why diff data of the page is also written (and by default flushed to disk) to **REDO logs**
 - The data in REDO logs will be read only in case of recovery
 - During recovery, the modified pages will be reconstructed with the modified data



InnoDB Redo Log

- What people mean by InnoDB log (e.g., in my.cnf)
 - `Innodb_log_file_size=1G`
 - `Innodb_log_files_in_group=2`
- Two or more redo logs pre-allocated and used in a circular fashion
 - `ib_logfile1` `ib_logfile2` files in `/path/to/test-data`
- Information necessary to redo (or re-apply) changes to data stored in InnoDB
- Used to reconstruct changes if necessary (i.e., crash recovery)



Log Sequence Number (LSN)

- A 64-bit unsigned integer representing a point in time in the redo log system
 - The total number of log data that have been generated since log initialization
- An ever-increasing number similar to the number of bytes logged
- You can leverage the **LSN** to locate a log data in a log file
 - LSN: *lsn_t* type variable



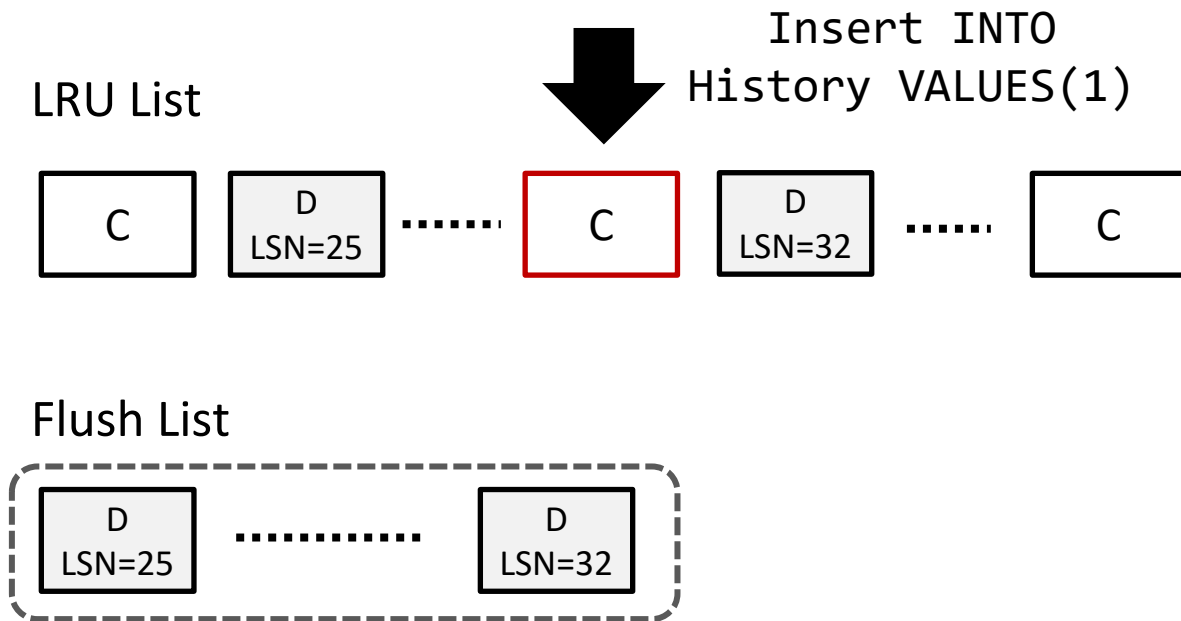
Checkpoint

- The LSN value of the latest changes written to the data files
- Checkpoint means that changes made prior to the checkpoint LSN have been flushed
- Once a checkpoint has been completed, redo logs prior to the checkpoint are no longer needed

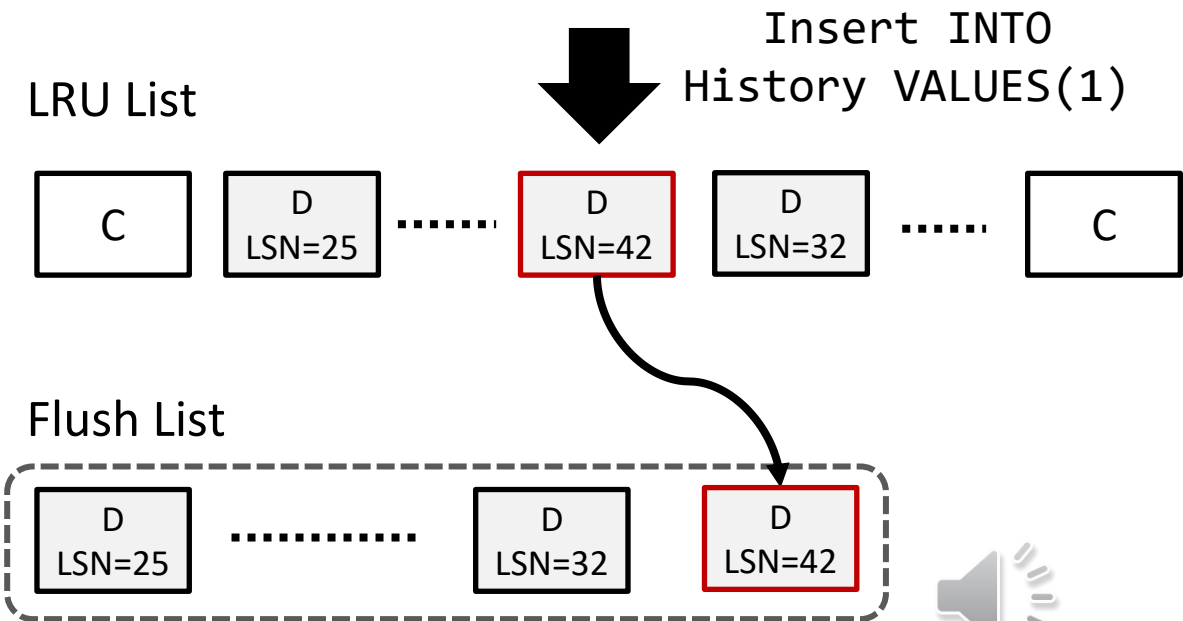


Flush List

- **List of dirty pages** that have been changed in the memory but not written to disk
 - Strictly ordered on the *oldest modification LSN*




Before page modification



After page modification

Flush List Flush

- **Flushing to advance the oldest modified LSN**
 - Also known as fuzzy checkpoint



Types of Checkpoints

- *Sharp checkpoint (at shutdown):*
 - Flushes all modified pages for committed transactions to disk
 - Writes down the LSN of the most recently committed transaction
 - All flushed pages is consistent as of a single point in time (the checkpoint LSN) → “sharp”
- *Fuzzy checkpoint (at normal time):*
 - Flushes pages as time passes (i.e., **flush list flushing**)
 - Flushed pages might not all be consistent with each other as of a single point in time → “fuzzy”



Page Cleaner Thread(s)

- buf/buf0flu.cc: buf_flush_page_cleaner_coordinator()

```
switch (recv_sys->flush_type) {  
case BUF_FLUSH_LRU:  
    /* Flush pages from end of LRU if required */  
    pc_request(0, LSN_MAX);  
    while (pc_flush_slot() > 0) {}  
    pc_wait_finished(&n_flushed_lru, &n_flushed_list);  
    break;  
  
case BUF_FLUSH_LIST:  
    /* Flush all pages */  
    do {  
        pc_request(ULINT_MAX, LSN_MAX);  
        while (pc_flush_slot() > 0) {}  
    } while (!pc_wait_finished(&n_flushed_lru,  
                               &n_flushed_list));  
    break;
```

LRU tail flush

Flush list flush



Flush List Flush

- buf/buf0flu.cc: buf_do_flush_list_batch()

```
static
uint
buf_do_flush_list_batch(
    buf_pool_t*    buf_pool,
    uint           min_n,
    lsn_t          lsn_limit)
{
    uint          count = 0;
    uint          scanned = 0;

    ut_ad(buf_pool_mutex_own(buf_pool));

    /* Start from the end of the list looking for a suitable
       block to be flushed. */
    buf_flush_list_mutex_enter(buf_pool);
    uint len = UT_LIST_GET_LEN(buf_pool->flush_list);
```

Flush the dirty pages inside the
flush list until lsn_limit

Acquire flush list mutex



Flush List Flush

- buf/buf0flu.cc: buf_do_flush_list_batch()

```
for (buf_page_t* bpage = UT_LIST_GET_LAST(buf_pool->flush_list);
    count < min_n && bpage != NULL && len > 0
    && bpage->oldest_modification < lsn_limit;
    bpage = buf_pool->flush_hp.get(),
    ++scanned) {

    buf_page_t* prev;

    ut_a(bpage->oldest_modification > 0);
    ut_ad(bpage->in_flush_list);

    prev = UT_LIST_GET_PREV(list, bpage);
    buf_pool->flush_hp.set(prev);
    buf_flush_list_mutex_exit(buf_pool);
```

Iterate for loop starting from the last bpage of the flush list until bpage lsn(oldest modification) is smaller than lsn_limit

oldest_modification ==0 → clean
oldest_modification >0 → dirty



Flush List Flush

- buf/buf0flu.cc: buf_do_flush_list_batch()

```
#ifdef UNIV_DEBUG
    bool flushed =
#endif /* UNIV_DEBUG */
    buf_flush_page_and_try_neighbors(
        bpage, BUF_FLUSH_LIST, min_n, &count);

    buf_flush_list_mutex_enter(buf_pool);

    ut_ad(flushed || buf_pool->flush_hp.is_hp(prev));

    --len;
}
```

For all flushable pages within the flush area, flush them asynchronously

Decrement the length of a flush list



Flush List Flush: Complete I/O

- After all the work for flushing is complete, the following function is called last to complete I/O
- buf/buf0buf.cc: buf_page_io_complete()

```
bool
buf_page_io_complete(
/*=====*/
    buf_page_t* bpage, /*!< in: pointer to the block in question */
    bool        evict) /*!< in: whether or not to evict the page
                        from LRU list. */
{
    enum buf_io_fix io_type;
    buf_pool_t* buf_pool = buf_pool_from_bpage(bpage);
    const ibool uncompressed = (buf_page_get_state(bpage)
                                == BUF_BLOCK_FILE_PAGE);

    ut_a(buf_page_in_file(bpage));
```

```
/* We do not need protect io_fix here by mutex to read
it because this is the only function where we can change the value
from BUF_IO_READ or BUF_IO_WRITE to some other value, and our code
ensures that this is the only thread that handles the i/o for this
block. */
```

```
io_type = buf_page_get_io_fix(bpage);
ut_ad(io_type == BUF_IO_READ || io_type == BUF_IO_WRITE);
```

io_type is either read or write



Flush List Flush: Complete I/O

- buf/buf0buf.cc: buf_page_io_complete()

case BUF_IO_WRITE:

/* Write means a io_type is write the completion routine in the flush system */

buf_flush_write_complete(bpage);

```
if (uncompressed) {  
    rw_lock_sx_unlock_gen(&((buf_block_t*) bpage)->lock,  
                          BUF_IO_WRITE);  
}
```

buf_pool->stat.n_pages_written++;

Do not free the flushed page!
Keep it in LRU list as a clean page.

/* We decide whether or not to evict the page from the LRU list based on the flush_type.

* BUF_FLUSH_LIST: don't evict

* BUF_FLUSH_LRU: always evict

* BUF_FLUSH_SINGLE_PAGE: eviction preference is passed by the caller explicitly. */

```
if (buf_page_get_flush_type(bpage) == BUF_FLUSH_LRU) {  
    evict = true;  
}
```

If LRU list flush, then free the page and return it to the free list

```
if (evict) {  
    mutex_exit(buf_page_get_mutex(bpage));  
    buf_LRU_free_page(bpage, true);  
} else {  
    mutex_exit(buf_page_get_mutex(bpage));  
}
```

break;



InnoDB Configuration Options (my.cnf)

- You can readjust innodb parameters according to your environment
 - E.g., `innodb_io_capacity`, `innodb_buffer_pool_size`...
- Optimizing InnoDB by just modifying these parameters can lead to significant performance improvement
- Refer to this link for details about innodb parameters:
<https://dev.mysql.com/doc/refman/5.7/en/innodb-parameters.html>



PA 1

- **PA1: Improve TpmC by changing IO related InnoDB parameters**
- For the assignment, your task is to change/add the I/O related InnoDB parameters in my.cnf to improve TpmC
- By looking at MySQL source code and MySQL document, investigate which parameter affects I/O operation process in InnoDB and how you can improve the TpmC by readjusting it/them.
- You can add/modify multiple InnoDB options.



PA 1

- After achieving performance gain, present an experiment result before and after changing `my.cnf`
- Then, elaborate the reason why it leads to performance improvement based on MySQL source code and document.
- The PA will be graded based on the following criteria: **TpmC improvement (40%)**, **Result Analysis (60%)**.
- You will get zero marks on both criteria if you forge your results.
- Refer to week5 for the PA 1 experiment guide
<https://github.com/LeeBohyun/SWE3033-S2023>



References

- [1] MySQL document: <https://dev.mysql.com>
- [2] Mijin An, MySQL Buffer Management, <https://www.slideshare.net/meeeejin/mysql-buffer-management>
- [3] An et.al., “Avoiding Read Stalls on Flash Storage”, SIGMOD2022, <https://dl.acm.org/doi/pdf/10.1145/3514221.3526126>

