

RocksDB: An Introduction

Bo-Hyun Lee
lia323@skku.edu



VLDB
Lab.



RocksDB Introduction

- A persistent **key-value** store for fast storage
 - Exploiting the full potential of high read/write rates offered by *flash or RAM*
- A **log structured** database engine, written entirely in C++
- **Open-source**, based on LevelDB 1.5
- Focusing on **performance** and **scalability**
 - Optimized for Server workloads



Three Basic Components of RocksDB

1. Memtable

- In-memory data structure
- A buffer, temporarily host the incoming writes

2. Logfile

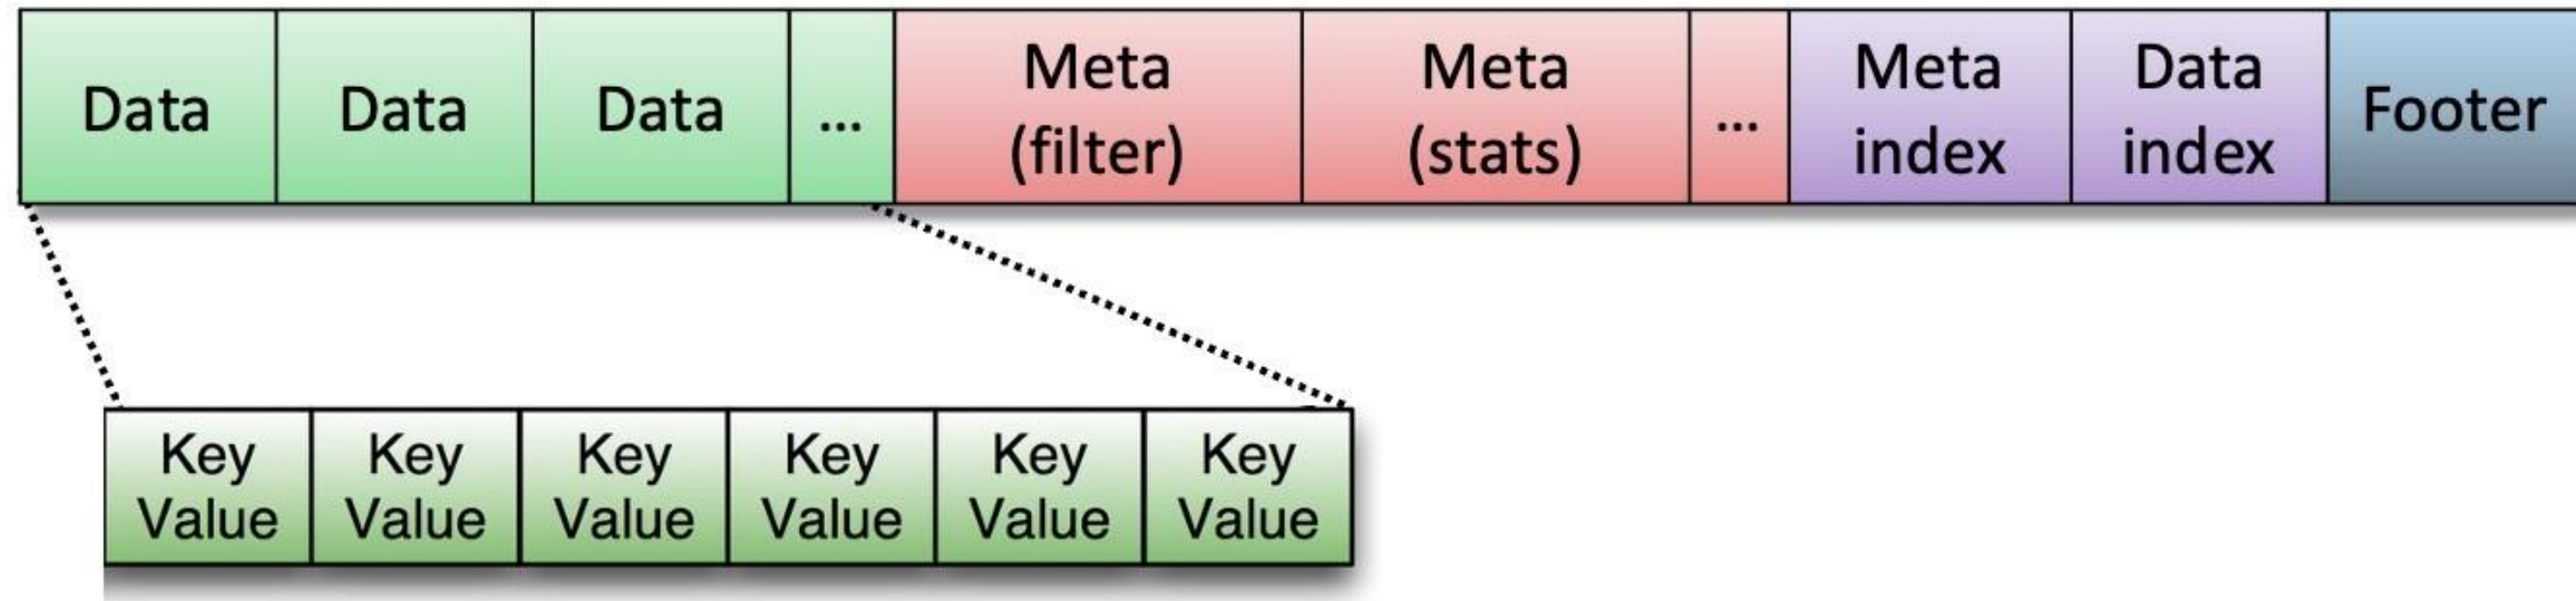
- Sequentially written file on storage

3. SSTable (=SST file)

- Sorted sequence table on storage → to facilitate easy lookup of keys
- A file which contains a set of arbitrary, sorted **key-value pairs** inside
- Organized in levels
- Immutable in its lifetime



Block Based Table



- The **default SST table format** in RocksDB
- The sequence of key/value pairs in the file are sorted in key order
 - They are partitioned into a sequence of data blocks

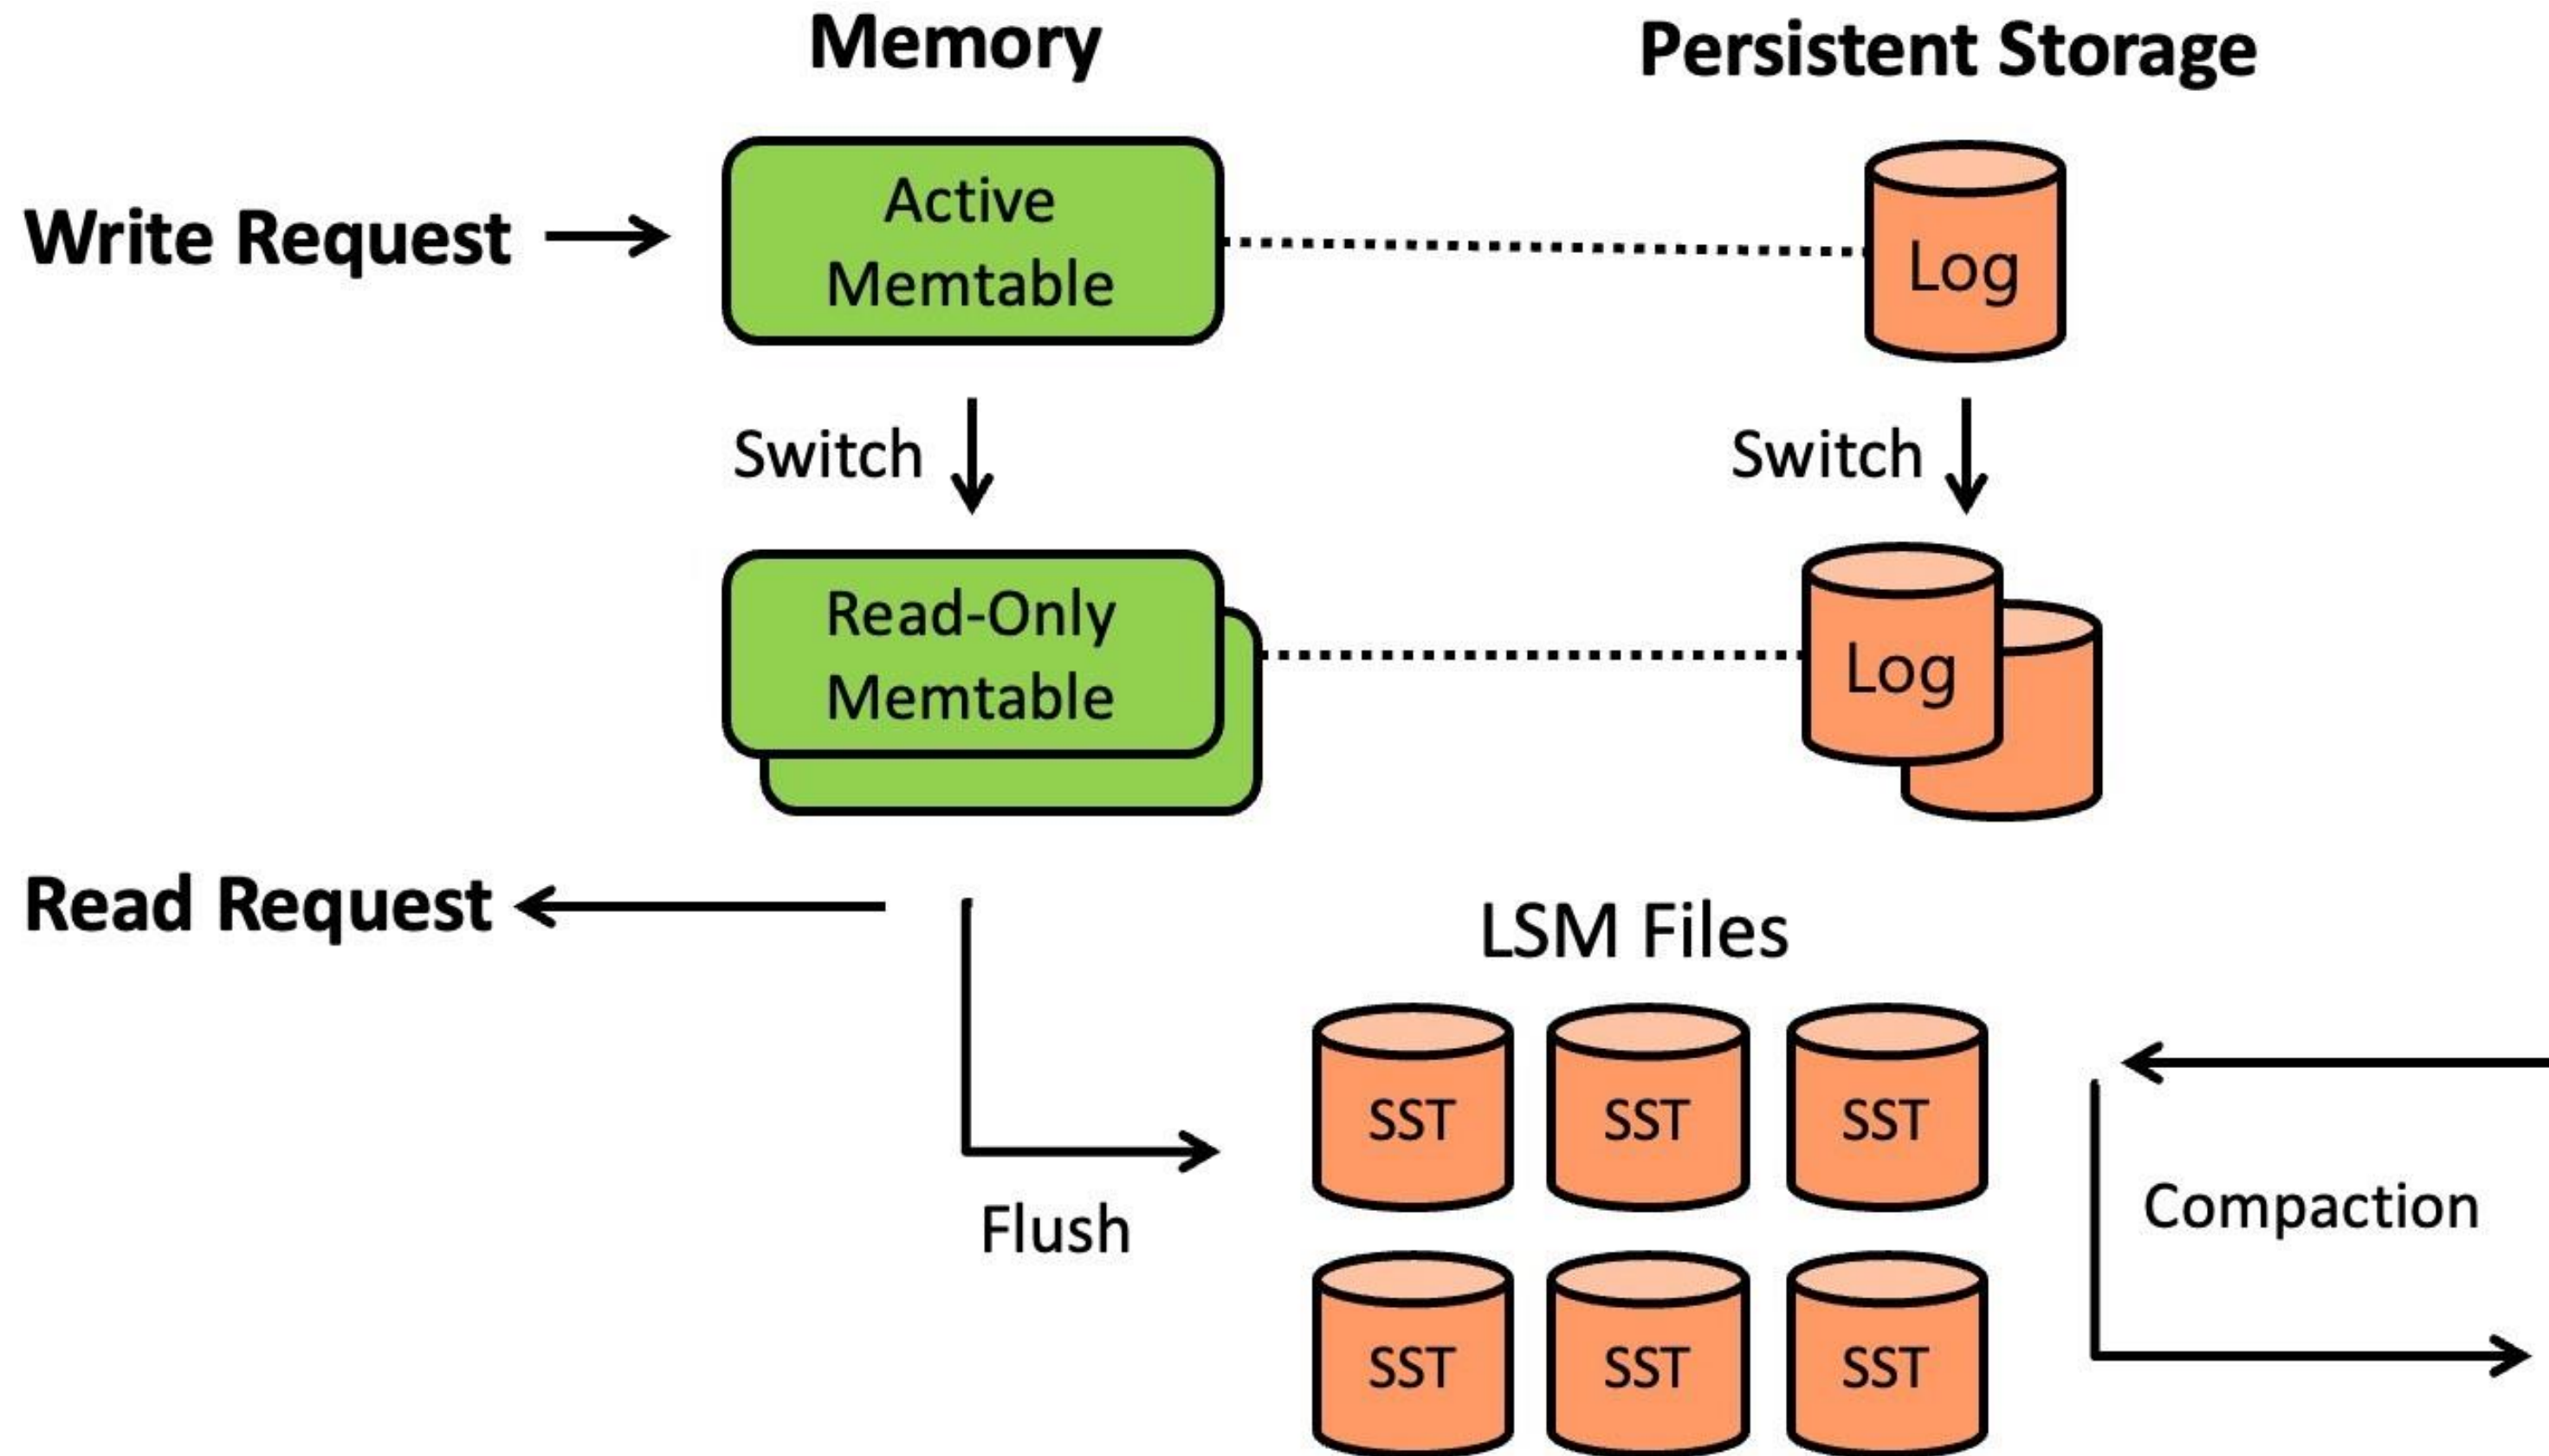


Memtable and SST File

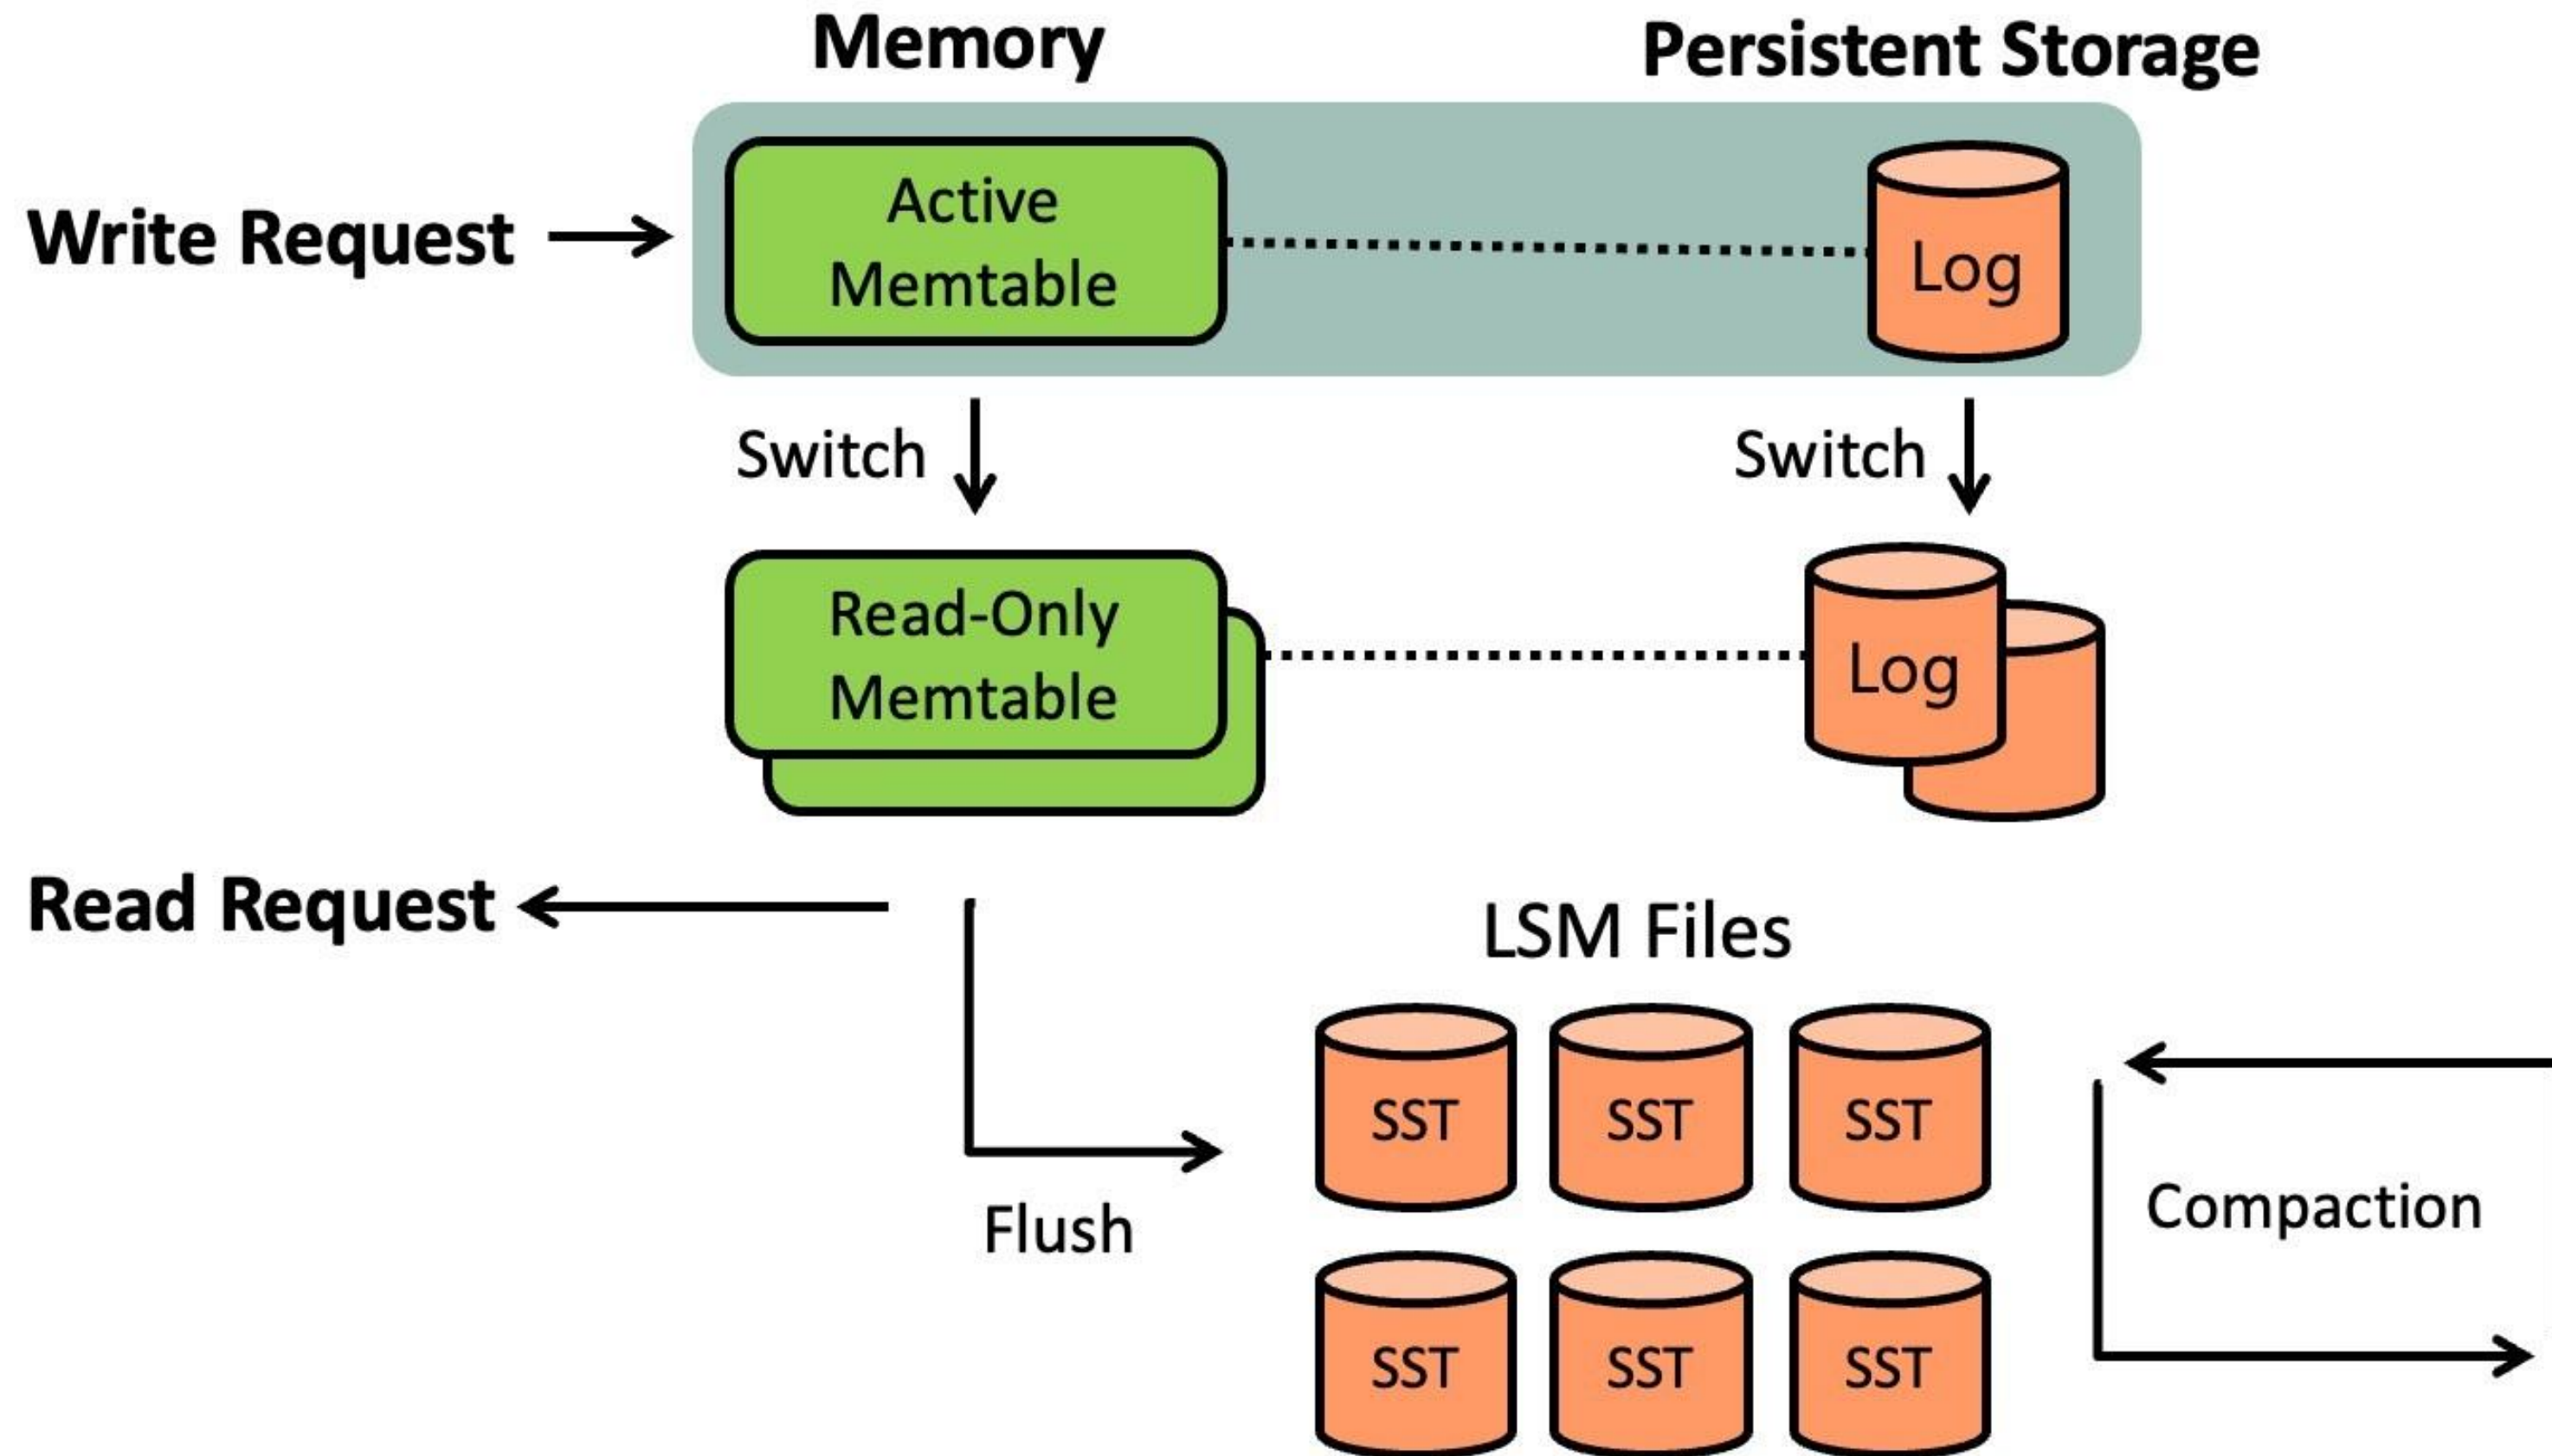
- On-disk SSTable indexes are always loaded into memory
- All writes go directly to the memtable index
- **Reads** check the memtable first → then, the SSTable indexes
- Periodically, the memtable is flushed to disk as an SSTable
- Periodically, on-disk SSTables are merged
 - Update/delete records will overwrite/remove the older data



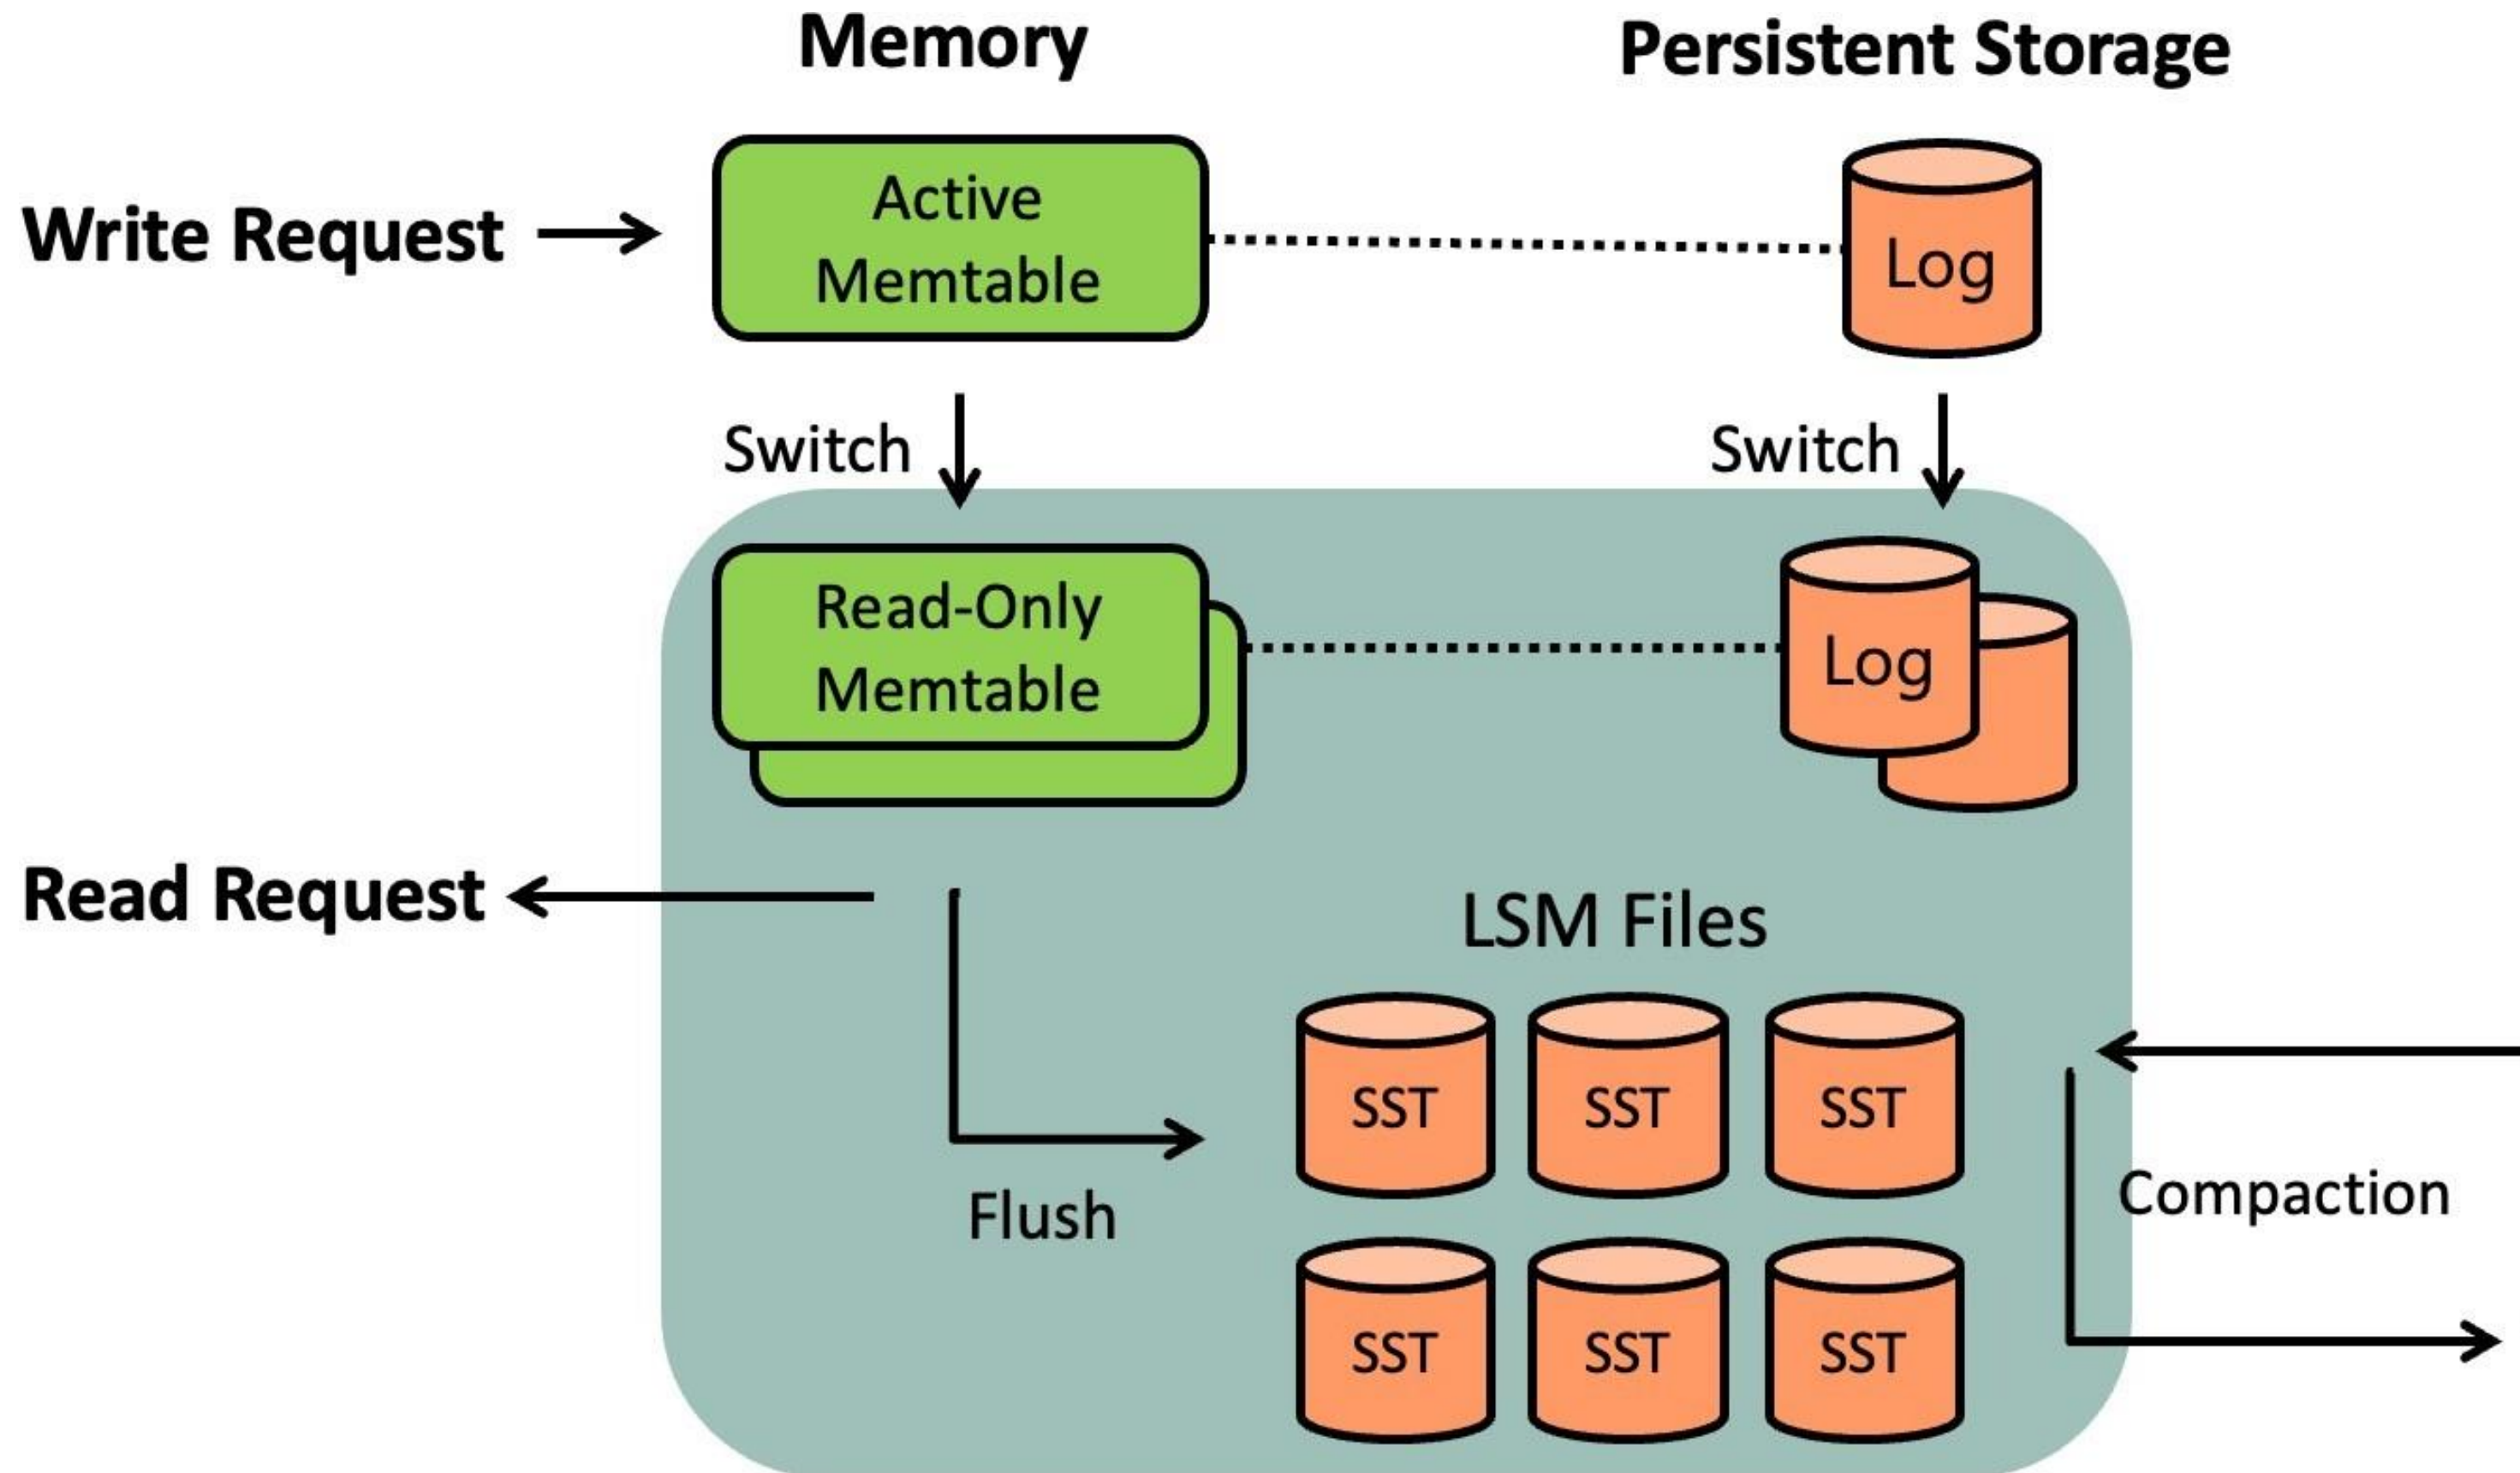
RocksDB Architecture



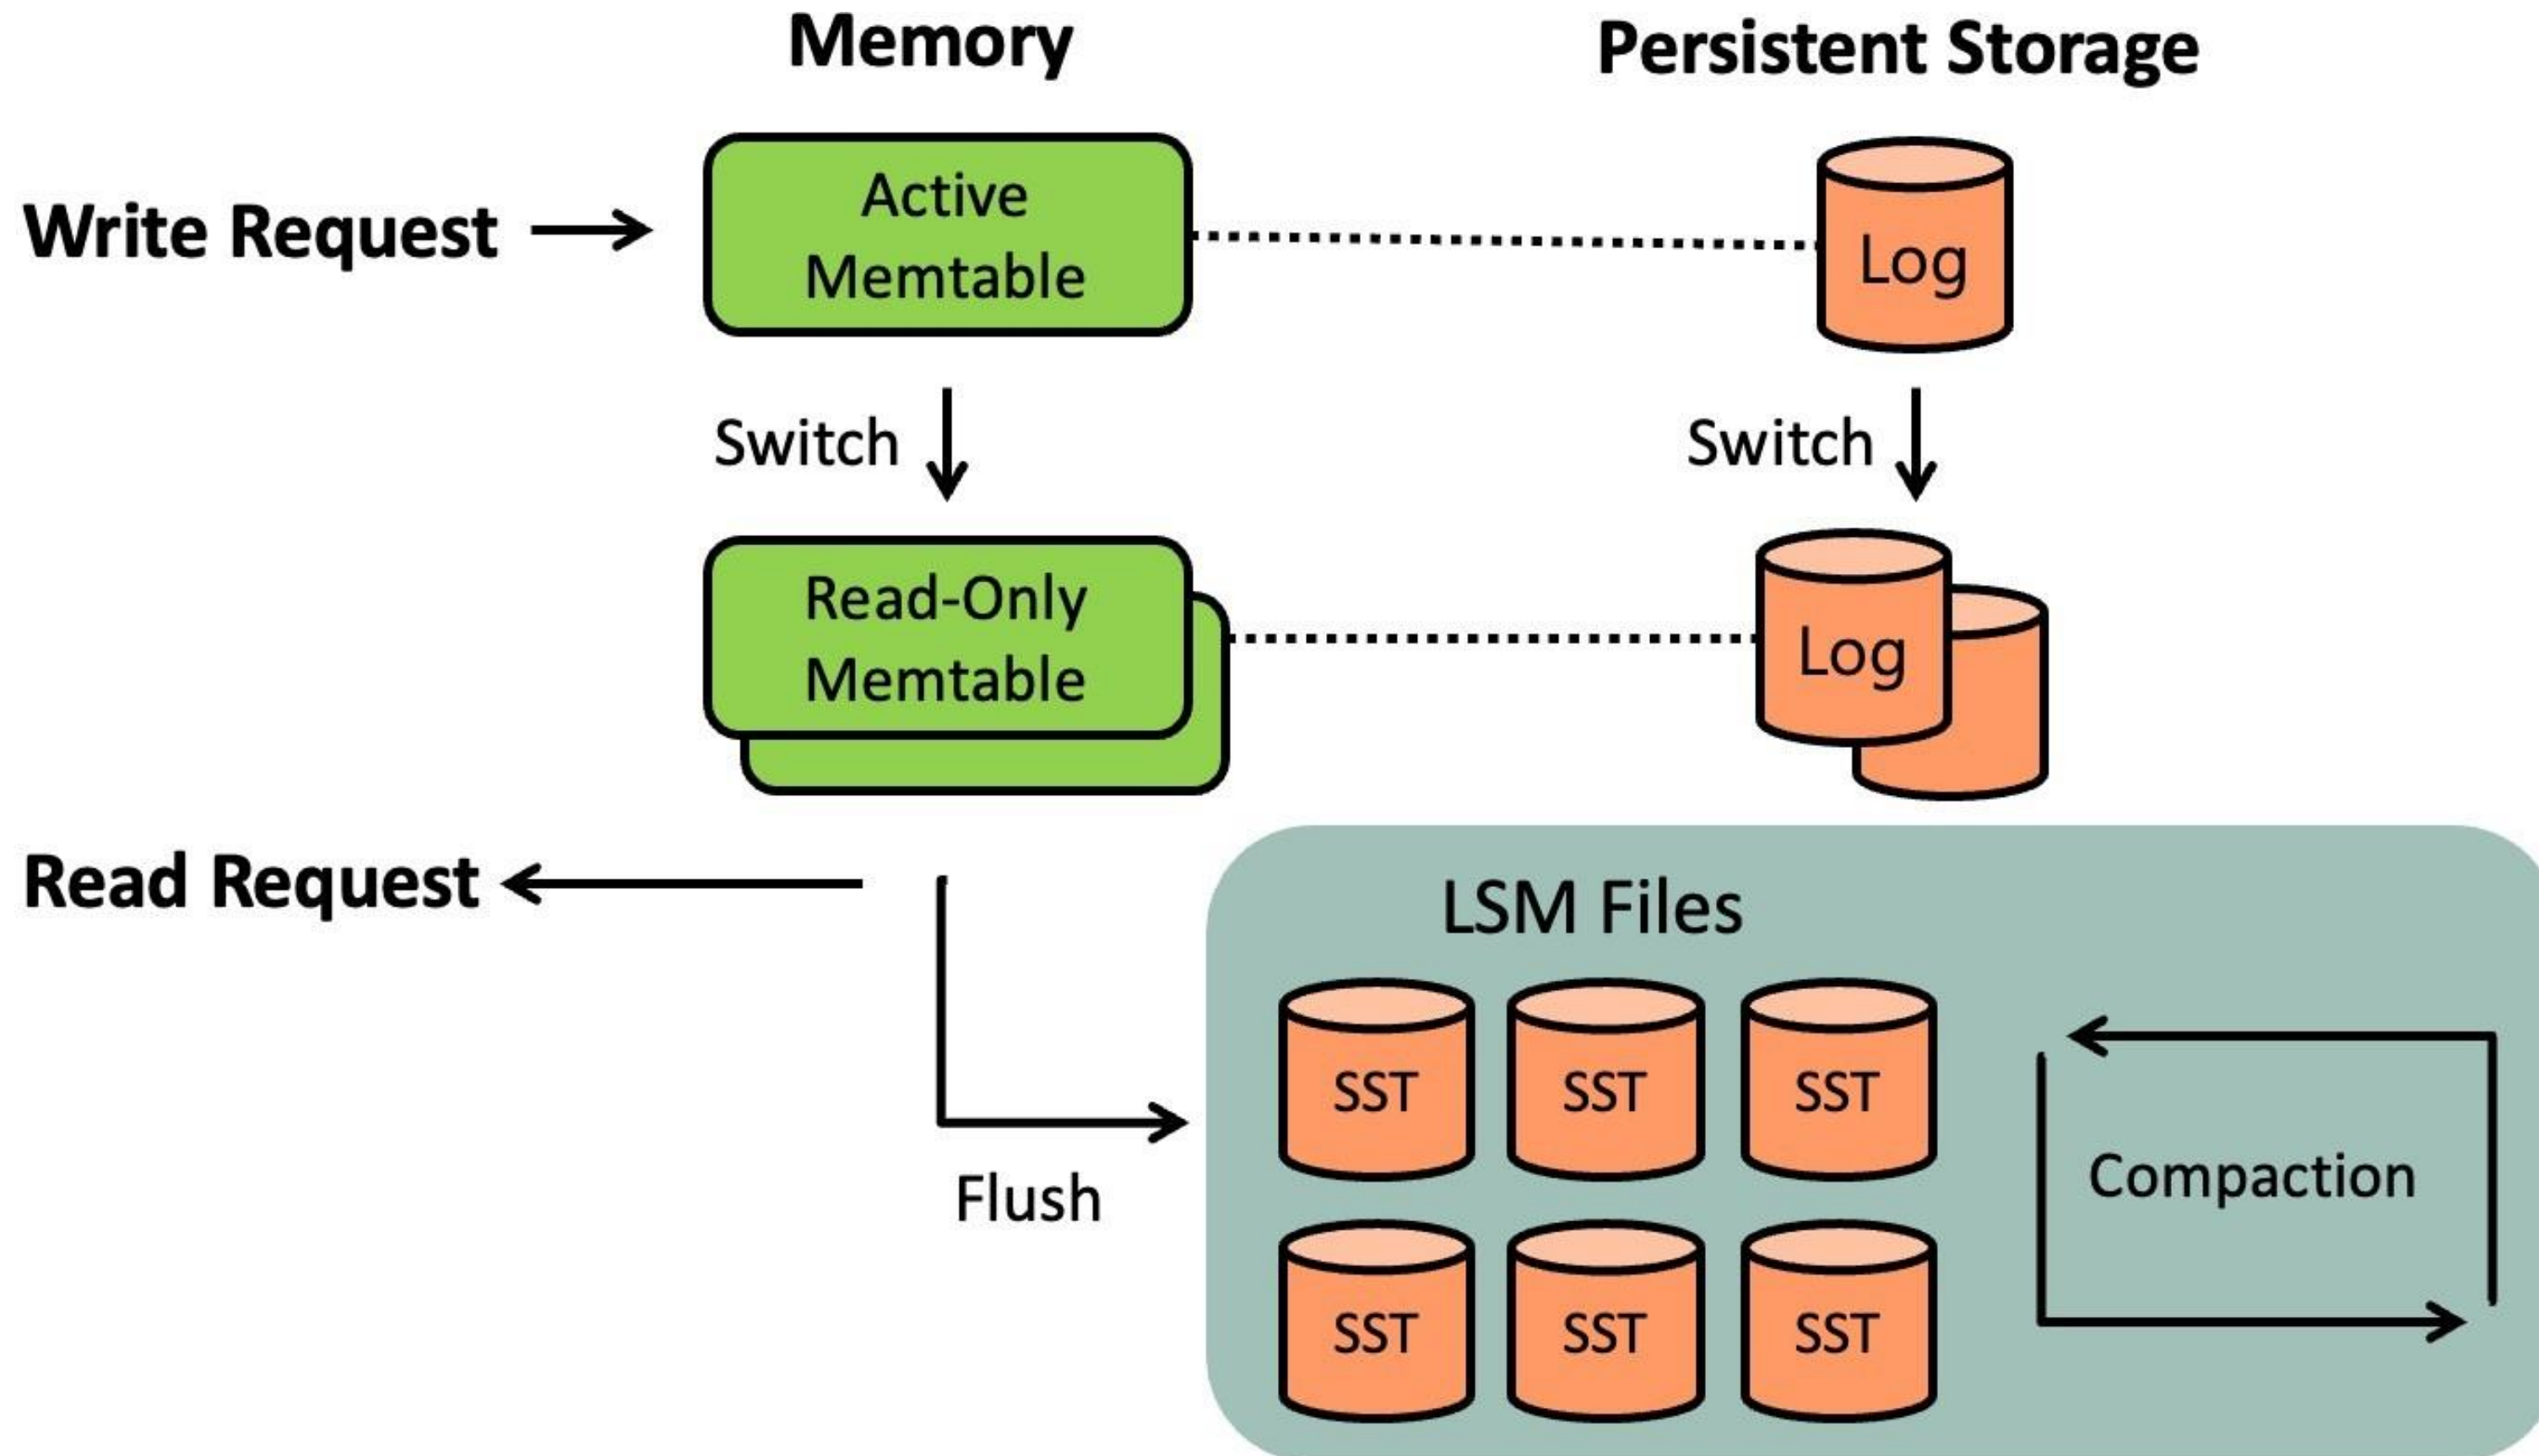
RocksDB Architecture




RocksDB Architecture



RocksDB Architecture




B+ Tree



- *B+Tree* directly overwrites old records to store new updates → **In-place update**
- This structure is **read-optimized** since only the most recent version of each record is stored
- However, this design sacrifices write performance as **updates incur random I/Os**



Log-Structured Merge Tree



- LSM-Tree → **Out-of-place update**
 - **N-level merge trees**
 - Splitting a **logical tree** into **several physical pieces**
 - So that the *most-recently-updated* portion of data is in a tree in **memory**
 - **Transform random writes into sequential writes** using *logfile* and in-memory store (*memtable*)



Install RocksDB and Run DB_Bench

- This week, you will install RocksDB and run DB_Bench on your own system
- Refer to github link for week 6 experiment guide
- Please send an email to lia323@skku.edu if you have any question about this week's lecture or experiment



Reference

- Facebook, “RocksDB”, <http://rocksdb.org>
- O'Neil, Patrick E., Edward Y. C. Cheng, Dieter Gawlick and Elizabeth J. O'Neil. “The log-structured merge-tree (LSM-tree).” *Acta Informatica* 33 (2009): 351-385.
- Chen Luo, Michael J. Carey, “LSM-based Storage Techniques: A Survey”. *VLDB Journal*, 2019
- Dhruba Borthakur and Hadbo Xu, “The Story of RocksDB”, SlideShare, <https://www.slideshare.net/HiveData/tech-talk-rocksdb-slides-by-dhruba-borthakur-haobo-xu-of-facebook>
- Mijin An, “RocksDB detail”, SlideShare, <https://www.slideshare.net/meeeejin/rocksdb-detail>
- Facebook, “RocksDB Wiki”, GitHub, <https://github.com/facebook/rocksdb/wiki>