

# Concurrent Programming

Prof. Jin-Soo Kim(jinsookim@skku.edu)  
TA – Taekyun Roh(tkroh0198@skku.edu)  
Sewan Ha(hsewan2495@gmail.com)  
Computer Systems Laboratory  
Sungkyunkwan University  
<http://csl.skku.edu>



# Announcement

- **PA1 is graded**

- Uploaded on i-campus
- If any questions, please contact **tkroh0198@skku.edu** or ask **Taekyun Roh** directly.
- Room # : 85545(Internet Management Technology Lab)  
Time : 2018/11/12 (Monday) 16:00~18:00

- **PA2 will be uploaded until this Sunday**

# Echo Server Revisited


```
int main (int argc, char *argv[]) {
    ...
    listenfd = socket(AF_INET, SOCK_STREAM, 0);

    bzero((char *)&saddr, sizeof(saddr));
    saddr.sin_family = AF_INET;
    saddr.sin_addr.s_addr = htonl(INADDR_ANY);
    saddr.sin_port = htons(port);
    bind(listenfd, (struct sockaddr *)&saddr, sizeof(saddr));

    listen(listenfd, 5);
    while (1) {
        connfd = accept(listenfd, (struct sockaddr *)&caddr, &clen);
        while ((n = read(connfd, buf, MAXLINE)) > 0) {
            printf ("got %d bytes from client.\n", n);
            write(connfd, buf, n);
        }
        close(connfd);
    }
}
```


# Iterative Servers (1)

- One request at a time



# Iterative Servers (2)

## ■ Fundamental flaw



## ■ Solution: use concurrent servers instead

- Use multiple concurrent flows to serve multiple clients at the same time.

# Creating Concurrent Flows

## ■ Processes

- Kernel automatically interleaves multiple logical flows.
- Each flow has its own private address space.

## ■ Threads

- Kernel automatically interleaves multiple logical flows.
- Each flow shares the same address space.
- Hybrid of processes and I/O multiplexing

## ■ I/O multiplexing with `select()`

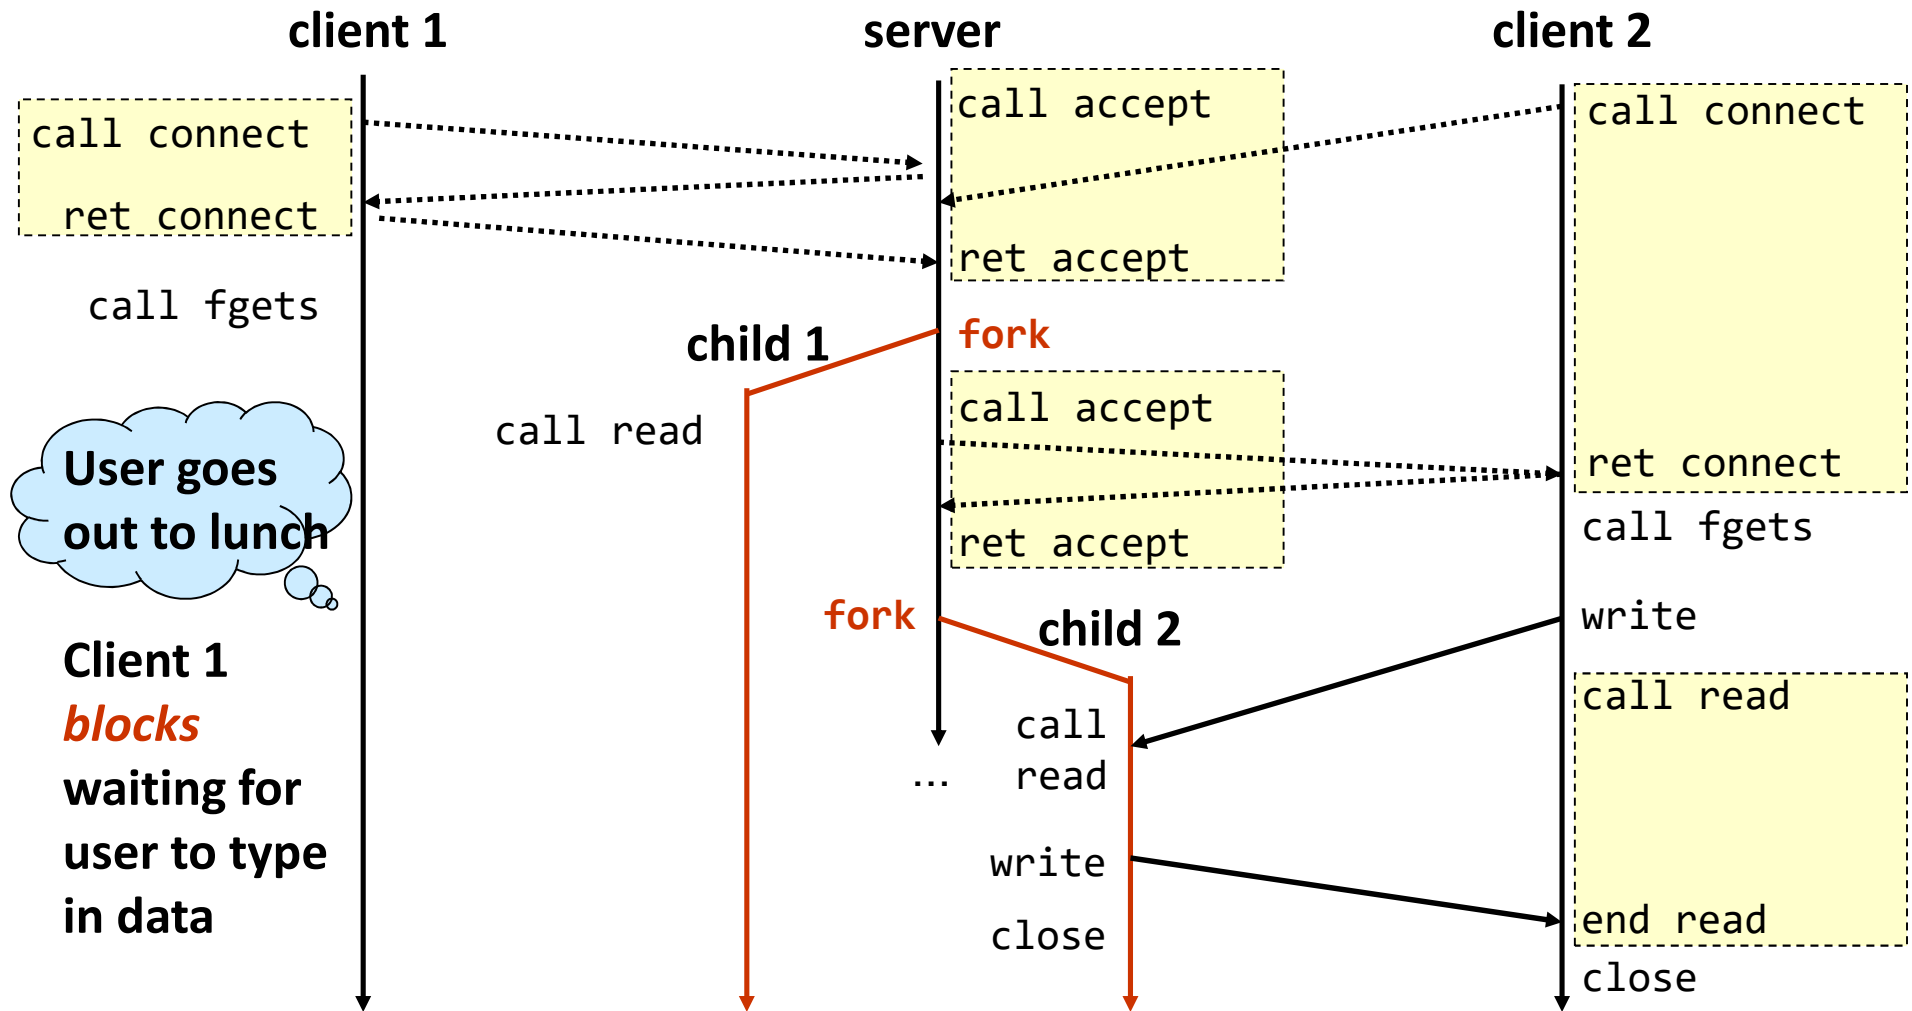
- User manually interleaves multiple logical flows
- Each flow shares the same address space
- Popular for high-performance server designs.



**Concurrent Programming**

**Process-based**

# Process-based Servers





# Implementation Issues

- **Servers should restart `accept()` if it is interrupted by a transfer of control to the `SIGCHLD` handler**
  - Not necessary for systems with POSIX signal handling.
  - Required for portability on some older Unix systems.
- **Server must reap zombie children**
  - to avoid fatal memory leak
- **Server must close its copy of `connfd`.**
  - Kernel keeps reference for each socket.
  - After **`fork()`**, **`refcnt(connfd) = 2`**
  - Connection will not be closed until **`refcnt(connfd) = 0`**

# Process-based Designs

## ■ Pros

- Handles multiple connections concurrently.
- Clean sharing model.
  - Descriptors (no), file tables (yes), global variables (no)
- Simple and straightforward.

## ■ Cons

- Additional overhead for process control.
  - Process creation and termination
  - Process switching
- Nontrivial to share data between processes.
  - Requires IPC (InterProcess Communication) mechanisms: FIFO's, System V shared memory and semaphores

# Echo Server

## ■ Iterative version

```
int main (int argc, char *argv[])
{
    . . .

    while (1) {
        connfd = accept (listenfd, (struct sockaddr *)&caddr,
                        &caddrlen));

        while ((n = read(connfd, buf, MAXLINE)) > 0) {
            printf ("got %d bytes from client.\n", n);
            write(connfd, buf, n);
        }

        close(connfd);
    }
}
```

# Echo Server: Process-based

```
int main (int argc, char *argv[])
{
    . . .
    signal (SIGCHLD, handler);

    while (1) {
        connfd = accept (listenfd, (struct sockaddr *)&caddr,
                        &caddrlen);
        if (fork() == 0) {
            close(listenfd);
            while ((n = read(connfd, buf, MAXLINE)) > 0) {
                printf ("got %d bytes from client.\n", n);
                write(connfd, buf, n);
            }
            close(connfd);
            exit(0);
        }
        close(connfd);
    }
}
```

```
void handler(int sig) {
    pid_t pid;
    int stat;
    while ((pid = waitpid(-1, &stat,
                        WNOHANG)) > 0);

    return;
}
```

# Exercise #1

- With your own code, make calculator server
  - At server side, calculate string(**only one arithmetic operation(+,-,\*,/,%)**) transmitted from client
  - At client side, print the result of calculation
  - Hint)
    - You don't have to change client-side code
    - To change string to integer and vice versa, you can use **sprintf()** and **sscanf()**

## Sample Input)

```
server connected to localhost (127.0.0.1)
connection terminated.
```

## Sample Output)

```
host : 127.0.0.1
2+3
result : 5
3+4
result : 7
5*7
result : 35
20-30
result : -10
80/3
result : 26
9%2
result : 1

invalid calculation
^C
```

# Exercise #2

- **With your own code, make process-based echo server**
  - At the same time, multiple client can be served by echo server
- **There should be no memory leakage**
  - There should be some codes that handle zombie process
  - How about closing files?

- **Print process ID**

**Sample Result)**

```
Server connected to localhost (127.0.0.1)
pid(23810) : got 6 bytes from client.
Server connected to localhost (127.0.0.1)
pid(23812) : got 2 bytes from client.
pid(23812) : got 6 bytes from client.
pid(23812) : got 4 bytes from client.
```