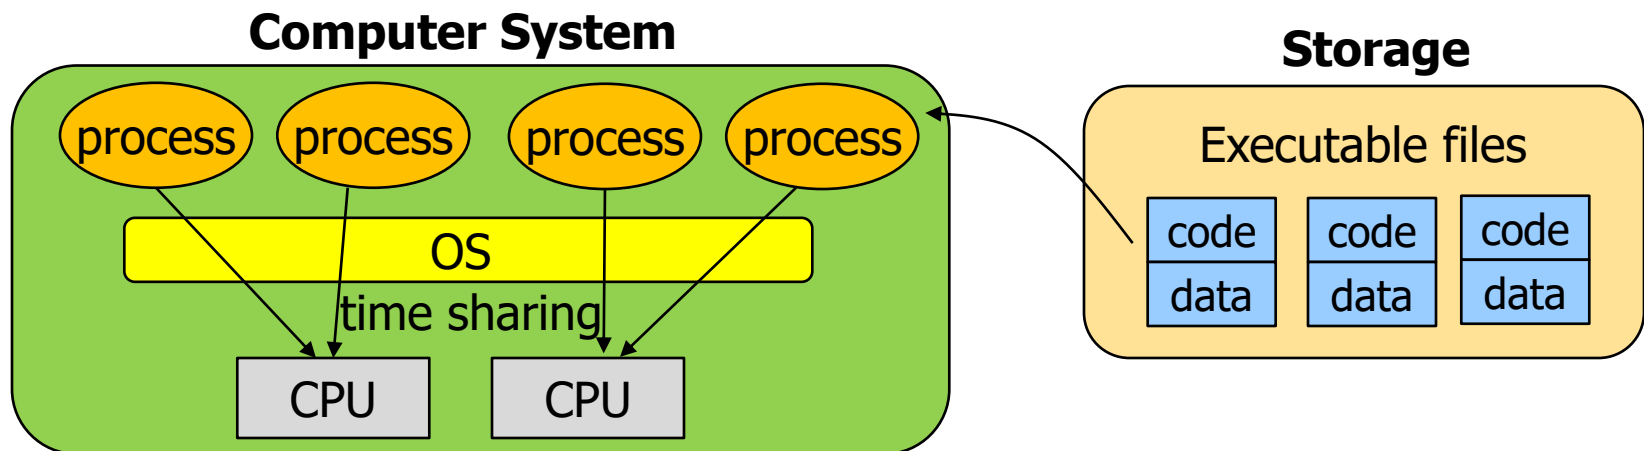


---

# **Chap 4, 5: Process**

# Process Concept

- Process (= running program)
  - An entity that is registered to kernel for execution
  - Kernel manages the processes to improve overall system performance
- A typical system may be seemingly running tens or even hundreds of processes at the same time.
- CPU Virtualization
  - There are only a few physical CPUs available,
  - The OS can promote the illusion that many virtual CPUs exist via **time sharing**
  - By running one process, then stopping it and running another.



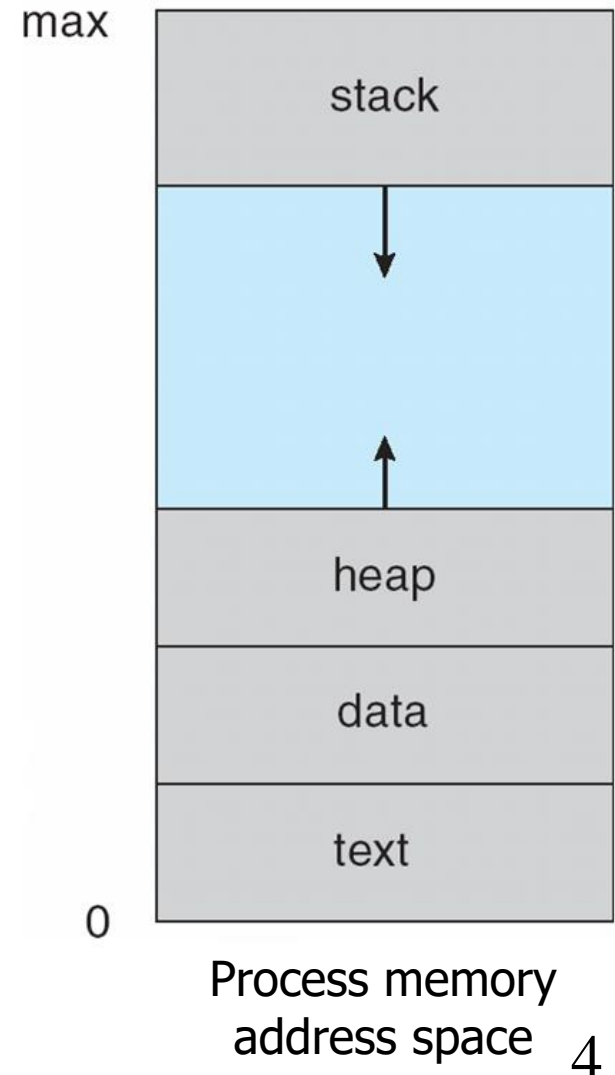
# Process Concept

---

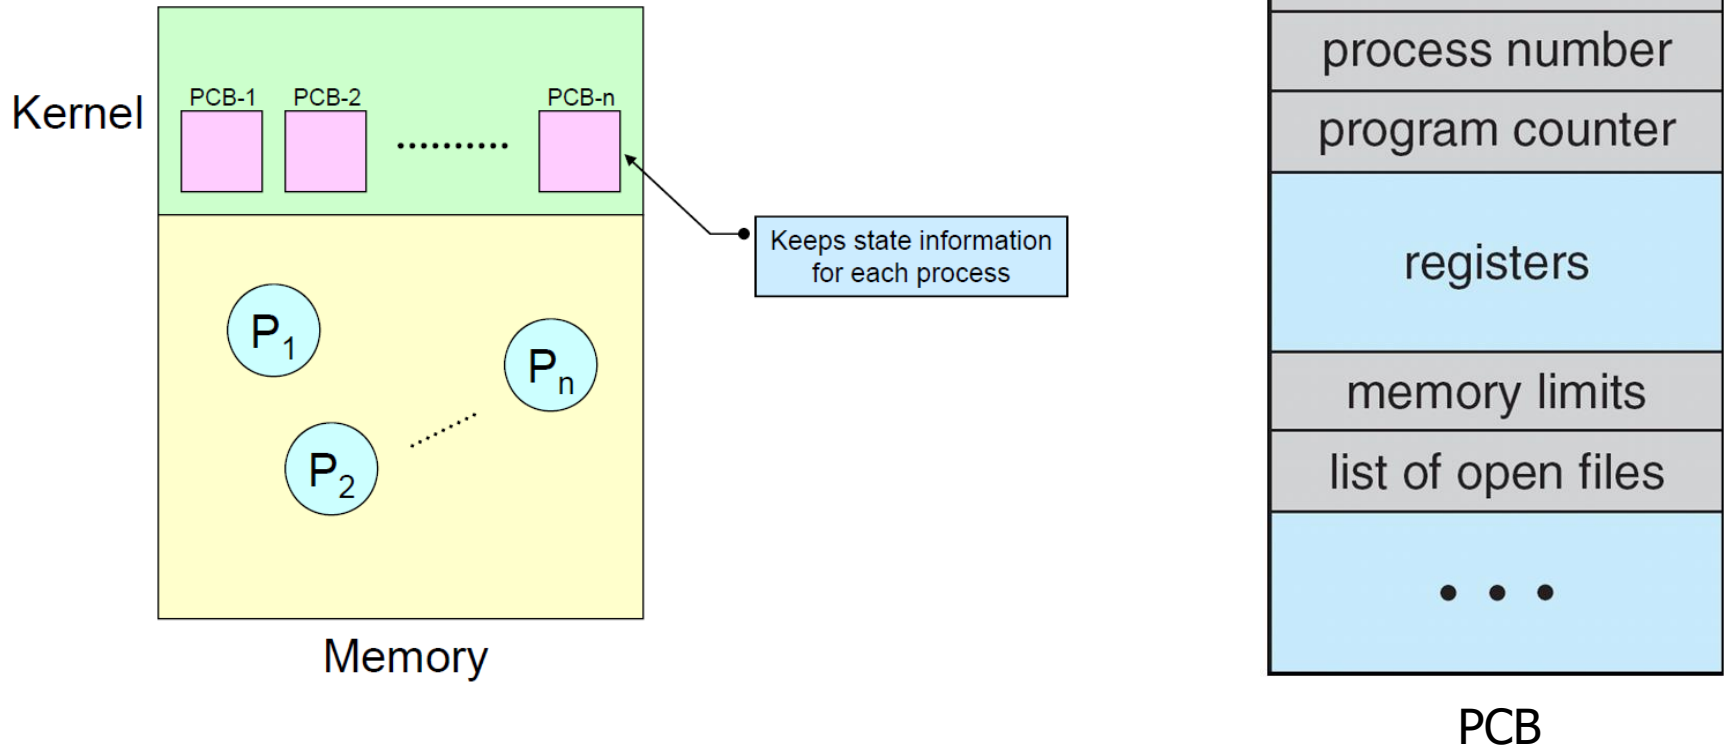
- A process is
  - A program in execution
  - An entity that is registered and being managed by kernel
  - An **active** entity
    - Request/allocate/release system resources during execution
  - An entity that is allocated the PCB
- **PCB** (Process Control Block)
  - Keeps several information about each process that is registered to the kernel
  - Maintained in kernel space

# Machine State of Process

- A process includes
  - Memory
    - Program code → text
    - Global data → data
    - Temporary data → stack
      - Local variables, function parameters, return addresses
    - Heap
      - Memory area that is dynamically allocated during execution
  - Registers
    - Values of the processor registers
    - Include **program counter, stack pointer, frame pointer**
  - I/O information
    - a list of the files the process currently has open



# PCB



- **Information in PCB (task control block)**
  - PID (Process Identification Number)
  - Process state: running, waiting, etc
  - Program counter: location of instruction to next execute
  - CPU registers: contents of all process-centric registers
  - Scheduling information
    - Process priority, scheduling parameters, etc.
  - Memory management information
    - Base/limit registers, page tables, segment tables, etc.
  - I/O status information
    - List of I/O devices allocated, list of open files, etc.
  - Accounting information
    - CPU used, clock time elapsed since start, time limits
  - Context save area
    - Space for saving register context of the process

- **Notes**

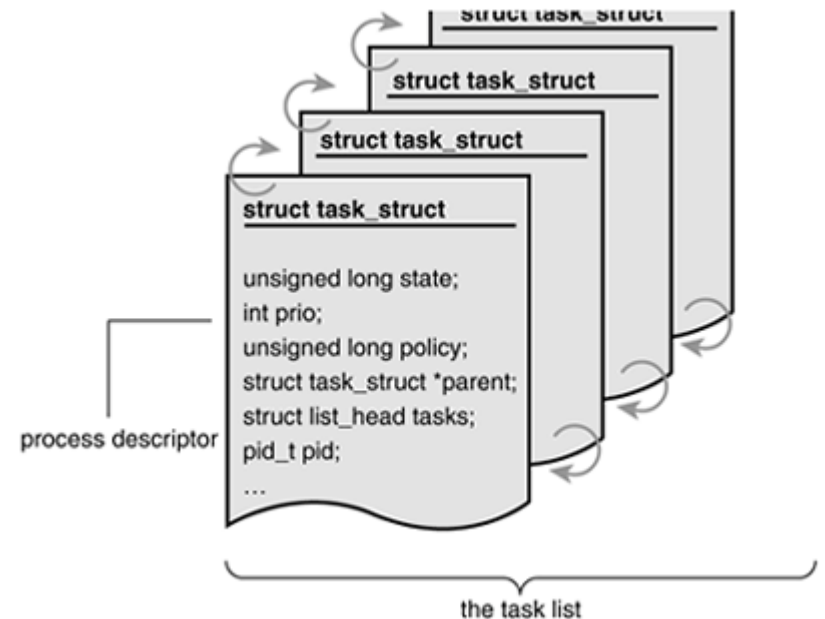
- PCB information is different for each OS
- PCB access speed is important for overall system performance

- **In Unix**

- Process table slot
- U-area

- **In Linux**

- Process descriptor ([task\\_struct](#))



# Process API

---

- **Create**
  - Some method to create new processes.
  - OS is invoked to create a new process to run the program you have indicated.
- **Destroy**
  - an interface to destroy processes forcefully.
- **Wait**
  - Wait for a process to stop running
- **Miscellaneous Control**
  - suspend a process (stop it from running for a while) and then resume it (continue it running).
- **Status**
  - get some status information about a process
  - such as how long it has run for, or what state it is in.

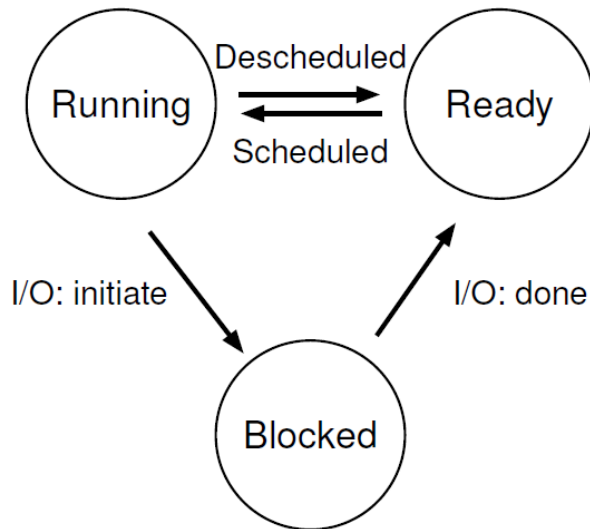


# Process States

---

- **Running**
  - a process is running on a processor, executing instructions.
- **Ready**
  - a process is ready to run but for some reason the OS has chosen not to run it at this given moment.
- **Blocked**
  - a process has performed some kind of operation that makes it not ready to run until some other event takes place.
  - A common example: when a process initiates an I/O request to a disk, it becomes blocked and thus some other process can use the processor.

# Process State Transition

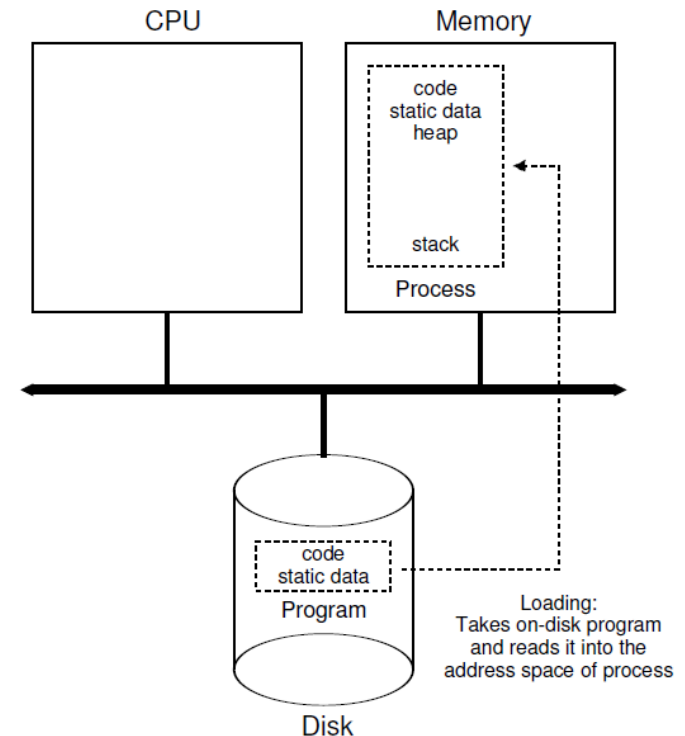


| Time | Process <sub>0</sub> | Process <sub>1</sub> | Notes                         |
|------|----------------------|----------------------|-------------------------------|
| 1    | Running              | Ready                |                               |
| 2    | Running              | Ready                |                               |
| 3    | Running              | Ready                |                               |
| 4    | Running              | Ready                | Process <sub>0</sub> now done |
| 5    | –                    | Running              |                               |
| 6    | –                    | Running              |                               |
| 7    | –                    | Running              |                               |
| 8    | –                    | Running              | Process <sub>1</sub> now done |

| Time | Process <sub>0</sub> | Process <sub>1</sub> | Notes  |
|------|----------------------|----------------------|--|
| 1    | Running              | Ready                |  |
| 2    | Running              | Ready                |  |
| 3    | Running              | Ready                | Process <sub>0</sub> initiates I/O                               |
| 4    | Blocked              | Running              | Process <sub>0</sub> is blocked,<br>so Process <sub>1</sub> runs |
| 5    | Blocked              | Running              |  |
| 6    | Blocked              | Running              |  |
| 7    | Ready                | Running              | I/O done   |
| 8    | Ready                | Running              | Process <sub>1</sub> now done                                    |
| 9    | Running              | –                    |  |
| 10   | Running              | –                    | Process <sub>0</sub> now done                                    |

# Process Creation

1. Programs initially reside on **disk** in some kind of **executable format**
2. **load** a program's code and any static data into memory, into the address space of the process.
  - Eager or lazy loading (paging and swapping)
3. Allocate run-time stack & initialize it with arguments (argc, argv)
4. Allocate heap
5. Do other initialization tasks related to I/O setup
6. jump to the main() routine, transfers control of the CPU to the newly-created process



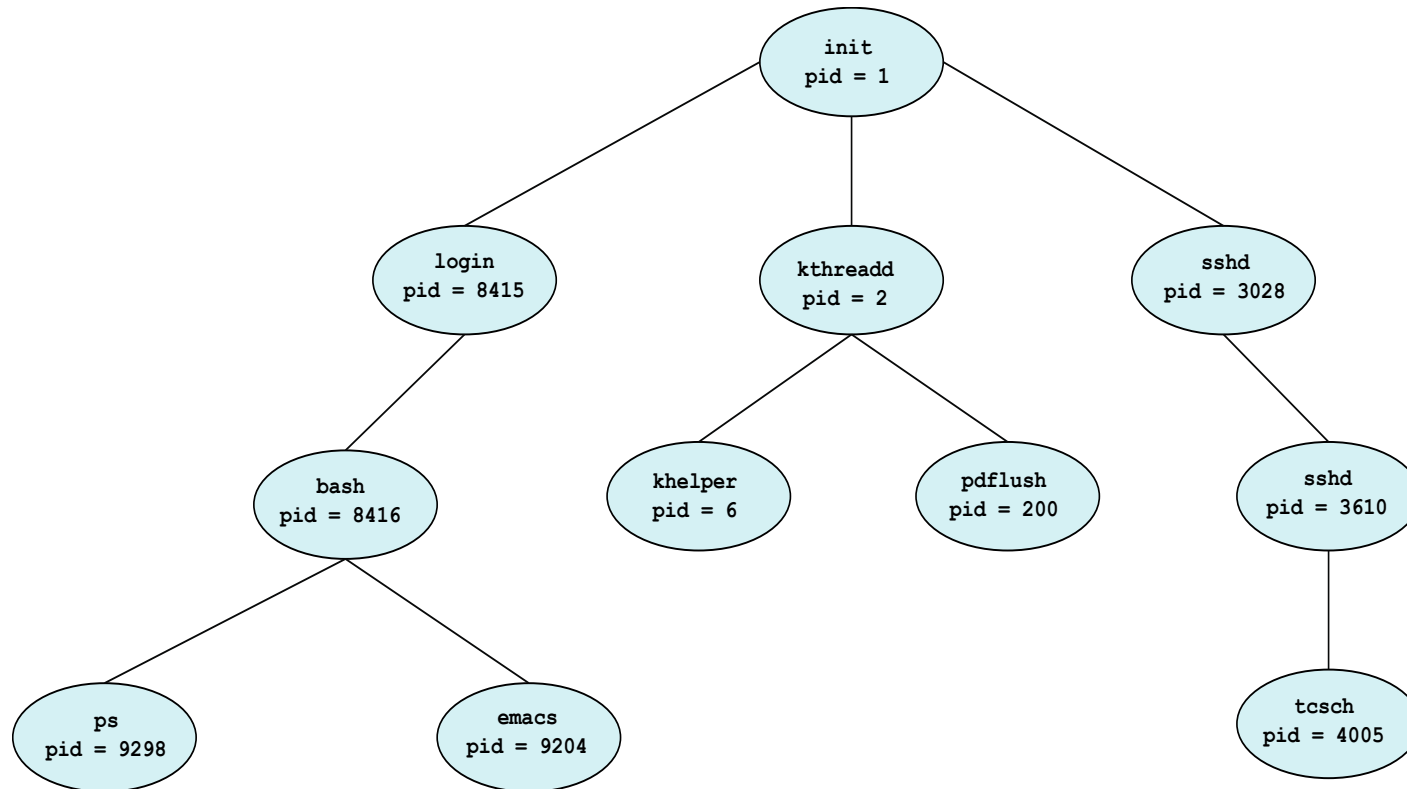
# Process Creation

---

- Parent process create children processes, which, in turn create other processes, forming a tree of processes
- Generally, process identified and managed via a process identifier (pid)
- Resource sharing options
  - Parent and children share all resources
  - Children share subset of parent's resources
  - Parent and child share no resources
- Execution options
  - Parent and children execute concurrently
  - Parent waits until children terminate

# Process Tree in Linux

---



# 'PS' command in Linux

```
1. cs423@localhost:/usr/src/kernels/2.6.40.3-0.fc15.x86_64/include/linux (ssh)
[cs423@localhost linux]$ ps -el
 F S      UID      PID      PPID      C      PRI      NI      ADDR      SZ      WCHAN      TTY      TIME      CMD
 4 S      0         1         0      0      80      0      -      13880  epoll_ ?    00:00:29 systemd
 1 S      0         2         0      0      80      0      -           0 kthrea ?    00:00:00 kthreadd
 1 S      0         3         2      0      80      0      -           0 run_ks ?    00:00:04 ksoftirqd/0
 1 S      0         6         2      0     -40      -      -           0 cpu_st ?    00:00:00 migration/0
 5 S      0         7         2      0     -40      -      -           0 watchd ?    00:00:05 watchdog/0
 1 S      0         8         2      0      60     -20      -           0 rescue ?    00:00:00 cpuset
 1 S      0         9         2      0      60     -20      -           0 rescue ?    00:00:00 khelper
 5 S      0        10         2      0      80      0      -           0 devtmp ?    00:00:00 kdevtmpfs
 1 S      0        11         2      0      60     -20      -           0 rescue ?    00:00:00 netns
 1 S      0        12         2      0      80      0      -           0 bdi_sy ?    00:00:02 sync_supers
 1 S      0        13         2      0      80      0      -           0 bdi_fo ?    00:00:00 bdi-default
 1 S      0        14         2      0      60     -20      -           0 rescue ?    00:00:00 kintegrityd
 1 S      0        15         2      0      60     -20      -           0 rescue ?    00:00:00 kblockd
 1 S      0        16         2      0      60     -20      -           0 rescue ?    00:00:00 ata_sff
 1 S      0        17         2      0      80      0      -           0 hub_th ?    00:00:00 khubd
 1 S      0        18         2      0      60     -20      -           0 rescue ?    00:00:00 md
 1 S      0        20         2      0      80      0      -           0 watchd ?    00:00:01 khungtaskd
 1 S      0        21         2      0      80      0      -           0 kswapd ?    00:00:42 kswapd0
 1 S      0        22         2      0      85      5      -           0 ksm_sc ?    00:00:00 ksmd
 1 S      0        23         2      0      99     19      -           0 khugep ?    00:00:12 khugepaged
 1 S      0        24         2      0      80      0      -           0 fsnoti ?    00:00:00 fsnotify_mark
 1 S      0        25         2      0      60     -20      -           0 rescue ?    00:00:00 crypto
 1 S      0        31         2      0      60     -20      -           0 rescue ?    00:00:00 kthrotld
 1 S      0        34         2      0      80      0      -           0 scsi_e ?    00:01:50 scsi_eh_0
 1 S      0        35         2      0      80      0      -           0 scsi_e ?    00:00:00 scsi_eh_1
 1 S      0        38         2      0      60     -20      -           0 rescue ?    00:00:00 kpsmouse
 1 S      0       221         2      0      60     -20      -           0 rescue ?    00:00:00 mpt_poll_0
 1 S      0       222         2      0      60     -20      -           0 rescue ?    00:00:00 mpt/0
 1 S      0       227         2      0      80      0      -           0 scsi_e ?    00:00:00 scsi_eh_2
 1 S      0       281         2      0      80      0      -           0 kjourn ?    00:00:13 jbd2/sda1-8
 1 S      0       282         2      0      60     -20      -           0 rescue ?    00:00:00 ext4-dio-unwrit
 1 S      0       303         2      0      80      0      -           0 bdi_wr ?    00:00:10 flush-8:0
 1 S      0       310         2      0      80      0      -           0 kaudit ?    00:00:00 kauditd
 4 S      0       319         1      0      80      0      -      5249  epoll_ ?    00:00:00 systemd-logger
 4 S      0       328         1      0      76     -4      -      4998  poll_s ?    00:00:00 udevd
 0 S     38       538         1      0      80      0      -      7633  poll_s ?    00:00:07 ntpd
 4 S      0       539         1      0      80      0      -     31724  poll_s ?    00:00:00 abrttd
 4 S      0       541         1      0      80      0      -      6967  poll_s ?    00:00:00 avahi-daemon
 4 S      0       545         1      0      80      0      -      4198  hrttime ?    00:00:00 atd
 4 S      0       547         1      0      80      0      -     29552  hrttime ?    00:00:14 crond
 4 S      0       548         1      0      80      0      -      4715  hrttime ?    00:00:00 smartd
```

systemd = init

child of kthreadd (kernel thread)

child of systemd (user thread)

# Process Creation API – fork()

- UNIX examples
- **fork()** system call creates new process
  - Child duplicate of parent address space, exact copy of the calling process

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
int main(int argc, char *argv[])
{
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) { // fork failed; exit
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) { // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
    } else { // parent goes down this path (main)
        printf("hello, I am parent of %d (pid:%d)\n", rc, (int) getpid());
    }
    return 0;
}
```

Return value of fork()

- Parent: child PID
- Child: 0


```
prompt> ./p1
hello world (pid:29146)
hello, I am parent of 29147 (pid:29146)
hello, I am child (pid:29147)
prompt>
```

or

```
prompt> ./p1
hello world (pid:29146)
hello, I am child (pid:29147)
hello, I am parent of 29147 (pid:29146)
prompt>
```

# Process Creation API – wait()

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <unistd.h>
4  #include <sys/wait.h>
5
6  int
7  main(int argc, char *argv[])
8  {
9      printf("hello world (pid:%d)\n", (int) getpid());
10     int rc = fork();
11     if (rc < 0) {          // fork failed; exit
12         fprintf(stderr, "fork failed\n");
13         exit(1);
14     } else if (rc == 0) { // child (new process)
15         printf("hello, I am child (pid:%d)\n", (int) getpid());
16     } else {               // parent goes down this path (main)
17         int wc = wait(NULL);
18         printf("hello, I am parent of %d (wc:%d) (pid:%d)\n",
19               rc, wc, (int) getpid());
20     }
21     return 0;
22 }
```



The return value of wait() is the child PID.

```
prompt> ./p2
hello world (pid:29266)
hello, I am child (pid:29267)
hello, I am parent of 29267 (wc:29267) (pid:29266)
prompt>
```

The child will always print first.



# Process Creation API – exec()

- **exec()** system call used after a **fork()** to replace the process' memory space with a new program
- On Linux, there are six variants of exec()
  - **execl, execlp(), execle(), execv(), execvp(), execvpe()**

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <unistd.h>
4  #include <string.h>
5  #include <sys/wait.h>
6
7  int
8  main(int argc, char *argv[])
9  {
10     printf("hello world (pid:%d)\n", (int) getpid());
11     int rc = fork();
12     if (rc < 0) {          // fork failed; exit
13         fprintf(stderr, "fork failed\n");
14         exit(1);
15     } else if (rc == 0) { // child (new process)
16         printf("hello, I am child (pid:%d)\n", (int) getpid());
17         char *myargs[3];
18         myargs[0] = strdup("wc"); // program: "wc" (word count)
19         myargs[1] = strdup("p3.c"); // argument: file to count
20         myargs[2] = NULL;          // marks end of array
21         execvp(myargs[0], myargs); // runs word count
22         printf("this shouldn't print out");
23     } else {                // parent goes down this path (main)
24         int wc = wait(NULL);
25         printf("hello, I am parent of %d (wc:%d) (pid:%d)\n",
26                rc, wc, (int) getpid());
27     }
28     return 0;
29 }
```

→ a successful call to exec() never returns

```
prompt> ./p3
hello world (pid:29383)
hello, I am child (pid:29384)
    29      107    1030 p3.c
hello, I am parent of 29384 (wc:29384) (pid:29383)
prompt>
```

# Process Termination

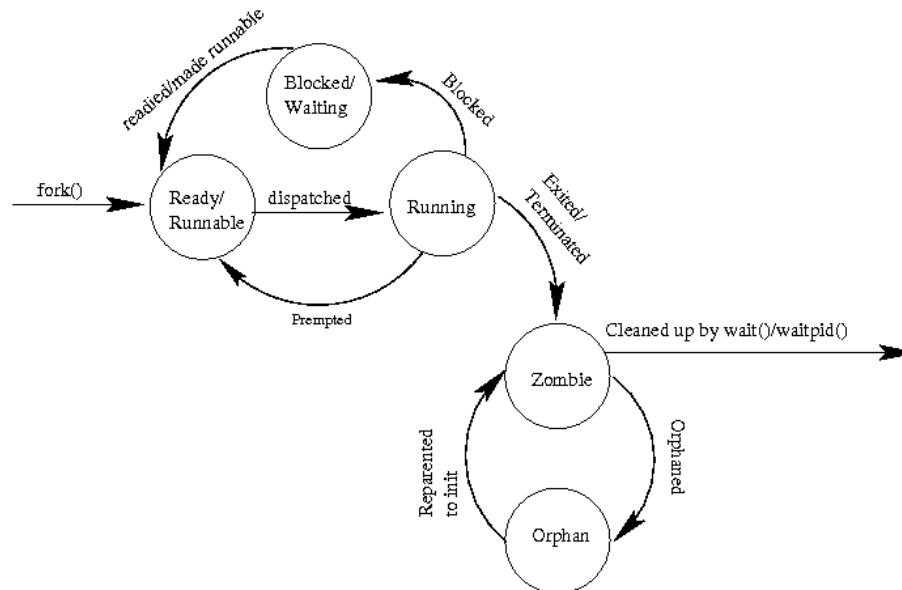
---

- **Process executes last statement and then asks the operating system to delete it using the `exit()` system call.**
  - Returns status data from child to parent (via `wait()`)
  - The resources of the now-extinct process are deallocated by operating system
- **Parent may terminate the execution of children processes using the `abort()` system call. Some reasons for doing so:**
  - Child has exceeded allocated resources
  - Task assigned to child is no longer required
  - Some operating systems does not allow a child to continue if its parent terminates.
    - All its children must also be terminated.
    - Cascading termination (All children, grandchildren, etc)
    - The termination is initiated by the operating system.

# Zombie & Orphan

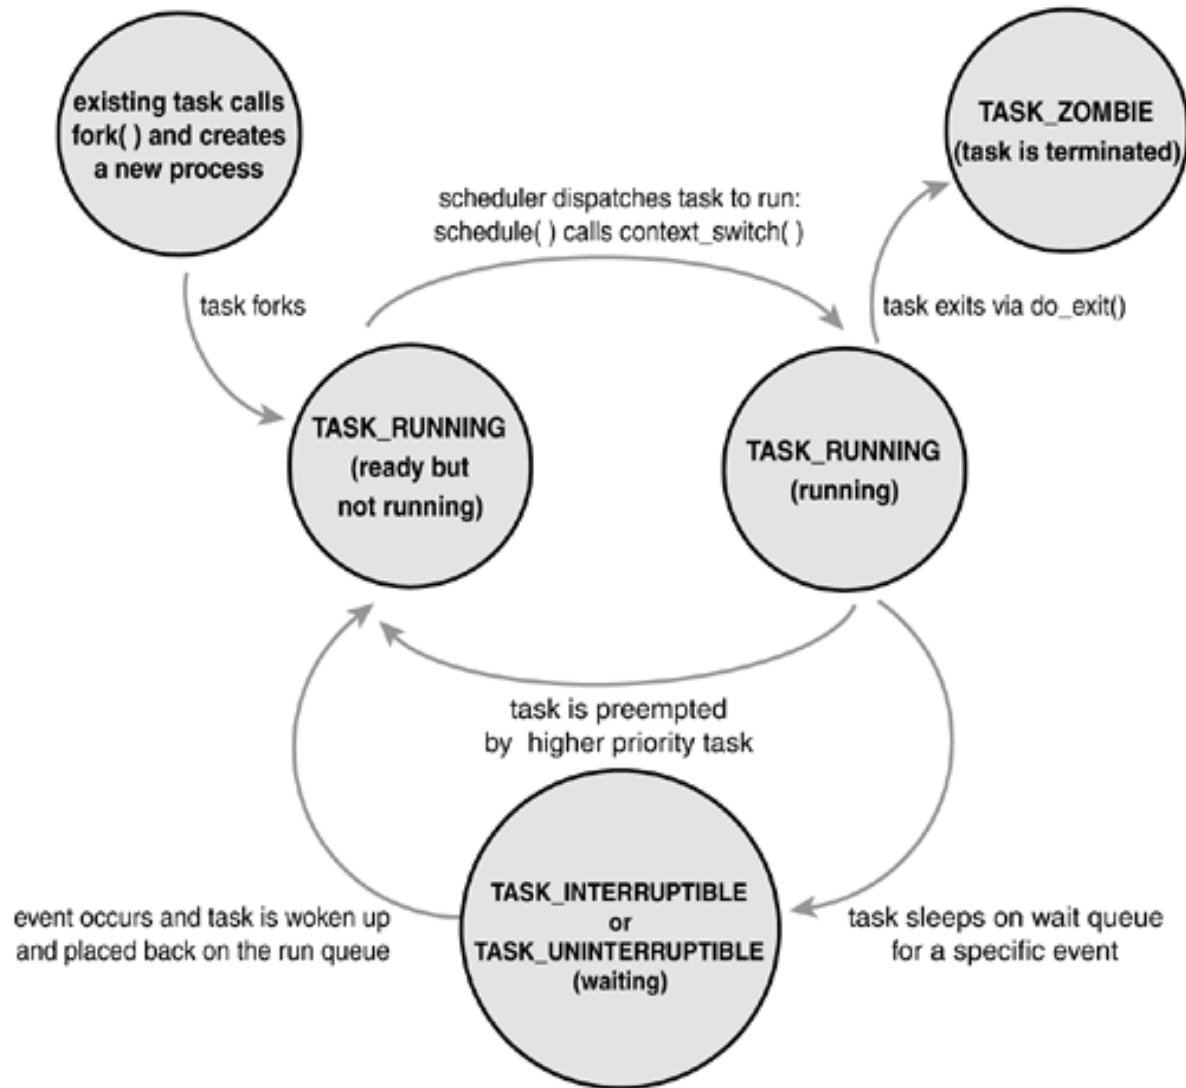
- The parent process may wait for termination of a child process by using the `wait()` system call. The call returns status information and the pid of the terminated process

```
pid = wait(&status);
```
- If no parent waiting (did not invoke `wait()`) process is a **zombie**
  - It has exited but has not yet been cleaned up
- If parent terminated without invoking `wait`, process is an **orphan**
  - May be reparented and cleaned up by `init`.
  - In modern Linux systems, an orphan process may be reparented to a "subreaper" process instead of `init`. (see `prctl(2)`)



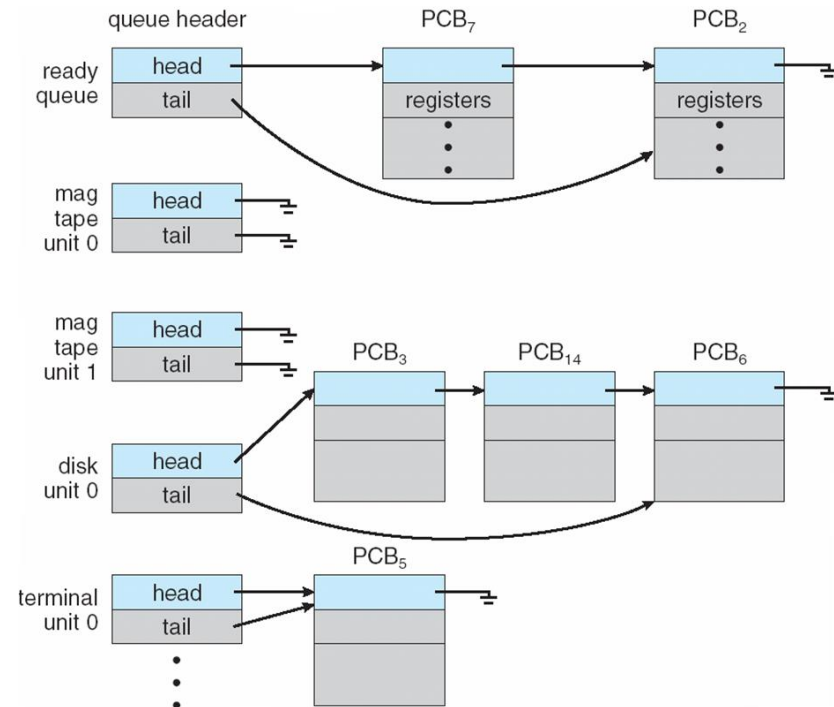
reaper

# Process States in Linux



# Scheduling Queues

- **Ready queue (ready list)**
  - Processes in ready state
    - Requesting processor, all other resources allocated
  - Process scheduling
    - Selecting a process from the ready list and dispatch it when the processor is available
- **I/O queue (device queue)**
  - Processes in blocked state
    - requesting I/O resources and waiting for their availability
    - Move to ready queue when device responds (via interrupt)
  - Separate list for each resource



# Scheduling Queues

---

