

---

# **Page Replacement**

## **Chap 21, 22**

# Virtual Memory Concept

---

- **Virtual memory**

- Concept

- A technique that allows the execution of processes that are not completely in memory
      - Partition each user's program into multiple blocks
      - Load into memory the blocks that is necessary at each time during execution
        - Only part of the program needs to be in memory for execution
        - Noncontiguous allocation
      - Logical memory size is not constrained by the amount of physical memory that is available

- Separation of logical memory as perceived by users from physical memory

# Virtual Memory Concept

---

- **Virtual memory**

- Benefits

- Easier programming

- Programmer no longer needs to worry about the amount of physical memory available

- Allows address spaces to be shared by several processes

- Higher multiprogramming degree

- Increase in CPU utilization and throughput (not in response time and turnaround time)

- Less I/O for loading and swapping processes into memory

- Faster execution of processes


# Virtual Memory Concept

---

- **Virtual memory**
  - Drawbacks
    - Address mapping overhead
    - Page fault handling overhead
    - Not adequate for real-time (embedded) systems

# Virtual Memory That is Larger Than Physical Memory

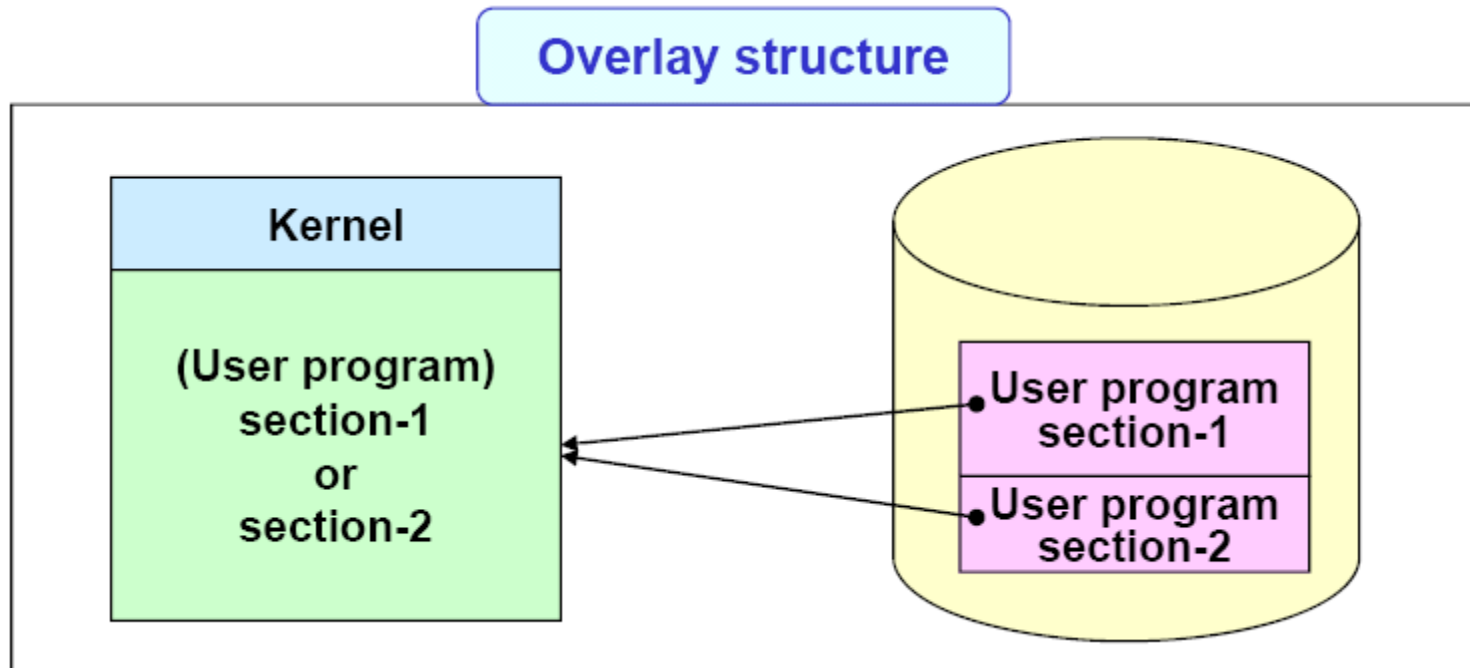
**OS make use of a larger, slower device to transparently provide the illusion of a large virtual address space**



# Memory Overlay (old system)

---


- **Program-size > memory-size**
  - w/o OS support
  - Requires support from compiler/linker/loader



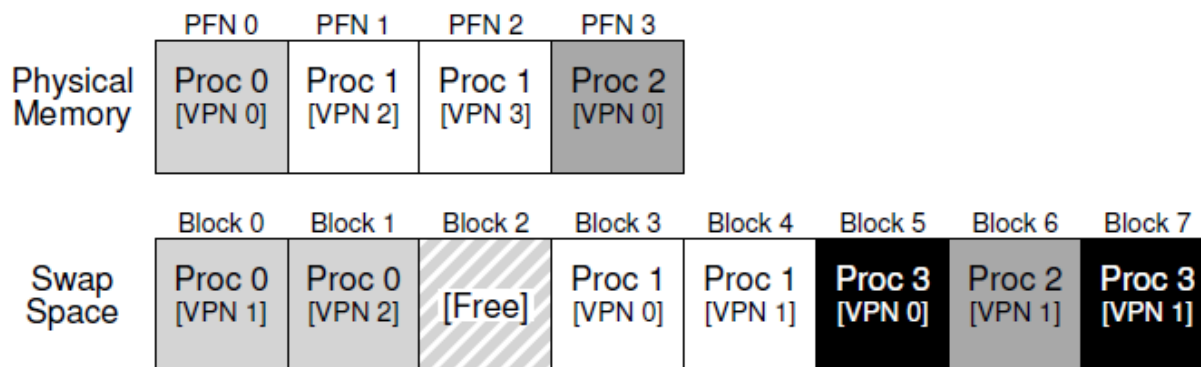
# Swapping

- A process can be swapped temporarily out of memory to a **backing store (swap device)**

Process-level  
swapping



Page-level  
swapping



# Swapping

---

- **Notes on swapping**

- Time quantum vs swap time
  - Time quantum should be substantially larger than swap time (context switch time) for efficient CPU utilization
- Memory areas to be swapped out
  - Swap only what is actually used
- Pending I/O
  - If the I/O is asynchronously accessing the user memory for I/O buffers, then the process cannot be swapped
  - Solutions
    - Never swap a process with pending I/O
    - Execute I/O operations only into kernel buffers (and deliver it to the process memory when the process is swapped in)



# Demand Paging

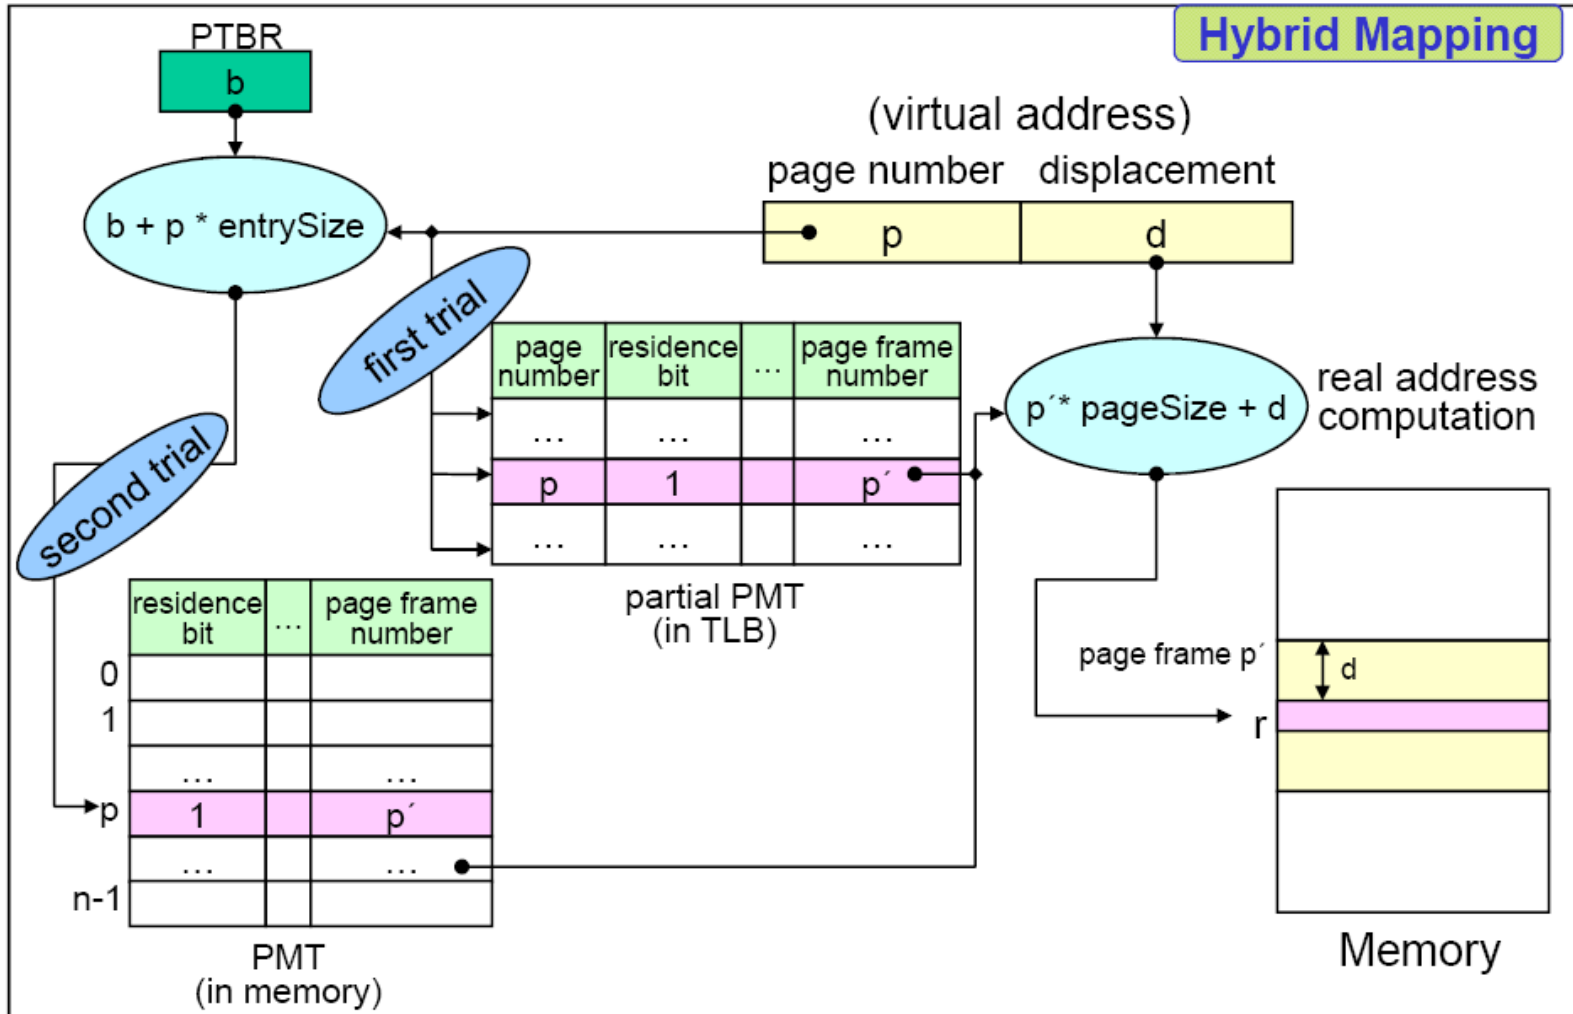
---

- **Paging (Demand paging) system**

- Partition the program into the same size blocks (pages)
- Loading of executable program
  - Initially, load pages only as they are needed
  - During execution, load the pages when they are demanded (referenced)
  - Pages that are never accessed are never loaded into physical memory
- With each page table entry a **present (residence)** bit is associated
  - **Present = true:** in-memory, **memory resident**
  - **Present = false:** not-in-memory
- Initially present bit is set to **false** on all entries
- During MMU address translation, if present bit in page table entry is **false**  $\Rightarrow$  page fault

# Demand Paging

- Address Mapping



# Page Fault

---

- If there is a reference to a page, first reference to that page will trap to operating system:  
**page fault**
  1. Operating system looks at another table to decide:
    - Invalid reference  $\Rightarrow$  abort
    - Just not in memory
  2. Find free frame
  3. Swap page into frame via scheduled disk operation
  4. Reset tables to indicate page now in memory  
Set **present** bit = **T**
  5. Restart the instruction that caused the page fault

# Steps in Handling a Page Fault

```
1  VPN = (VirtualAddress & VPN_MASK) >> SHIFT
2  (Success, TlbEntry) = TLB_Lookup(VPN)
3  if (Success == True)    // TLB Hit
4      if (CanAccess(TlbEntry.ProtectBits) == True)
5          Offset    = VirtualAddress & OFFSET_MASK
6          PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
7          Register = AccessMemory(PhysAddr)
8      else
9          RaiseException(PROTECTION_FAULT)
10 else                    // TLB Miss
11     PTEAddr = PTBR + (VPN * sizeof(PTE))
12     PTE = AccessMemory(PTEAddr)
13     if (PTE.Valid == False)
14         RaiseException(SEGMENTATION_FAULT)
15     else
16         if (CanAccess(PTE.ProtectBits) == False)
17             RaiseException(PROTECTION_FAULT)
18         else if (PTE.Present == True)
19             // assuming hardware-managed TLB
20             TLB_Insert(VPN, PTE.PFN, PTE.ProtectBits)
21             RetryInstruction()
22         else if (PTE.Present == False)
23             RaiseException(PAGE_FAULT)
```

```
1  PFN = FindFreePhysicalPage()
2  if (PFN == -1)                // no free page found
3      PFN = EvictPage()        // run replacement algorithm
4  DiskRead(PTE.DiskAddr, pfn) // sleep (waiting for I/O)
5  PTE.present = True           // update page table with present
6  PTE.PFN     = PFN            // bit and translation (PFN)
7  RetryInstruction()           // retry instruction
```

# Stages in Demand Paging

---

- 1. Trap to the operating system**
- 2. Save the user registers and process state**
- 3. Determine that the interrupt was a page fault**
- 4. Check that the page reference was legal and determine the location of the page on the disk**
- 5. Issue a read from the disk to a free frame:**
  1. Wait in a queue for this device until the read request is serviced
  2. Wait for the device seek and/or latency time
  3. Begin the transfer of the page to a free frame
- 6. While waiting, allocate the CPU to some other user**
- 7. Receive an interrupt from the disk I/O subsystem (I/O completed)**
- 8. Save the registers and process state for the other user**
- 9. Determine that the interrupt was from the disk**
- 10. Correct the page table and other tables to show page is now in memory**
- 11. Wait for the CPU to be allocated to this process again**
- 12. Restore the user registers, process state, and new page table, and then resume the interrupted instruction**

# Performance of Demand Paging

---

- **Effective access time**

- Memory access time
  - 10 ~ 200 nanoseconds (Assume 200ns)
- Average paging service time: about 8 ms
- Page fault rate:  $p$  ( $0 \leq p \leq 1$ )
- EAT(Effective Access Time)
  - $$\begin{aligned} \text{EAT} &= (1-p) \cdot m_a + p \cdot \text{PagingTime} \\ &= (1-p) \cdot 200 + p \cdot 8,000,000 \\ &= 200 + 7,999,800 \cdot p \end{aligned}$$
  - When  $p = 1/1000$ ,  $\text{EAT} = 8.2 \text{ us}$  ( $40 \times m_a$ )
- If we want less than 10% degradation,
  - $\text{EAT} = 200 + 7,999,800 \cdot p < 220$
  - $P < 0.0000025$  ( $= 1/400,000$ )

# Demand Paging Optimizations

---

- **Swap space I/O faster than file system I/O even if on the same device**
  - Swap allocated in larger chunks, less management needed than file system
- **Pages from program binary on disk can be simply discarded rather than paging out when freeing frame**
  - Still need to write to swap space
    - Pages not associated with a file (like stack and heap) – **anonymous memory**
    - Pages modified in memory but not yet written back to the file system
- **Prefetching**
  - OS could guess that a page is about to be used, and thus bring it in ahead of time
- **Mobile systems**
  - Typically don't support swapping
  - Instead, Low Memory Killer
  - cf. **zswap**

# Page Replacement

---

- Prevent over-allocation of memory
- Use **modify** (**update**, **dirty**) bit to reduce overhead of page transfers
  - only modified pages are written to disk
  - If `modify == 1`, the contents of the page in memory and in disk are not same
    - Write-back (to disk) is necessary for the page
- Large virtual memory can be provided on a smaller physical memory



# Page Replacement


---

1. Find the location of the desired page on disk
2. Find a free frame:
  - If there is a free frame, use it
  - If there is no free frame, use a page replacement algorithm to select a **victim frame**
    - Write victim frame to disk if dirty
3. Bring the desired page into the (newly) free frame; update the page and frame tables
4. Continue the process by restarting the instruction that caused the trap

Note now potentially 2 page transfers for page fault – increasing EAT

# Graph of Page Faults Versus The Number of Frames

---




# When Replacements Really Occur?

---

- OS keeps a small portion of memory free more proactively
- **Watermark** scheme
  - high watermark (HW) and low watermark (LW)
  - When OS notices that there are fewer than LW pages available, a background thread (**swap daemon** or **page daemon**) that is responsible for freeing memory runs.
  - The thread evicts pages until there are HW pages available.
  - The background thread then goes to sleep.
  - many systems will **cluster** or **group** a number of pages and write them out at once to the swap partition, thus increasing the efficiency of the disk

# Memory Reclamation at Linux

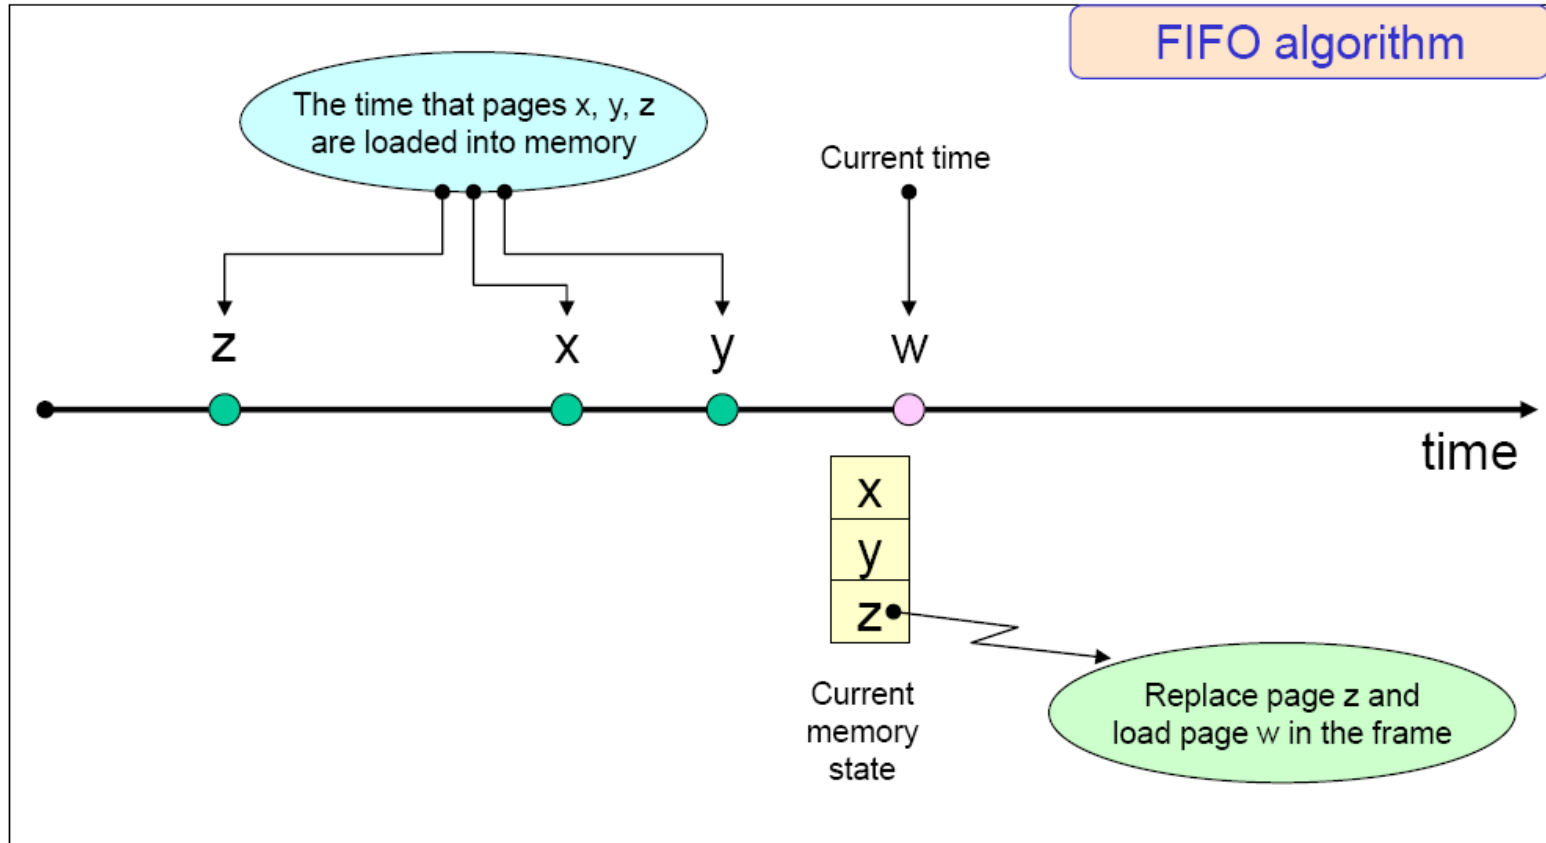


# First-In-First-Out (FIFO) Algorithm

---

- Choose the page to be replaced based on when the page is previously loaded into memory
- Scheme
  - Replace the oldest page
- Requirements
  - Timestamping (memory load time for each page) is necessary
- Characteristics
  - May replace frequently used pages
- **FIFO anomaly (Belady's anomaly)**
  - In FIFO algorithm, page fault frequency may increase even if more memory frames are allocated

# First-In-First-Out (FIFO) Algorithm



# First-In-First-Out (FIFO) Algorithm

- FIFO algorithm: Example**

- 4 page frames allocated, initially empty

$\omega = 1\ 2\ 6\ 1\ 4\ 5\ 1\ 2\ 1\ 4\ 5\ 6\ 4\ 5$

Memory state change (FIFO)

| Time         | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 |
|--------------|---|---|---|---|---|---|---|---|---|----|----|----|----|----|
| Ref. string  | 1 | 2 | 6 | 1 | 4 | 5 | 1 | 2 | 1 | 4  | 5  | 6  | 4  | 5  |
| Memory state | 1 | 1 | 1 | 1 | 1 | 5 | 5 | 5 | 5 | 5  | 5  | 5  | 4  | 4  |
|              |   | 2 | 2 | 2 | 2 | 2 | 1 | 1 | 1 | 1  | 1  | 1  | 1  | 5  |
|              |   |   | 6 | 6 | 6 | 6 | 6 | 2 | 2 | 2  | 2  | 2  | 2  | 2  |
|              |   |   |   |   | 4 | 4 | 4 | 4 | 4 | 4  | 4  | 6  | 6  | 6  |
| Page fault   | F | F | F |   | F | F | F | F |   |    |    | F  | F  | F  |

▪ Number of page faults: 10

# First-In-First-Out (FIFO) Algorithm

- FIFO algorithm: Anomaly example

$\omega = 1\ 2\ 3\ 4\ 1\ 2\ 5\ 1\ 2\ 3\ 4\ 5$

Number of page frames: 3

| Time         | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 |
|--------------|---|---|---|---|---|---|---|---|---|----|----|----|
| Ref. string  | 1 | 2 | 3 | 4 | 1 | 2 | 5 | 1 | 2 | 3  | 4  | 5  |
| Memory state | 1 | 1 | 1 | 4 | 4 | 4 | 5 | 5 | 5 | 5  | 5  | 5  |
|              |   | 2 | 2 | 2 | 1 | 1 | 1 | 1 | 1 | 3  | 3  | 3  |
|              |   |   | 3 | 3 | 3 | 2 | 2 | 2 | 2 | 2  | 4  | 4  |
| Page fault   | F | F | F | F | F | F | F |   |   | F  | F  |    |

▪ Number of page faults: 9

Number of page frames: 4

| Time         | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 |
|--------------|---|---|---|---|---|---|---|---|---|----|----|----|
| Ref. string  | 1 | 2 | 3 | 4 | 1 | 2 | 5 | 1 | 2 | 3  | 4  | 5  |
| Memory state | 1 | 1 | 1 | 1 | 1 | 1 | 5 | 5 | 5 | 5  | 4  | 4  |
|              |   | 2 | 2 | 2 | 2 | 2 | 2 | 1 | 1 | 1  | 1  | 5  |
|              |   |   | 3 | 3 | 3 | 3 | 3 | 3 | 2 | 2  | 2  | 2  |
|              |   |   | 4 | 4 | 4 | 4 | 4 | 4 | 4 | 3  | 3  | 3  |
| Page fault   | F | F | F | F |   |   |   | F | F | F  | F  | F  |

▪ Number of page faults: 10




# MIN algorithm (OPT algorithm)

---

- Proposed by Belady in 1966
- Minimizes page fault frequency (proved)
- Scheme
  - Replace the page that will not be used for the longest period of time
  - Tie-breaking rule
    - Page with greatest (or smallest) page number
- Unrealizable
  - Can be used only when the process's reference string is known a priori
- Usage
  - Performance measurement tool for replacement schemes

# MIN algorithm (OPT algorithm)



# MIN algorithm (OPT algorithm)

- **MIN algorithm: Example**

- 4 page frames allocated, initially empty

$\omega = 1\ 2\ 6\ 1\ 4\ 5\ 1\ 2\ 1\ 4\ 5\ 6\ 4\ 5$

Memory state change (MIN)

| Time         | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 |
|--------------|---|---|---|---|---|---|---|---|---|----|----|----|----|----|
| Ref. string  | 1 | 2 | 6 | 1 | 4 | 5 | 1 | 2 | 1 | 4  | 5  | 6  | 4  | 5  |
| Memory state | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1  | 1  | 6  | 6  | 6  |
|              |   | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2  | 2  | 2  | 2  | 2  |
|              |   |   | 6 | 6 | 6 | 5 | 5 | 5 | 5 | 5  | 5  | 5  | 5  | 5  |
|              |   |   |   |   | 4 | 4 | 4 | 4 | 4 | 4  | 4  | 4  | 4  | 4  |
| Page fault   | F | F | F |   | F | F |   |   |   |    |    | F  |    |    |


▪ Number of page faults: 6

# Least Recently Used (LRU) Algorithm

---

- Choose the page to be replaced based on the reference time
- Scheme
  - Replace the page that has not been used for the longest period of time
- Requirements
  - Timestamping (page reference time) is necessary
- Characteristics
  - Based on program locality
  - Approximates to the performance of MIN algorithm
- **Used in most practical systems**
- Drawbacks
  - Timestamping overhead at every page reference
  - Number of page faults increases steeply when the process executes large loop with insufficiently allocated memory

# Least Recently Used (LRU) Algorithm



# Least Recently Used (LRU) Algorithm

- LRU algorithm: Example

- 4 page frames allocated, initially empty

$\omega = 1\ 2\ 6\ 1\ 4\ 5\ 1\ 2\ 1\ 4\ 5\ 6\ 4\ 5$

Memory state change (LRU)

| Time         | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 |
|--------------|---|---|---|---|---|---|---|---|---|----|----|----|----|----|
| Ref. string  | 1 | 2 | 6 | 1 | 4 | 5 | 1 | 2 | 1 | 4  | 5  | 6  | 4  | 5  |
| Memory state | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1  | 1  | 1  | 1  | 1  |
|              |   | 2 | 2 | 2 | 2 | 5 | 5 | 5 | 5 | 5  | 5  | 5  | 5  | 5  |
|              |   |   | 6 | 6 | 6 | 6 | 6 | 2 | 2 | 2  | 2  | 6  | 6  | 6  |
|              |   |   |   |   | 4 | 4 | 4 | 4 | 4 | 4  | 4  | 4  | 4  | 4  |
| Page fault   | F | F | F |   | F | F |   | F |   |    |    | F  |    |    |

▪ Number of page faults: 7

# Implementation of LRU algorithm

---

- By counter
  - Use PMT with count field
  - Increment processor clock or counter for each memory access
  - Record the value of processor clock or counter in the corresponding PMT entry for each page reference
  - Can get the relative order of recent access to each page
  - PMT search for selecting a page to be replaced

# Implementation of LRU algorithm

---

- By stack
  - Stack
    - Stack for each process, whose entry is page number
    - Maintains the stack elements (page numbers) in the order of recent access
    - Can delete an element in the middle of the stack
  - When no page fault
    - Deletes the referenced page number from the stack, and inserts it on top of the stack
  - When page fault
    - Displaces the page whose number is at the bottom of the stack, deletes it from the stack, and inserts incoming page number on top of the stack