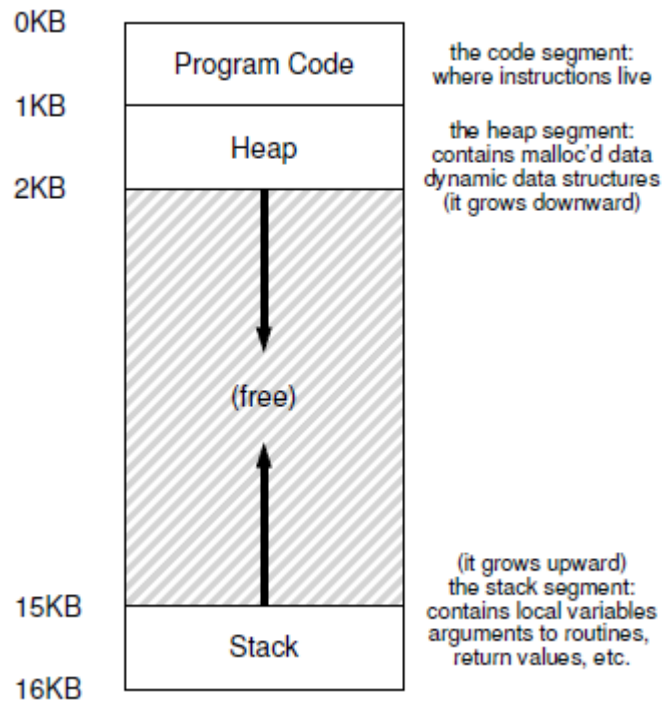

Concurrency, Thread


Thread

- Classic view
 - a single point of execution within a program
 - a single PC where instructions are being fetched from and executed),
- Multi-threaded program
 - Has more than one point of execution
 - multiple PCs, each of which is being fetched and executed from.
- Threads **share** the same address space and thus can access the same data
- Each thread has its own private set of registers (including PC)
- When switching from running one (T1) to running the other (T2), a **context switch** must take place
 - thread control blocks (TCBs)
 - the address space remains the same (i.e., no need to switch the page table).

Multiple Stacks

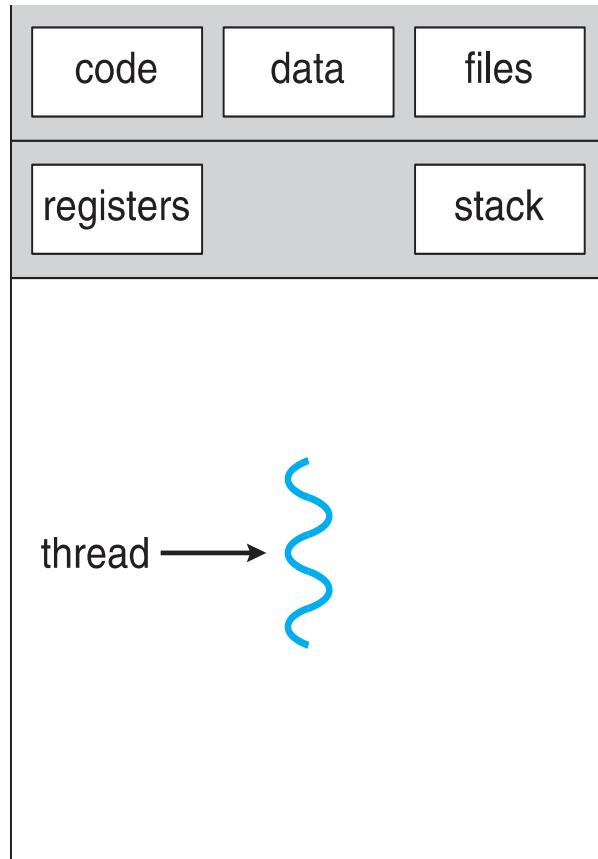


Single-Threaded

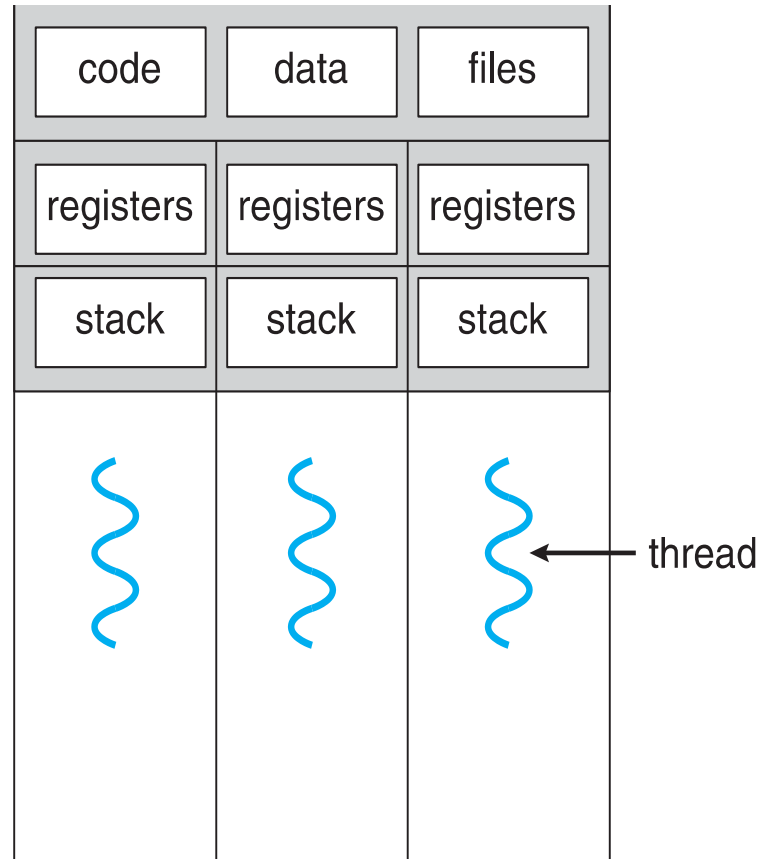


Multi-Threaded

Single and Multithreaded Processes



single-threaded process



multithreaded process

Why Use Threads?

- **Parallelism**
 - One thread per CPU can make programs run faster on multiple processors
- **I/O overlapping (even in single processor)**
 - Avoid blocking program progress due to slow I/O
 - While one thread in your program waits (i.e., blocked waiting for I/O), the CPU scheduler can switch to other threads, which are ready to run and do something useful.
 - Similar to the effect of multiprogramming
- You could use multiple **processes** instead of **threads**.
 - Threads share an address space and thus make it easy to share data
 - Processes are a more sound choice for logically separate tasks

Thread

- Most modern applications are multithreaded
- Threads run within application
- Multiple tasks with the application can be implemented by separate threads
 - Update display
 - Fetch data
 - Spell checking
 - Answer a network request
- Process creation is heavy-weight while thread creation is **light-weight**
- Can simplify code, increase efficiency
- Kernels are generally multithreaded

Benefits

- **Responsiveness**
 - may allow continued execution if part of process is blocked, especially important for user interfaces
- **Resource Sharing**
 - threads share resources of process, easier than shared memory or message passing
- **Economy**
 - cheaper than process creation, thread switching lower overhead than context switching
- **Scalability**
 - process can take advantage of multiprocessor architectures

Thread Creation

```
1  #include <stdio.h>
2  #include <assert.h>
3  #include <pthread.h>
4
5  void *mythread(void *arg) {
6      printf("%s\n", (char *) arg);
7      return NULL;
8  }
9
10 int
11 main(int argc, char *argv[]) {
12     pthread_t p1, p2;
13     int rc;
14     printf("main: begin\n");
15     rc = pthread_create(&p1, NULL, mythread, "A"); assert(rc == 0);
16     rc = pthread_create(&p2, NULL, mythread, "B"); assert(rc == 0);
17     // join waits for the threads to finish
18     rc = pthread_join(p1, NULL); assert(rc == 0);
19     rc = pthread_join(p2, NULL); assert(rc == 0);
20     printf("main: end\n");
21     return 0;
22 }
```

Once a thread is created, it may start running right away, or ready

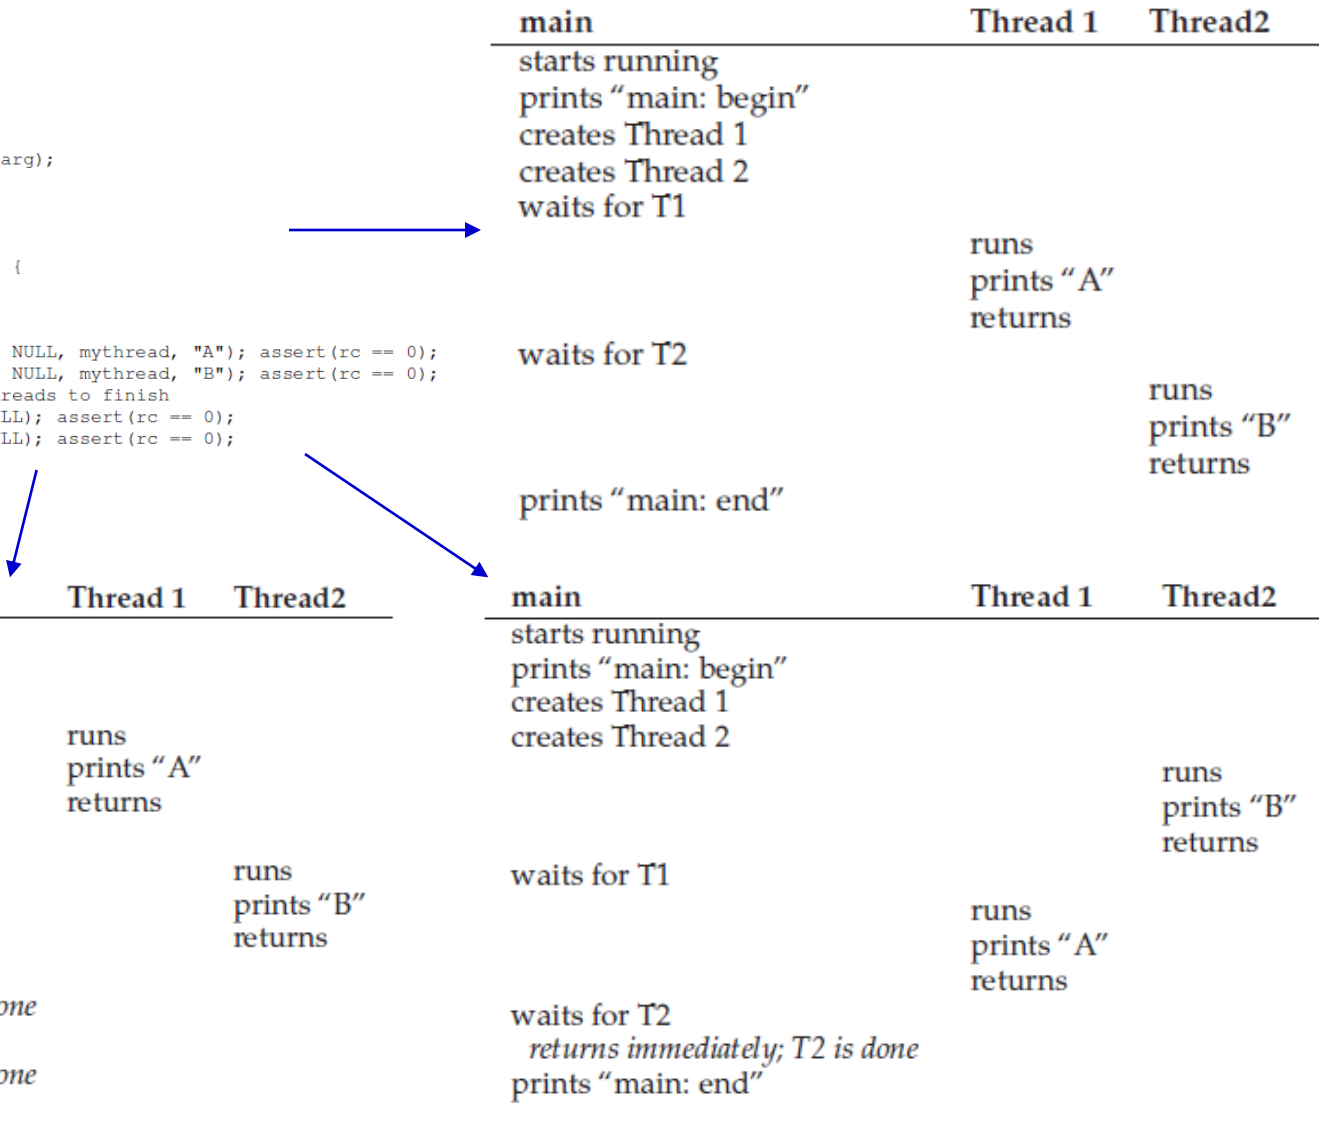

pthread join() waits for a particular thread to complete

Indeterminate executions

```

1  #include <stdio.h>
2  #include <assert.h>
3  #include <pthread.h>
4
5  void *mythread(void *arg) {
6      printf("%s\n", (char *) arg);
7      return NULL;
8  }
9
10 int
11 main(int argc, char *argv[]) {
12     pthread_t p1, p2;
13     int rc;
14     printf("main: begin\n");
15     rc = pthread_create(&p1, NULL, mythread, "A"); assert(rc == 0);
16     rc = pthread_create(&p2, NULL, mythread, "B"); assert(rc == 0);
17     // join waits for the threads to finish
18     rc = pthread_join(p1, NULL); assert(rc == 0);
19     rc = pthread_join(p2, NULL); assert(rc == 0);
20     printf("main: end\n");
21     return 0;
22 }

```



Why It Gets Worse: Shared Data

```
1  #include <stdio.h>
2  #include <pthread.h>
3  #include "mythreads.h"
4
5  static volatile int counter = 0;
6
7  //
8  // mythread()
9  //
10 // Simply adds 1 to counter repeatedly, in a loop
11 // No, this is not how you would add 10,000,000 to
12 // a counter, but it shows the problem nicely.
13 //
14 void *
15 mythread(void *arg)
16 {
17     printf("%s: begin\n", (char *) arg);
18     int i;
19     for (i = 0; i < 1e7; i++) {
20         counter = counter + 1;
21     }
22     printf("%s: done\n", (char *) arg);
23     return NULL;
24 }
25
```

```
prompt> gcc -o main main.c -Wall -pthread
prompt> ./main
main: begin (counter = 0)
A: begin
B: begin
A: done
B: done
main: done with both (counter = 20000000)
```

[expected]

```
26 //
27 // main()
28 //
29 // Just launches two threads (pthread_create)
30 // and then waits for them (pthread_join)
31 //
32 int
33 main(int argc, char *argv[])
34 {
35     pthread_t p1, p2;
36     printf("main: begin (counter = %d)\n", counter);
37     Pthread_create(&p1, NULL, mythread, "A");
38     Pthread_create(&p2, NULL, mythread, "B");
39
40     // join waits for the threads to finish
41     Pthread_join(p1, NULL);
42     Pthread_join(p2, NULL);
43     printf("main: done with both (counter = %d)\n", counter);
44     return 0;
45 }
```

```
prompt> ./main
main: begin (counter = 0)
A: begin
B: begin
A: done
B: done
main: done with both (counter = 19345221)
```

even indeterminate



[real]

Uncontrolled Scheduling

counter = counter + 1;



```
100    mov 0x8049a1c, %eax
105    add $0x1, %eax
108    mov %eax, 0x8049a1c
```

OS	Thread 1	Thread 2	(after instruction)			
			PC	%eax	counter	
	<i>before critical section</i>		100	0	50	each thread has its own private registers
	mov 0x8049a1c, %eax		105	50	50	
	add \$0x1, %eax		108	51	50	
interrupt	<i>save T1's state</i>					
	<i>restore T2's state</i>		100	0	50	
		mov 0x8049a1c, %eax	105	50	50	
		add \$0x1, %eax	108	51	50	
		mov %eax, 0x8049a1c	113	51	51	
interrupt	<i>save T2's state</i>					
	<i>restore T1's state</i>		108	51	51	
	mov %eax, 0x8049a1c		113	51	51	

Race Condition

- **Race condition**
 - Results depend on the timing execution of the code, indeterminate
- **Critical section**
 - a piece of code that accesses a shared variable (or more generally, a shared resource) and must not be concurrently executed by more than one thread.
- **Mutual exclusion.**
 - guarantees that if one thread is executing within the critical section, the others will be prevented from doing so.



Edsger Dijkstra
Turing award 1972

Wish For Atomicity

- We need more powerful instructions
 - do exactly whatever we needed done in a single step → **atomic**
 - could not be interrupted mid-instruction, because that is precisely the guarantee we receive from the hardware
 - remove the possibility of an untimely interrupt
 - E.g. `memory-add 0x8049a1c, $0x1`
- Instead of many complex atomic instructions, hardware provides a few useful instructions upon which we can build a general set of **synchronization primitives**.

One More Problem: Waiting For Another

- Another common interaction
 - One thread must wait for another to complete some action before it continues
 - e.g., when a process performs a disk I/O and is put to sleep; when the I/O completes, the process needs to be roused from its slumber so it can continue.
- **Condition variables**

Thread Libraries

- **Thread library provides programmer with API for creating and managing threads**
- **Two primary ways of implementing**
 - Library entirely in user space
 - Function call, not system call
 - Kernel-level library supported by the OS

Pthreads

- May be provided either as user-level or kernel-level
- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- *Specification*, not *implementation*
- API specifies behavior of the thread library, implementation is up to development of the library
- Common in UNIX operating systems (Solaris, Linux, Mac OS X)
- Compile
 - prompt> gcc -o main main.c -Wall -pthread

Thread API

- Thread Creation
- Thread Completion
- Locks
- Condition Variables

Thread Creation

```
#include <pthread.h>
int
pthread_create(      pthread_t *      thread,
                    const pthread_attr_t * attr,
                    void *              (*start_routine) (void*),
                    void *              arg);
```

- Arguments
 - **thread**: a pointer to a structure of type `pthread_t`
 - Used to interact with this thread
 - **attr**: specify any attributes this thread might have
 - e.g., stack size or scheduling priority
 - initialized with a separate call to `pthread_attr_init()`
 - **function pointer**: which function should this thread start running in?
 - **arg**: the argument to be passed to the function where the thread begins execution
 - void pointer allows us to pass in any type of argument

Thread Creation

```
1  #include <pthread.h>
2
3  typedef struct __myarg_t {
4      int a;
5      int b;
6  } myarg_t;
7
8  void *mythread(void *arg) {
9      myarg_t *m = (myarg_t *) arg;
10     printf("%d %d\n", m->a, m->b);
11     return NULL;
12 }
13
14 int
15 main(int argc, char *argv[]) {
16     pthread_t p;
17     int rc;
18
19     myarg_t args;
20     args.a = 10;
21     args.b = 20;
22     rc = pthread_create(&p, NULL, mythread, &args);
23     ...
24 }
```

unpack the arguments as desired

package into a single type

Thread Completion

```
int pthread_join(pthread_t thread, void **value_ptr);
```

- Wait for a thread to complete
- Arguments
 - pthread_t: specify which thread to wait for.
 - a pointer to the return value

```
1  #include <stdio.h>
2  #include <pthread.h>
3  #include <assert.h>
4  #include <stdlib.h>
5
6  typedef struct __myarg_t {
7      int a;
8      int b;
9  } myarg_t;
10
11 typedef struct __myret_t {
12     int x;
13     int y;
14 } myret_t;
15
16 void *mythread(void *arg) {
17     myarg_t *m = (myarg_t *) arg;
18     printf("%d %d\n", m->a, m->b);
19     myret_t *r = Malloc(sizeof(myret_t));
20     r->x = 1;
21     r->y = 2;
22     return (void *) r;
23 }
24
25 int
26 main(int argc, char *argv[]) {
27     pthread_t p;
28     myret_t *m;
29
30     myarg_t args = {10, 20};
31     Pthread_create(&p, NULL, mythread, &args);
32     Pthread_join(p, (void **) &m);
33     printf("returned %d %d\n", m->x, m->y);
34     free(m);
35     return 0;
36 }
```

Thread Completion

- don't return a pointer which refers to something allocated on the thread's call stack
 - automatically deallocated

```
1 void *mythread(void *arg) {
2     myarg_t *m = (myarg_t *) arg;
3     printf("%d %d\n", m->a, m->b);
4     myret_t r; // ALLOCATED ON STACK: BAD!
5     r.x = 1;
6     r.y = 2;
7     return (void *) &r;
8 }
```

Locks

```
int pthread_mutex_lock(pthread_mutex_t *mutex);  
int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

- Providing mutual exclusion to a critical section via **locks**

```
pthread_mutex_t lock;  
pthread_mutex_lock(&lock);  
x = x + 1; // or whatever your critical section is  
pthread_mutex_unlock(&lock);
```

- If no other thread holds the lock, the thread will acquire the lock and enter the critical section.
- If another thread does indeed hold the lock, the thread trying to grab the lock will not return from the call until it has acquired the lock
- Lock initialization
 - Static `pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;`
 - Dynamic `int rc = pthread_mutex_init(&lock, NULL);
assert(rc == 0); // always check success!`
- Lock destroy
 - `pthread_mutex_destroy()`

trylock

```
int pthread_mutex_trylock(pthread_mutex_t *mutex);  
int pthread_mutex_timedlock(pthread_mutex_t *mutex,  
                             struct timespec *abs_timeout);
```

- Returns failure if the lock is already held
- the timedlock version of acquiring a lock returns after a timeout or after acquiring the lock, whichever happens first.
- avoid getting stuck (perhaps indefinitely) in a lock acquisition routine (prevent deadlock).

Condition Variables

```
int pthread_cond_wait(pthread_cond_t *cond, pthread_mutex_t *mutex);  
int pthread_cond_signal(pthread_cond_t *cond);
```

- To use a condition variable, one has to in addition have a lock that is associated with this condition
 - Prevent a race condition
- Wait
 - puts the calling thread to sleep, and thus waits for some other thread to signal it, usually when something in the program has changed that the now-sleeping thread might care about.
 - Release the lock before sleep, re-acquires the lock before returning
 - For safety, the waiting thread has to re-check the condition after wake-up (There could be a spurious wake-up)

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;  
pthread_cond_t cond = PTHREAD_COND_INITIALIZER;  
  
Pthread_mutex_lock(&lock);  
while (ready == 0)  
    Pthread_cond_wait(&cond, &lock);  
Pthread_mutex_unlock(&lock);
```

```
Pthread_mutex_lock(&lock);  
ready = 1;  
Pthread_cond_signal(&cond);  
Pthread_mutex_unlock(&lock);
```


Spin

```
while (ready == 0); // spin
```

instead of wait()

```
ready = 1;
```

instead of signal()

- A simple flag instead of a condition variable and associated lock
- Don't ever do this
 - Poor performance: spinning for a long time just wastes CPU cycles
 - Error prone

User Threads and Kernel Threads

- **User threads**

- management done by user-level threads library
- Three primary thread libraries:
 - POSIX **Pthreads**
 - Windows threads
 - Java threads

- **Kernel threads**

- Supported by the Kernel
- Examples – virtually all general-purpose operating systems, including: Windows, Solaris, Linux, Tru64 UNIX, Mac OS X

Implicit Threading

- **Growing in popularity as numbers of threads increase, program correctness more difficult with explicit threads**
- **Creation and management of threads done by compilers and run-time libraries rather than programmers**
- **Three methods explored**
 - Thread Pools
 - OpenMP
- **Other methods**
 - Intel Threading Building Blocks (TBB): Parallel C++ program
 - `java.util.concurrent` package

Thread Pools

- **Create a number of threads in a pool where they await work**
- **Advantages:**
 - Usually slightly faster to service a request with an existing thread than create a new thread
 - Allows the number of threads in the application(s) to be bound to the size of the pool
 - Separating task to be performed from mechanics of creating task allows different strategies for running task
 - i.e. Tasks could be scheduled to run periodically

```
public class ThreadExample {  
    public static void ThreadProc() {  
        ...  
    }  
  
    public static void Main() {  
        Thread t = new Thread(  
            new ThreadStart(ThreadProc));  
        t.Start();  
        ...  
        t.Join();  
    }  
}
```

```
public class Example {  
    public static void Main()  
    {  
        ThreadPool.QueueUserWorkItem(ThreadProc);  
    }  
  
    static void ThreadProc(Object stateInfo)  
    {  
        ...  
    }  
}
```

.NET example

OpenMP

- Set of compiler directives and an API for C, C++, FORTRAN
- Provides support for parallel programming in shared-memory environments
- Identifies **parallel regions** – blocks of code that can run in parallel

```
#pragma omp parallel
```

Create as many threads as there are cores

```
#pragma omp parallel for  
for(i=0;i<N;i++) {  
    c[i] = a[i] + b[i];  
}
```

Run for loop in parallel

```
#include <omp.h>  
#include <stdio.h>  
  
int main(int argc, char *argv[])  
{  
    /* sequential code */  
  
    #pragma omp parallel  
    {  
        printf("I am a parallel region.");  
    }  
  
    /* sequential code */  
  
    return 0;  
}
```

Linux Threads

- Linux refers to them as *tasks* rather than *threads*
- Thread creation is done through `clone()` system call
- `clone()` allows a child task to share the address space of the parent task (process)
 - for thread creation (cf> fork for process creation)
 - Flags control behavior

flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.

- `struct task_struct` points to process data structures (shared or unique)

Homework

- Homework in Chap 27 (Debugging Race/Deadlock w/ *helgrind*)