
Locks

Locks: The Basic Idea

- To implement a critical section
- A lock variable must be declared
- A lock variable holds the state of the lock
 - Available (unlocked, free)
 - Acquired (locked, held)
- Exactly one thread holds the lock, **owner**
- Once the owner of the lock calls `unlock()`, if there are waiting threads (stuck in `lock()`), one of them will (eventually) notice (or be informed of) this change of the lock's state, acquire the lock, and enter the critical section.

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;  
  
Pthread_mutex_lock(&lock); // wrapper for pthread_mutex_lock()  
balance = balance + 1;  
Pthread_mutex_unlock(&lock);
```

Evaluating Locks

- **Mutual exclusion**
 - whether the lock does its basic task
- **Fairness**
 - Does each thread contending for the lock get a fair shot at acquiring it once it is free?
 - No starvation
- **Performance**
 - time overheads added by using the lock
 - In the case of single thread (no contention)
 - In the case of multiple contending threads on a single CPU
 - In the case of multiple threads on multiple CPUs

Controlling Interrupts

- The earliest solutions used to provide mutual exclusion was to disable interrupts for critical sections (single-processor system)
 - + Simple
 - Requires to allow any calling thread to perform a **privileged** operation (interrupt on/off)
 - Does not work on multiprocessors
 - Lost interrupts
 - Inefficient: code that masks or unmask interrupts tends to be executed slowly by modern CPUs
- OS itself will use interrupt masking to guarantee atomicity when accessing its own data structures

```
void lock() {  
    DisableInterrupts();  
}  
void unlock() {  
    EnableInterrupts();  
}
```

A Failed Attempt: Just Using Loads/Stores

- Build a simple lock by using a single flag variable
- Correctness problem
- Performance problem
 - spin-waiting
 - high on a uniprocessor

```
1  typedef struct __lock_t { int flag; } lock_t;
2
3  void init(lock_t *mutex) {
4      // 0 -> lock is available, 1 -> held
5      mutex->flag = 0;
6  }
7
8  void lock(lock_t *mutex) {
9      while (mutex->flag == 1) // TEST the flag
10         ; // spin-wait (do nothing)
11     mutex->flag = 1;          // now SET it!
12 }
13
14 void unlock(lock_t *mutex) {
15     mutex->flag = 0;
16 }
```

Thread 1

call lock()

while (flag == 1)

interrupt: switch to Thread 2

flag = 1; // set flag to 1 (too!)

Thread 2

call lock()

while (flag == 1)

flag = 1;

interrupt: switch to Thread 1

Successful but Old Attempt (w/o HW support)

- **Dekker's algorithm ('68)**
- **Peterson's algorithm ('81)** P_0 enters CS only if
either $\text{flag}[1] = 0$ or $\text{turn} = 0$
- **Useless now**
 - H/W support for lock and relaxed memory consistency models

```
int flag[2];
int turn;

void init() {
    flag[0] = flag[1] = 0; // 1->thread wants to grab lock
    turn = 0; // whose turn? (thread 0 or 1?)
}

void lock() {
    flag[self] = 1; // self: thread ID of caller
    turn = 1 - self; // make it other thread's turn
    while ((flag[1-self] == 1) && (turn == 1 - self))
        ; // spin-wait
}

void unlock() {
    flag[self] = 0; // simply undo your intent
}
```

Building Working Spin Locks with test-and-set

- We need hardware support!
- test-and-set instruction (atomic exchange)
 - returns the old value pointed to by the `old_ptr` (test), and simultaneously updates said value to `new` (set).
 - this sequence of operations is performed **atomically** (uninterruptable)

```
int TestAndSet(int *old_ptr, int new) {  
    int old = *old_ptr; // fetch old value at old_ptr  
    *old_ptr = new; // store 'new' into old_ptr  
    return old; // return the old value  
}
```

Building Working Spin Locks with **test-and-set**

- As long as the lock is held by another thread, TestAndSet() will repeatedly return 1, and thus this thread will spin and spin until the lock is finally released.
- By making both the **test** and **set** a single atomic operation, we ensure that only one thread acquires the lock.
- To work correctly on a single processor, it requires a **preemptive scheduler**

```
1  typedef struct __lock_t {
2      int flag;
3  } lock_t;
4
5  void init(lock_t *lock) {
6      // 0 indicates that lock is available, 1 that it is held
7      lock->flag = 0;
8  }
9
10 void lock(lock_t *lock) {
11     while (TestAndSet(&lock->flag, 1) == 1)
12         ; // spin-wait (do nothing)
13 }
14
15 void unlock(lock_t *lock) {
16     lock->flag = 0;
17 }
```


Evaluating Spin Locks

- **Correctness?** Yes
 - allows a single thread to enter CS at a time
- **Fairness?** Bad
 - A thread spinning may spin forever
- **Performance?**
 - Single CPU case: high overhead
 - If the thread holding the lock is pre-empted within a critical section, other scheduled threads try to acquire the lock.
 - Each of those threads will spin for the duration of a time slice before giving up the CPU, a waste of CPU cycles.
 - Multiple CPUs: work reasonably well
 - Thread A on CPU 1 and Thread B on CPU 2, both contending for a lock.
 - Spinning to wait for a lock held on another processor doesn't waste many cycles in this case, and thus can be effective.

Compare-And-Swap (compare-and-exchange)

- Test whether the value at the address specified by `ptr` is equal to `expected`
 - If so, update the memory location pointed to by `ptr` with the `new` value.
 - If not, do nothing
- Similar behavior but more powerful instruction than test-and-set
 - Useful for **lock-free synchronization**

```
int CompareAndSwap(int *ptr, int expected, int new) {  
    int actual = *ptr;  
    if (actual == expected)  
        *ptr = new;  
    return actual;  
}
```

```
void lock(lock_t *lock) {  
    while (CompareAndSwap(&lock->flag, 0, 1) == 1)  
        ; // spin  
}
```

Load-Linked and Store-Conditional

- load-linked and store-conditional instructions can be used in tandem to build locks and other concurrent structures
- MIPS, Alpha, PowerPC, ARM
- store-conditional only succeeds if no intervening store to the address has taken place

```
int LoadLinked(int *ptr) {  
    return *ptr;  
}  
  
int StoreConditional(int *ptr, int value) {  
    if (no one has updated *ptr since the LoadLinked to this address) {  
        *ptr = value;  
        return 1; // success!  
    } else  
        return 0; // failed to update  
}
```

Load-Linked and Store-Conditional

```
void lock(lock_t *lock) {
    while (1) {
        while (LoadLinked(&lock->flag) == 1)
            ; // spin until it's zero
        if (StoreConditional(&lock->flag, 1) == 1)
            return; // if set-it-to-1 was a success: all done
        // otherwise: try it all over again
    }
}

void unlock(lock_t *lock) {
    lock->flag = 0;
}
```

or

```
void lock(lock_t *lock) {
    while (LoadLinked(&lock->flag) || !StoreConditional(&lock->flag, 1))
        ; // spin
}
```

Fetch-And-Add

```
int FetchAndAdd(int *ptr) {
    int old = *ptr;
    *ptr = old + 1;
    return old;
}
```

```
typedef struct __lock_t {
    int ticket;
    int turn;
} lock_t;
```

```
void lock_init(lock_t *lock) {
    lock->ticket = 0;
    lock->turn = 0;
}
```

```
void lock(lock_t *lock) {
    int myturn = FetchAndAdd(&lock->ticket);
    while (lock->turn != myturn)
        ; // spin
}
```

```
void unlock(lock_t *lock) {
    lock->turn = lock->turn + 1;
}
```

| | Ticket | myturn | Turn |
|-----------------|--------|--------|------|
| Initial | 0 | | 0 |
| T1 calls lock | 1 | 0 | 0 |
| T1 enters CS | 1 | | 0 |
| T2 calls lock | 2 | 1 | 0 |
| T3 calls lock | 3 | 2 | 0 |
| T1 calls unlock | 3 | | 1 |
| T2 enters CS | 3 | | 1 |
| T2 calls unlock | 3 | | 2 |
| T3 enters CS | 3 | | 2 |
| T3 calls unlock | 3 | | 3 |

Once a thread is assigned its ticket value, it will be scheduled at some point in the future



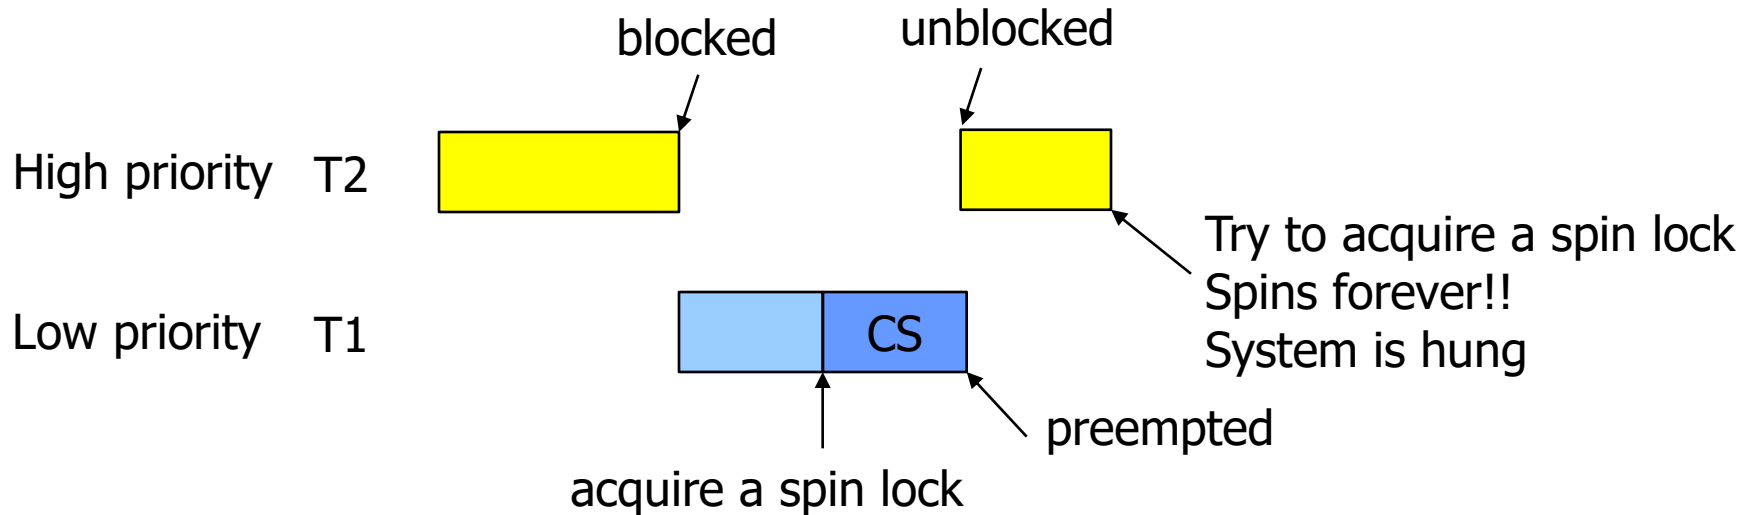
Ticket locks

Too Much Spinning

- Hardware-based locks are simple & work
 - But, can be quite inefficient
 - N threads on a single processor
 - Thread 0 is in a critical section w/ a lock, and unfortunately gets interrupted
 - Thread 1 is scheduled and tries to acquire the lock, and it begins to spin.
 - Thread 2 is scheduled and tries to acquire the lock, and it begins to spin.
 - ...
 - N-1 time slices may be wasted
- } waiting for the interrupted (lock-holding) thread to be run again
- Hardware support alone cannot solve the problem.
 - We'll need OS support too!

More Reason to Avoid Spinning

- **Priority-Driven Scheduling**



How about if just avoid the use of spin locks?

A Simple Approach: Just Yield

- **yield** system call moves the caller from the running state to the ready state, deschedules itself
- High context switch cost
 - each thread calling lock() will execute run-and-yield pattern before the thread holding the lock gets to run again
 - High context switch cost
- Starvation problem
 - A thread may get caught in an endless yield loop while other threads repeatedly enter and exit the critical section.

```
void init() {  
    flag = 0;  
}  
void lock() {  
    while (TestAndSet(&flag, 1) == 1)  
        yield(); // give up the CPU  
}  
void unlock() {  
    flag = 0;  
}
```

Spinlock: No Context Switching

Using Queues: Sleeping Instead Of Spinning

- Too much left to chance
 - The schedule determines who runs next; if it makes a bad choice – yield immediately or sleep
 - Let's get some control over who gets to acquire the lock next
- Need a queue to keep track of which threads are waiting to acquire the lock.
- A lock puts a caller to sleep if it tries to acquire a held lock and wakes it when the lock is free.

Lock With Queues, Test-and-set, Yield, And Wakeup

```
1  typedef struct __lock_t {
2      int flag;
3      int guard;
4      queue_t *q;
5  } lock_t;
6
7  void lock_init(lock_t *m) {
8      m->flag = 0;
9      m->guard = 0;
10     queue_init(m->q);
11 }
12
13 void lock(lock_t *m) {
14     while (TestAndSet(&m->guard, 1) == 1)
15         ; //acquire guard lock by spinning
16     if (m->flag == 0) {
17         m->flag = 1; // lock is acquired
18         m->guard = 0;
19     } else {
20         queue_add(m->q, gettid());
21         m->guard = 0;
22         park();
23     }
24 }
25
26 void unlock(lock_t *m) {
27     while (TestAndSet(&m->guard, 1) == 1)
28         ; //acquire guard lock by spinning
29     if (queue_empty(m->q))
30         m->flag = 0; // let go of lock; no one wants it
31     else
32         unpark(queue_remove(m->q)); // hold lock (for next thread!)
33     m->guard = 0;
34 }
```

Isn't that a
race condition?

Guard is a spin-lock around the flag
and queue manipulations
the time spent spinning is quite limited

Guard lock must be released before park()

put a calling thread to sleep

wake a waiting thread; pass the lock directly to the next
thread acquiring it; flag is not set to 0 in-between

wakeup/waiting race

- A thread (T0) will be about to park, assuming that it should sleep until the lock is no longer held.
- A switch at that time to the thread (T1) holding the lock.
- T1 releases the lock.
- The subsequent park by T0 would then sleep forever.
- Solaris solution
 - setpark()
 - a thread can indicate it is *about to* park
 - After this, if the thread is interrupted and another calls unpark before the park is called, parks returns immediately
- Another solution
 - pass the guard into the kernel.
 - kernel could take precautions to atomically release the lock and dequeue the running thread.

```
queue_add(m->q, gettid());  
setpark(); // new code  
m->guard = 0;  
park();
```

Different OS, Different Support

- Linux: **futex** (Fast Userspace MuTEX)
 - each futex has associated with it a specific **physical memory location**, as well as a **per-futex in-kernel queue**.
 - Provides atomic compare-and-block operation
 - futex is a lower-level construct
 - Used as building blocks for mutex, condition variables, semaphores
 - **futex_wait(address,expected)**
 - puts the calling thread to sleep if `mem[address] = expected`
 - If it is not equal, the call returns immediately.
 - **futex_wake(address)**
 - wakes one thread that is waiting on the queue.
 - Use a 32-bit integer
 - The leftmost bit (the +/- sign) tracks the lock state
 - 0: free, 1: locked
 - Remaining 31 bits: the number of waiters on the lock

Linux-based Futex Locks

```
1  void mutex_lock (int *mutex) {
2      int v;
3      /* Bit 31 was clear, we got the mutex (this is the fastpath) */
4      if (atomic_bit_test_set (mutex, 31) == 0)
5          return;
6      atomic_increment (mutex);
7      while (1) {
8          if (atomic_bit_test_set (mutex, 31) == 0) {
9              atomic_decrement (mutex);
10             return;
11         }
12         /* We have to wait now. First make sure the futex value
13            we are monitoring is truly negative (i.e. locked). */
14         v = *mutex;
15         if (v >= 0)
16             continue;
17         futex_wait (mutex, v);
18     }
19 }
20
21 void mutex_unlock (int *mutex) {
22     /* Adding 0x80000000 to the counter results in 0 if and only if
23        there are not other interested threads */
24     if (atomic_add_zero (mutex, 0x80000000))
25         return;
26
27     /* There are other threads waiting for this mutex,
28        wake one of them up. */
29     futex_wake (mutex);
30 }
```

Two-Phase Locks

- Spinning can be useful, particularly if the lock is about to be released
- **Two-Phase Lock: hybrid between spin-locks and yielding**
- First phase
 - the lock spins for a while, hoping that it can acquire the lock.
 - if the lock is not acquired, a second phase is entered
- Second phase
 - the caller is put to sleep
 - only woken up when the lock becomes free later.